

# Lean Proof Strategy Proposal for the Scheduling Rules Document

Stuart Swan, Platform Architect, Siemens EDA

16 June 2026

## 0. Goal and scope

The goal is to formalize the scheduling rules document's **mathematical proof core** in Lean: trace compatibility, witness construction, cutpoint correspondence, elaboration, and one-way liveness/deadlock preservation. The Lean proof would certify the formalized theorem stack under the stated assumptions, not every informal prose statement in the document.

The scheduling rules document's own framing is conditional: the formal results depend on assumptions, observation-boundary choices, witness-selection conditions, and liveness/progress premises.

The Lean development should therefore be organized around explicit theorem instances and assumption bundles, rather than hidden global axioms.

Notation and signature status. The Lean-like declarations in this strategy are intended to fix proof architecture, theorem dependencies, and object domains. Unless a declaration is explicitly identified as a final Lean signature, it should be read as schematic pseudo-Lean. In particular, channel, process, side, prefix, and bucket types must be unified in the concrete development through the selected `ComparisonInstance`, the side-parametric `SystemOfSide` map, and the shared proof-internal alphabet `C.Sigma`. This note is not a relaxation of the proof obligations; it prevents schematic notation such as `Channel`, `C.Sys_ref.Channel`, `Prefix`, and `InfiniteExecutionBucketed` from being mistaken for already type-checking Lean code.

Version alignment note. This revision updates the Lean proof strategy to align with the scheduling-rules v45 Appendix K changes, while retaining the earlier v42-aligned strategy as the base. It retains the v42 alignment points: (i) Appendix G/J separates the informational causal-dependency graph from the strict Appendix J required-order relation  $\prec_J$ ; (ii) Theorem J.1 produces an explicitly selected `Sys_ref` witness prefix and is organized along the fixed RTL-time ideal chain rather than by arbitrary  $\prec_J$ -minimal extension; (iii) Corollary J.1 and Appendix K/R cutpoint transfers use an explicit `RTL_B-to-source`

abstraction relation  $J.Abs$  and a corresponding implementation obligation  $J.AbsConf$ ; and (iv) Appendix K A5 is now the conservative serial-blocking source-liveness assumption, while Lemma K.2b supplies the serial-source dominance bridge from an RTL-matched / overlapped  $Sys\_ref$  witness deadlock to a serial-source deadlock.

The retained v42-specific alignment points are: (v) R2C combinational signal IO is now modeled as component-local, stutter-canonical fixed-point signal-observation units with selected observable slices  $Obs(C\_comb)$ , not as an unconditional per-cycle observable bucket; R2C matching is by component identity and local R2C occurrence ordinal, with cycle-locked components handled as an explicit stronger observation mode. (vi)  $ChannelEnabled$  is eligibility/availability for scheduler or arbiter selection, not same-cycle commit by itself; E4 is the safety/legality rule for committed transfers, while B2 supplies eventual selection and commit for continuously  $ChannelEnabled$  actions. (vii) hidden arbitration, grant, or back-pressure may not be smuggled in as an extra E4 availability conjunct at an ordinary boundary; it must be represented as fair selection among  $ChannelEnabled$  actions or as explicit  $Sys\_B^{elab}$  boundary structure. (viii) K.0/K.0a frontier reasoning now explicitly proves full  $NextFront$  minimality over synchronization, anchored signal/R2C, and message items, and rendezvous endpoint-request transfer relies on  $J.Abs/J.AbsConf$  raw endpoint-request facts. (ix) K.1 uses fair drain-normalization, rather than pure idle-stutter, to justify  $Exec\_post$  applicability at the relevant deadlock cutpoint.

Revision update note (v42 base plus v45 Appendix K refinement pass). This revision is a refinement pass against the v42-aligned strategy plus the scheduling-rules v45 Appendix K update. It retains the earlier v42 items and repairs several internal inconsistencies: (x) the Common Formal Definitions side-condition discharge record is given an explicit Lean-side audit structure tying WC, NB-CF, the Appendix L witness selector  $W$ ,  $Sys\_B^{elab}(E)$ ,  $J.AbsConf$ , and  $K\&$  to their discharge evidence (Section 1.1); (xi) B-faithfulness of each back-annotated capacity  $B(c)$  is recorded as a named tool/RTL-flow compliance hypothesis alongside E4 channel-realization conformance and  $J.AbsConf$  (Section 20); (xii) the v42 B3 modularity note is reflected: for a single explicit abstract boundary with ordinary E4 buffered or rendezvous semantics, the B3  $P\_push/P\_pop$  discharge is derivable from B2 plus E4, and Lean should prove that derivation while retaining B3 as the modular per-channel obligation checked against channel realizations and elaborated  $Sys\_B^{elab}$  boundary chains (Section 6); (xiii) the Appendix K terminated-process convention is made an explicit derived theorem together with its scope warning: waiting on a terminated peer is not, by itself, an internal message-passing deadlock in the K.1 sense, and intended termination/EOF/done protocols must be modeled explicitly or verified outside Appendix K (Section 18); (xiv) initial channel occupancy is aligned with A10: document-aligned

theorem instances use reset-empty occupancy ( $\text{initialOcc} = 0$ ), and accounted nonzero initial occupancy is flagged as an extension outside the main document’s empty-FIFO model (Section 7); (xv) RULE 1 / RULE 2 source well-formedness is recorded as part of the R2 obligation for proof-backed flows (Section 9); (xvi) ObsSigma gains the R2C component-local observation-unit slice already required by the document’s Definition I.Obs (Section 15); (xvii) JAssumptions gains the elaboration-projection well-formedness field that the J.1 proof sketch already consumed, with the corresponding KAssumptions import field (Sections 17 and 18); (xviii) ProcessStalledAtFront is unified with FrontierBlocked so the Appendix K layer uses the Section 13 conjunction form of “blocked on message passing” (Section 18); (xix) RequiredJOrder acyclicity is proved exactly as in the v42 Theorem J.1 argument, by the per-execution commit-cycle measure over the fixed legal prefix (Section 17); (xx) the A5 source-liveness field is identified definitionally with the serial execution-scoped predicate InternalMPDeadlockFreeOnSerialExecRef, and Theorem K.1 uses Lemma K.2b to bridge the witness-relative Exec\_ref\_liberal deadlock produced by Corollary J.1 to the serial-source deadlock needed to contradict A5 (Section 18); (xxi) START\_P / FINISH\_P anchors are read as per-reset-epoch dynamic occurrences, matching the v42 dynamic-reset treatment of direct inputs (Section 3); (xxii) Corollary L.4 (cadence-sensitive polling handled by isolation) is given an explicit named theorem alias (Section 14); (xxiii) a scope note records that the formalization targets the single-global-clock semantics of Appendices G–K (below); and (xxiv) the v45 Lemma K.2b proof spine is formalized as an eager-count reduction: serial reference executions are explicitly liberal/well-formed reference executions plus the no-eager restriction, and the one-step reduction is supported by endpoint-source-order prefix suppression, joint value/control consistency, occupancy/disabledness domination, no-third-party-enable monotonicity, and terminal-SCC persistence.

M5 alignment note (standing B1 in the main theorem stack). The latest scheduling-rules Appendix G/J wording treats B1 as a standing RTL\_B determinism assumption for the main Appendix G, Theorem J.1, Corollary J.1, and Appendix K theorem instances under the fixed stimulus. Lean therefore treats PostDeterministicSigmaTrace C as carried by the B-rule bundle used by JAssumptions and KAssumptions, not as an optional premise introduced only when uniqueness is invoked. B4/B5 remain separate additional premises used for  $\varepsilon$ -normalization,  $\varepsilon$ -quiescent endpoint reasoning, and uniqueness up to finite  $\varepsilon$ -stutter.

Terminology note on v42/v45 labels. References to “v42” in this strategy identify alignment decisions inherited unchanged from the earlier v42 pass. Unless explicitly stated otherwise, the theorem instances, assumptions, and appendix references in this revision target the attached v45 Scheduling Rules document.

Scope note (single clock). The formal core targeted by this strategy is the single-global-clock semantics of Appendices G–K of the scheduling-rules document. Appendix O of that document is an informative multi-clock sketch only. A multi-clock Lean development would additionally require the global event order or partial order over local-domain cycles, precise CDC commit semantics, cross-domain occupancy sampling rules, synchronizer fairness/metastability abstractions, and CDC progress obligations that Appendix O explicitly defers; those layers are out of scope for this strategy unless separately supplied.

The primary scheduling rules document on which this doc is based is available here:

[https://github.com/Stuart-Swan/Matchlib-Examples-Kit-For-Accellera-Synthesis-WG/blob/master/matchlib\\_examples/doc/catapult\\_user\\_view\\_scheduling\\_rules.pdf](https://github.com/Stuart-Swan/Matchlib-Examples-Kit-For-Accellera-Synthesis-WG/blob/master/matchlib_examples/doc/catapult_user_view_scheduling_rules.pdf)

The latest version of this Lean proof strategy document is available here:

[https://github.com/Stuart-Swan/Matchlib-Examples-Kit-For-Accellera-Synthesis-WG/blob/master/matchlib\\_examples/doc/Lean\\_proof\\_strategy.pdf](https://github.com/Stuart-Swan/Matchlib-Examples-Kit-For-Accellera-Synthesis-WG/blob/master/matchlib_examples/doc/Lean_proof_strategy.pdf)

The Matchlib examples kit that contains training slides, examples, etc., is available here:

<https://github.com/Stuart-Swan/Matchlib-Examples-Kit-For-Accellera-Synthesis-WG/tree/master>

## 1. Central comparison object

The proof should be parameterized by a selected comparison instance:

structure ComparisonInstance where

Sys\_ref : System

RTL\_B : System

-- Shared proof-internal observable alphabet for the Sys\_ref / RTL\_B comparison.

-- Both sides emit events in this same  $\Sigma$  schema.

Sigma : Alphabet

-- Optional outer DUT/testbench-visible projection alphabet.

-- In ordinary cases this may be the identity projection from Sigma.

Sigma\_DUT : Alphabet

sigmaDUTProjection : ProjectionMap Sigma Sigma\_DUT

env : ReactiveEnvironment Sigma

refChoice : SysRefChoice

-- The explicit Sys\_ref field and the reference-choice record must name the

-- same source reference system. This prevents theorem instances from using

```

-- C.Sys_ref in one definition and C.refChoice.Sys_ref in another while silently
-- referring to different systems.
sys_ref_consistent :
Sys_ref = refChoice.Sys_ref

-- Optional modular combinational model used by R2C.
-- The v42 R2C semantics are stated over explicit combinational observation
-- components. Each component has a selected observable slice Obs(C_comb), the
-- composed network has a unique  $\varepsilon$ -fixed point  $\mu_t$  in each compared cycle, and
--  $\Sigma$ -visible R2C events are emitted only as stutter-canonical fixed-point
-- signal-observation units. Matching is by component identity and local R2C
-- occurrence ordinal, not by global cycle number or by equality of unrelated
-- CycleBucket contents.
combModel? : Option CombComposition

assumptions : AssumptionBundle
witness : WitnessBundle
contracts : ModelingContracts

```

The comparison relation between `Sys_ref` and `RTL_B` is defined over the single shared proof-internal alphabet `C.Sigma`. The reference and post-HLS executions may have different internal  $\varepsilon$ -structure, bucket decomposition, or implementation state, but their observable events are encoded in the same  $\Sigma$  event schema before `CompatibleP / CompatibleSys` is applied.

`Sigma_DUT` is not a second post-side alphabet. It is only the chosen outer DUT/testbench-visible projection of `C.Sigma`. In elaborated cases, introduced proof-internal staged/bundle/join events may be hidden or accounted for by `sigmaDUTProjection /  $\Pi_{\text{elab}}$`  when projecting to `Sigma_DUT`, while the proof-internal compatibility theorem still ranges over `C.Sigma`.

`Sys_ref` is the chosen source reference model for the theorem instance:  
ordinary bounded-FIFO case: `Sys_ref = Sys_B`;  
elaborated case: `Sys_ref = Sys_Belab(E)` for a well-formed elaboration package `E`.

The `refChoice` field records why this `Sys_ref` is the prescribed source reference model, and `sys_ref_consistent` enforces that `C.Sys_ref` and `C.refChoice.Sys_ref` are the same system. This follows the updated Appendix Q direction: the formal

theorems do not require a unique algorithm that constructs  $\text{Sys\_B}^{\text{elab}}$ ; instead, the chosen  $\text{Sys\_ref}$  is part of the theorem instance, but the theorem instance may not contain two inconsistent reference systems.

### 1.1 Side-condition discharge record

The scheduling-rules v42 Common Formal Definitions add an explicit side-condition discharge record: before Theorem J.1 or the Appendix K liveness-preservation theorem is applied to a concrete design, the verification flow shall state, for each named side condition, whether the condition is inapplicable, satisfied by a simple syntactic/design restriction, discharged by tool/RTL evidence, or discharged by a source/RTL witness construction. The named conditions are WC, NB-CF, the Appendix L snooping / witness selector W,  $\text{Sys\_B}^{\text{elab}}$ , J.AbsConf, and  $\text{K}\epsilon$ . The record is not an additional semantic assumption beyond the premises already used by Appendices I–K; its purpose is to make theorem application auditable.

The Lean encoding should be a bookkeeping structure over the same evidence objects the strategy already defines, so that an instantiated theorem instance carries its own audit trail:

```
inductive SideConditionStatus
```

```
| inapplicable
```

```
| syntacticRestriction
```

```
| toolOrRTLEvidence
```

```
| witnessConstruction
```

```
structure SideConditionDischargeRecord
```

```
(C : ComparisonInstance)
```

```
(W : WitnessBundle C) where
```

```
-- WC — witness-compatible  $\epsilon$  choices / non-blocking outcomes.
```

```
wc_status : SideConditionStatus
```

```
wc_evidence :
```

```
wc_status  $\neq$  SideConditionStatus.inapplicable  $\rightarrow$ 
```

```
WC C W  $\vee$ 
```

```
( $\forall$  P, NBCF C P)  $\vee$ 
```

```
Nonempty (SnoopingWrapperAssumptions C W)
```

```
-- NB-CF — non-blocking control-flow neutrality, when used as the simple
```

```
-- sufficient condition that avoids nontrivial witness selection.
```

```
nb_cf_status : SideConditionStatus
```

nb\_cf\_evidence :

nb\_cf\_status  $\neq$  SideConditionStatus.inapplicable  $\rightarrow$

$\forall P, \text{NBCF } C \ P$

-- Appendix L snooping / witness selector W, when NB-CF does not hold but a

-- source execution must be selected to match the RTL\_B  $\varepsilon$ /latency choices.

snooping\_status : SideConditionStatus

snooping\_evidence :

snooping\_status  $\neq$  SideConditionStatus.inapplicable  $\rightarrow$

Nonempty (SnoopingWrapperAssumptions C W)

-- Sys\_B<sup>elab</sup> — elaborated source reference model, when ordinary Sys\_B

-- capacities at the outer boundaries do not represent the RTL realization.

elab\_status : SideConditionStatus

elab\_evidence :

C.refChoice.kind = ReferenceKind.elaborated  $\rightarrow$

$\exists E, C.\text{refChoice.elab?} = \text{some } E \wedge \text{WellFormedElaborationPackage } E$

-- J.AbsConf — RTL\_B cutpoint-abstraction conformance for the same selected

-- ( $\rho_{\text{post}}, \sigma_{\text{ref}}, w$ ) triple used by Corollary J.1 / Appendix K.

jabsconf\_status : SideConditionStatus

jabsconf\_evidence :

jabsconf\_status  $\neq$  SideConditionStatus.inapplicable  $\rightarrow$

Nonempty (JAbsConf C W)

-- K $\varepsilon$  — WFG-inert  $\varepsilon$  / no hidden micro-timing control decisions, required for

-- Appendix K liveness/deadlock preservation.

kepsilon\_status : SideConditionStatus

kepsilon\_evidence :

kepsilon\_status  $\neq$  SideConditionStatus.inapplicable  $\rightarrow$

A11\_Kepsilon\_WFGInert C

The record duplicates no semantics: each evidence field is one of the predicates or structures already defined in this strategy (WC, NBCF, SnoopingWrapperAssumptions, the elaboration package well-formedness obligations, JAbsConf, A11\_Kepsilon\_WFGInert). The WC field is intentionally guarded like the other optional audit entries: if wc\_status is inapplicable, no WC/NB-CF/snooping discharge is demanded by the audit object; otherwise the alternatives must cover the full proof-internal C.Sigma comparison, not

merely an outer Sigma\_DUT projection. Theorem instances for J.1 / K.1 may take the record as a convenience parameter, or the verification flow may construct it from the JAssumptions / KAssumptions bundles after the fact. Either way, the Appendix J Remark 2 side-condition audit — whether Sys\_ref is ordinary Sys\_B or elaborated Sys\_B^elab(E), whether failed non-blocking polls are covered by NB-CF or by a witness selector W satisfying Appendix L, whether the selected cutpoints discharge J.AbsConf for the same selected triple, and whether any liveness/deadlock use also discharges Kε and the Appendix K/N progress assumptions — is represented by an explicit object rather than left implicit.

## 2. Foundation layer: bucketed executions, not flat traces

The primary execution model should be **cycle-bucketed**, not a flat sequence of eps | tick | sigma e. The reason is that the document allows multiple same-cycle Σ-events without imposing artificial order among independent same-cycle events. The earlier review correctly noted that a flat trace would either impose a fake linearization or require an added bucket abstraction anyway.

Use:

structure State where

cycle : Nat  
 epsPos : Nat  
 store : Store

with explicit advancement rules:

- Intra-cycle ε-normalization steps advance epsPos only.
- Same-cycle Σ-events advance epsPos only.
- Clock ticks advance cycle and reset epsPos.
- Cycle-advancing silent progress is represented at the bucket level: a
- CycleBucket with sigmaEvents = ∅ may still consume a real clock cycle and
- may represent HLS-inserted FSM latency, arbitration/bookkeeping latency,
- failed non-blocking polls, or pipeline advances that emit no Σ-event.

Define buckets:

structure CycleBucket (Σ : Type) where

cycle : Nat  
 startState : State  
 epsPrefix : List EpsStep



```

-- The  $\Sigma$ -events committed in this clock cycle.
-- This is an unordered finite set. A CycleBucket does not carry an intrinsic
-- sequence or bucket-local partial order among same-cycle events.
sigmaEvents : Finset  $\Sigma$ 
epsSuffix : List EpsStep
endState : State
endQuiescent : Prop
def ExecutionBucketed ( $\Sigma$  : Type) :=
List (CycleBucket  $\Sigma$ )
def InfiniteExecutionBucketed ( $\Sigma$  : Type) :=
Nat  $\rightarrow$  CycleBucket  $\Sigma$ 

```

Bucket well-formedness predicates, not the Finset representation alone, enforce per-cycle scalar constraints such as at most one Push\_c and at most one Pop\_c per channel in a cycle. Required ordering is not stored in CycleBucket. It is derived globally from observable units and from the Appendix J required-order relation  $\prec_J$ , not from bucket membership or bucket-local sequencing.

The scheduling-rules v42 text intentionally separates two relations: causal-dependency graph ( $\rightarrow$ ), an information-flow graph over observable IO actions that is useful for explaining functional dependencies but need not be acyclic in every legal execution; and RequiredJOrder /  $\prec_J$ , the well-founded strict order over observable units used by the Appendix J finite-ideal induction and by CompatibleP / CompatibleSys.

This provides the right foundation for:

- $\varepsilon$ -quiescent cutpoints;
- rendezvous same-cycle commits;
- R2C component-local fixed-point signal-observation units;
- J.1 ideal induction over observable units and the mandated  $\prec_J$  relation;
- B2/B3 temporal reasoning over infinite bucketed executions.

Important separation: ExecutionBucketed is the semantic carrier for state advancement,  $\varepsilon$ -quiescence, occupancy updates, R2C fixed points, R2C stutter-canonical unit emission, and fairness/progress reasoning. It is not itself the cross-model trace-compatibility predicate.

In particular, system-level compatibility  $\tau_{\text{ref},\text{system}} \approx_{\text{sys}} \tau_{\text{post},\text{system}}$  must not be

defined as equality of CycleBucket lists, equality of bucket cycle numbers, or bucket-by-bucket equality of sigmaEvents. The document's Appendix J compatibility relation is an ordered cross-model observable-compatibility predicate. It preserves the explicitly required  $\leq_J$  constraints, atomic multi-endpoint units, E1–E5 obligations, FIFO/synchronization/signal-anchor obligations, and matched value labels, while allowing independent observable units to commute or to be grouped into different cycle buckets in Sys\_ref and RTL\_B. Same-cycle membership in a CycleBucket is therefore a timing fact for one execution, not a source of intrinsic order.

Thus:

bucket equality may be required inside explicitly atomic units, such as rendezvous Push/Pop pairs and SyncChannel/ac\_sync handshakes;

bucket decomposition may be used to prove local state-transition and enablement facts;

but CompatibleSys /  $\approx_{\text{sys}}$  is defined over observable units and the mandated  $\leq_J$  relation, not over identical bucket decompositions.

### 3. Events, operation instances, and source order

Define dynamic operation universes:

inductive MsgKind

| Push

| Pop

| PushNB

| PopNB

def MsgKind.isBlocking : MsgKind → Bool

| MsgKind.Push => true

| MsgKind.Pop => true

| MsgKind.PushNB => false

| MsgKind.PopNB => false

structure DynMsgInstance where

proc : Process

channel : Channel

kind : MsgKind -- Push, Pop, PushNB, PopNB

sourceLoc : SourceLoc

instanceId : Nat

```
def Blocking (q : DynMsgInstance) : Prop :=
  q.kind.isBlocking = true
```

```
def NonBlocking (q : DynMsgInstance) : Prop :=
  q.kind.isBlocking = false
```

```
structure FrontierItem where
  proc      : Process
  itemKind  : FrontierItemKind
  sourceLoc : SourceLoc
  instanceld : Nat
```

Core source-order objects are witness-specific. They range only over dynamic operation instances that actually occur in the compared execution / witness. Untaken branches and unexecuted loop iterations are not members of this dynamic universe.

```
Qexec : Process → Set DynMsgInstance
QOmega : Process → Set DynMsgInstance
Omega : Process → Set FrontierItem
```

$Qexec, P$  is the execution-level dynamic-instance universe. It includes executed failed non-blocking polls that may produce no  $\Sigma$ -event.

$\Omega, P$  is the frontier/order universe used by Front, NextFront, WFG, E3/E5, BoundaryReq attribution, and source-order reasoning. Failed NB-only executed polls are not included in  $\Omega, P$  merely because they executed.

$Q\Omega, P$  denotes the executed dynamic message-operation-instance slice induced by  $\Omega, P$ . Since  $Qexec, P$  is a set of  $DynMsgInstance$  and  $\Omega, P$  is a set of  $FrontierItem$ , this is not a direct set intersection. It is defined through the message-instance projection from frontier items:

```
QOmega, P :=
{ q : DynMsgInstance |
  q ∈ Qexec, P ∧
  ∃ x : FrontierItem,
  x ∈ Ω, P ∧
```

MsgInstanceOfFrontierItem C P x = some q }

```
def MsgInstanceOfFrontierItem
(C : ComparisonInstance)
(P : Process)
(x : FrontierItem) : Option DynMsgInstance :=
...
```

-- R2/E2 signal anchors include both explicit synchronization calls and the  
-- process-boundary synchronization events START\_P and FINISH\_P. This avoids  
-- making signal reads before the first explicit wait(), or signal writes after  
-- the last explicit wait(), ill-typed.

inductive SyncAnchor

```
| start (P : Process)
| call (s : SyncCall)
| finish (P : Process)
```

def SyncAnchorClock

```
(C : ComparisonInstance)
(W : WitnessBundle C)
(a : SyncAnchor) : Cycle :=
match a with
| SyncAnchor.start P => StartClock C W P
| SyncAnchor.call s => SyncClock C W s
| SyncAnchor.finish P => FinishClock C W P
```

-- A SyncAnchor may be used by pred\_S / succ\_S only when the corresponding  
-- synchronization event is actually present on the dynamic path represented by  
-- the witness universe. In particular, START\_P is not an implicit free anchor:  
-- it is usable only when the modeling convention includes START\_P as a real  
-- synchronization event in S. FINISH\_P is usable only for an actually  
-- terminating process whose FINISH\_P event is realized.

def SyncAnchorRealizedInUniverse

```
(C : ComparisonInstance)
(P : Process)
(W : WitnessBundle C)
(Ω : Set FrontierItem)
(a : SyncAnchor) : Prop :=
```

```

match a with
| SyncAnchor.start P' =>
P' = P ∧
∃ x,
x ∈ Ω ∧
StartSyncFrontierItem C P x ∧
ExecutedInWitness C W P x
| SyncAnchor.call s =>
SyncCallOwnedByProcess C P s ∧
∃ x,
x ∈ Ω ∧
SyncCallFrontierItem C P s x ∧
ExecutedInWitness C W P x
| SyncAnchor.finish P' =>
P' = P ∧
ProcessTerminatesInWitness C W P ∧
∃ x,
x ∈ Ω ∧
FinishSyncFrontierItem C P x ∧
ExecutedInWitness C W P x

```

```

structure DynamicUniverse
(C : ComparisonInstance)
(P : Process)
(W : WitnessBundle C) where
Q_exec : Set DynMsgInstance
Ω : Set FrontierItem
Q_Ω : Set DynMsgInstance

```

```

qomega_characterization :
∀ q,
q ∈ Q_Ω ↔
q ∈ Q_exec ∧
∃ x,
x ∈ Ω ∧
MsgInstanceOfFrontierItem C P x = some q

```

```

message_frontier_items_have_exec_instance :

```

$\forall x q,$   
 $x \in \Omega \rightarrow$   
 $\text{MsgInstanceOfFrontierItem } C \ P \ x = \text{some } q \rightarrow$   
 $q \in Q_{\text{exec}} \wedge q \in Q_{\Omega}$

$\text{qomega\_items\_have\_frontier\_item} :$   
 $\forall q,$   
 $q \in Q_{\Omega} \rightarrow$   
 $\exists x,$   
 $x \in \Omega \wedge$   
 $\text{MsgInstanceOfFrontierItem } C \ P \ x = \text{some } q$

$\text{executed\_only} :$   
 $\forall x,$   
 $x \in \Omega \rightarrow$   
 $\text{ExecutedInWitness } C \ W \ P \ x$

$\text{leP} :$   
 $\{x : \text{FrontierItem} \ // \ x \in \Omega\} \rightarrow$   
 $\{y : \text{FrontierItem} \ // \ y \in \Omega\} \rightarrow$   
 $\text{Prop}$

$\text{leP\_partial\_order} :$   
 $\text{PartialOrder } \text{leP}$

$\text{prefS} :$   
 $\text{SyncCall} \rightarrow$   
 $\text{Set } \{x : \text{FrontierItem} \ // \ x \in \Omega\}$   
 $\text{suffS} :$   
 $\text{SyncCall} \rightarrow$   
 $\text{Set } \{x : \text{FrontierItem} \ // \ x \in \Omega\}$

-- R2/E2 signal-anchor maps for sequential-process signal IO.  
--  $\text{pred\_S}$  is defined for ordinary sequential-process signal Reads.  
--  $\text{succ\_S}$  is defined for ordinary sequential-process signal Writes.  
-- The anchor type is  $\text{SyncAnchor}$ , not  $\text{SyncCall}$ , because  $\text{START\_P}$  may be the  
-- closest preceding anchor for an initial Read and  $\text{FINISH\_P}$  may be the closest

```

-- succeeding anchor for a terminal Write.
pred_S :
{x : FrontierItem // x ∈ Ω} →
Option SyncAnchor
succ_S :
{x : FrontierItem // x ∈ Ω} →
Option SyncAnchor

read_anchor_defined :
∀ x,
SequentialSignalReadItem x.val →
pred_S x ≠ none ∧
succ_S x = none
write_anchor_defined :
∀ x,
SequentialSignalWriteItem x.val →
succ_S x ≠ none ∧
pred_S x = none
pred_S_anchor_realized :
∀ x a,
SequentialSignalReadItem x.val →
pred_S x = some a →
SyncAnchorRealizedInUniverse C P W Ω a
succ_S_anchor_realized :
∀ x a,
SequentialSignalWriteItem x.val →
succ_S x = some a →
SyncAnchorRealizedInUniverse C P W Ω a
pred_S_is_closest_preceding_anchor :
∀ x a,
SequentialSignalReadItem x.val →
pred_S x = some a →
ClosestPrecedingDynamicSyncAnchor C P W a x.val
succ_S_is_closest_succeeding_anchor :
∀ x a,
SequentialSignalWriteItem x.val →
succ_S x = some a →
ClosestSucceedingDynamicSyncAnchor C P W a x.val

```

`signal_read_ordered_at_pred_S :`  
 $\forall x a,$   
`SequentialSignalReadItem x.val  $\rightarrow$`   
`pred_S x = some a  $\rightarrow$`   
`SignalItemClock C W x.val = SyncAnchorClock C W a`

`signal_write_ordered_at_succ_S :`  
 $\forall x a,$   
`SequentialSignalWriteItem x.val  $\rightarrow$`   
`succ_S x = some a  $\rightarrow$`   
`SignalItemClock C W x.val = SyncAnchorClock C W a`

`MsgInstanceOfFrontierItem` is defined only for frontier items that correspond to message-operation dynamic instances. Non-message frontier items, such as pure synchronization or signal-anchor items, map to none. Signal-anchor items instead use the explicit `pred_S` / `succ_S` fields above. For sequential processes, `pred_S` places each ordinary signal Read at its closest preceding `SyncAnchor`, and `succ_S` places each ordinary signal Write at its closest succeeding `SyncAnchor`. A `SyncAnchor` may be `START_P`, an explicit synchronization call, or `FINISH_P`. These anchors, not raw lexical position, determine the signal item's formal position for  $\leq_P$  /  $<_P$ , `Done_P`, `Remain_P`, `NextFront`, `E2SignalAnchorOrder`, and `CompatibleP` / `CompatibleSys` reasoning.

Per-reset-epoch anchor occurrences. In the scheduling-rules document, a new `START_P` synchronization event is emitted each time `P` exits reset, and v42 explicitly permits dynamic resets during normal HW operation, including resetting of direct-input-controlled regions. The constructors `SyncAnchor.start P` and `SyncAnchor.finish P` are therefore schematic: in the concrete development each must resolve to a particular dynamic occurrence in the witness universe — for example by carrying a reset-epoch / occurrence index, or by being identified with the realized start/finish frontier item returned by `SyncAnchorRealizedInUniverse`. `StartClock C W P` and `FinishClock C W P` are then read per dynamic occurrence, so that a Read anchored to `START_P` in a later reset epoch is anchored to that epoch's `START_P` event, not to the first one. This matches the `DirectInputContract` reset-epoch scoping in Section 5.

The field `Q_Ω` is therefore not a second independent universe: it is exactly the set of executed dynamic message-operation instances  $q \in Q_{\text{exec}}$  for which some



message frontier item  $x \in \Omega$  satisfies  $\text{MsgInstanceOfFrontierItem } C \ P \ x = \text{some } q$ . This keeps  $Q\_exec$ ,  $\Omega$ ,  $Q\_Q$ , and the signal-anchor universe type-distinct while making their relationships explicit for  $\text{BoundaryReq}$ ,  $\text{Front}$ ,  $\text{NextFront}$ ,  $\text{WFG}$ ,  $\text{E2}$ ,  $\text{E3/E5}$  issue-order reasoning, and  $\text{R2/R2C}$  signal visibility.

Thus source order  $\leq_P$  is not a static total order over syntactic program statements. It is a partial order over the witness-bundle-specific dynamic frontier-item universe  $\Omega$ . The selector component  $W.\text{selector}$  determines  $\varepsilon$ -modeled source choices, while  $W.\text{nbMap}$  records the dynamic-instance correspondence for failed non-blocking polls that produce no  $\Sigma$ -event.

Non-blocking calls need explicit modeling:

- failed  $\text{PushNB}$  /  $\text{PopNB}$  polls are executed dynamic instances in  $Q\_exec$ , but produce no  $\Sigma$ -event;
- successful non-blocking calls produce the corresponding committed  $\text{Push}_c(v)$  /  $\text{Pop}_c(v)$  event;
- repeated NB calls are distinct dynamic source instances even if a request signal remains physically asserted.

This matches the document's issue/commit treatment of NB calls.

```
structure NonBlockingInstanceWF
(C : ComparisonInstance)
(P : Process)
(W : WitnessBundle C)
(U : DynamicUniverse C P W) where
```

```
nb_instances_are_nonpersistent :
 $\forall q,$ 
 $q \in U.Q\_exec \rightarrow$ 
 $\text{NonBlocking } q \rightarrow$ 
 $\neg \text{PersistentBlockingRequestInstance } C \ P \ q$ 
```

```
nb_reexecution_distinct_instances :
 $\forall q_1 \ q_2 \ t_1 \ t_2,$ 
 $q_1 \in U.Q\_exec \rightarrow$ 
 $q_2 \in U.Q\_exec \rightarrow$ 
 $\text{NonBlocking } q_1 \rightarrow$ 
 $\text{NonBlocking } q_2 \rightarrow$ 
```

SameSourceNBPollSite  $q_1 q_2 \rightarrow$

ExecutedAtCycle C W  $q_1 t_1 \rightarrow$

ExecutedAtCycle C W  $q_2 t_2 \rightarrow$

$t_1 \neq t_2 \rightarrow$

$q_1 \neq q_2$

nb\_physical\_assertion\_does\_not\_identify\_instances :

$\forall q_1 q_2 t_1 t_2,$

$q_1 \in U.Q\_exec \rightarrow$

$q_2 \in U.Q\_exec \rightarrow$

NonBlocking  $q_1 \rightarrow$

NonBlocking  $q_2 \rightarrow$

SameEndpoint  $q_1 q_2 \rightarrow$

EndpointRequestContinuouslyAsserted C W  $q_1.channel t_1 t_2 \rightarrow$

ExecutedAtCycle C W  $q_1 t_1 \rightarrow$

ExecutedAtCycle C W  $q_2 t_2 \rightarrow$

$t_1 \neq t_2 \rightarrow$

$q_1 \neq q_2$

nb\_not\_fair\_action :

$\forall q,$

$q \in U.Q\_exec \rightarrow$

NonBlocking  $q \rightarrow$

$\neg \exists h, \text{FairAction.push } q h \vee \text{FairAction.pop } q h$

The last field is only a type-level sanity condition for the fairness layer:  
B2 fairness ranges over persistent blocking Push/Pop requests and real  
synchronization calls. A successful NB call may still contribute the ordinary  
committed Push\_c / Pop\_c  $\Sigma$ -event, but the NB call instance itself does not become  
a persistent FairAction.

## 4. Reactive environment and modeling contracts

### 4.1 Reactive environment

“Fixed stimulus” should be modeled as a causal global strategy over  $\Sigma$ -history, not merely  
as a fixed stream of per-process inputs.

```

structure ReactiveEnvironment ( $\Sigma$  : Type) where
  inputAt : SigmaHistory  $\Sigma$   $\rightarrow$  ExternalInputs
  causal :
     $\forall h_1 h_2,$ 
      SamePrefix  $h_1 h_2 \rightarrow$ 
        inputAt  $h_1 =$  inputAt  $h_2$ 

```

This directly addresses the earlier issue that reactive stimulus must be interpreted as the same global  $\Sigma$ -causal behavior, not a function of local process history.

## 4.2 Modeling contracts

Direct-input contracts, array-access contracts, and latency-sensitive local-protocol isolation contracts should be parallel top-level modeling contracts, not partly folded into RRules and partly left separate.

The scheduling-rules document treats cadence-sensitive retry/timeout polling as latency-sensitive local protocol logic. In the Lean development this should be an explicit modeling contract, not a hidden theorem-side convention. Failed non-blocking polls still remain  $\varepsilon$ -steps with no committed Push\_c / Pop\_c  $\Sigma$ -event; the contract instead states that a polling site whose externally visible behavior depends on the number or cadence of failed polls is outside the ordinary latency-insensitive Appendix I/J source core unless it has been isolated, snoop-aligned, or supplied with an equivalent explicit cycle-accurate local protocol model at the chosen boundary.

```

inductive LocalProtocolHandling
| ordinaryLatencyInsensitiveCore
| isolatedBoundary
| snoopAligned
| explicitCycleAccurateModel

```

```

structure LatencySensitiveLocalProtocolContract

```

```

(C : ComparisonInstance) where

```

```

  handlingOf :

```

```

    Process  $\rightarrow$  SourceLoc  $\rightarrow$  LocalProtocolHandling

```

```

  cadence_sensitive_nb_sites_handled :

```

```

     $\forall P$  loc,

```

```

    CadenceSensitiveNBPollingSite C P loc  $\rightarrow$ 

```

```

    handlingOf P loc = LocalProtocolHandling.isolatedBoundary  $\vee$ 

```

handlingOf P loc = LocalProtocolHandling.snoopAligned V  
handlingOf P loc = LocalProtocolHandling.explicitCycleAccurateModel

ordinary\_core\_excludes\_unhandled\_cadence :

∀ P loc,

handlingOf P loc = LocalProtocolHandling.ordinaryLatencyInsensitiveCore →

¬ CadenceSensitiveNBPollingSite C P loc

isolated\_boundary\_has\_stable\_observable\_interface :

∀ P loc,

handlingOf P loc = LocalProtocolHandling.isolatedBoundary →

LatencyInsensitiveOrExplicitlyModeledBoundary C P loc

snoop\_aligned\_sites\_have\_admissible\_witness :

∀ P loc,

handlingOf P loc = LocalProtocolHandling.snoopAligned →

SnoopAlignedLocalProtocolSite C P loc

cycle\_accurate\_sites\_have\_explicit\_model :

∀ P loc,

handlingOf P loc = LocalProtocolHandling.explicitCycleAccurateModel →

ExplicitCycleAccurateLocalProtocolModel C P loc

structure ModelingContracts (C : ComparisonInstance) where

directInputs : DirectInputContract C

arrays : ArrayAccessMappingRules C

latencySensitiveLocalProtocols : LatencySensitiveLocalProtocolContract C

This keeps RRules close to the document's actual R-rule set and treats direct-input, external-array, and latency-sensitive local-protocol obligations as environment/modeling contracts. In particular, timeout/retry/poll-arbiter cadence is not added to the ordinary  $\Sigma$  alphabet; it is either local to an isolated block, selected by a snooping witness for that block, or covered by an explicit cycle-accurate local protocol model.

## 5. Direct-input contract and late sampling

The direct-input contract should contain only primitive environmental stability/update obligations. The fact that late sampling preserves the logical anchor should be a **derived theorem**, not a separate hypothesis.

The document says `hls_direct_input` permits late physical sampling because the environment holds the signal stable after reset, while `hls_direct_input_sync` permits controlled updates only at the designated synchronization handshake. The document also states that logical signal Read/Write events are timestamped at the E2 anchor, not at any later physical RTL sampling cycle; late sampling is  $\varepsilon$ -internal.

Define:

The direct-input stability rule applies only within a single reset epoch. The scheduling document explicitly permits dynamic resetting during normal HW operation, including resetting of direct-input-controlled regions.

```
def NoRelevantResetBetween
(C : ComparisonInstance)
(sig : Signal)
(t1 t2 : Cycle) : Prop :=
¬ ∃ r,
t1 < r ∧
r ≤ t2 ∧
ReceiverResetAffectsSignal C sig r
```

```
structure DirectInputContract (C : ComparisonInstance) where
stableWithinResetEpoch :
∀ sig t1 t2,
DirectInput sig →
AfterAllReceiversOutOfReset C sig t1 →
t1 ≤ t2 →
NoRelevantResetBetween C sig t1 t2 →
InputValue C.env sig t1 = InputValue C.env sig t2
syncControlledUpdateOnlyAtSync :
∀ sig s t,
DirectInputSyncControlled sig s →
ValueChangesAt C.env sig t →
SyncHandshakeOccursAt C s t
```

Then prove:

```
theorem lateSamplingPreservesLogicalAnchor
(C : ComparisonInstance)
(hDI : DirectInputContract C)
```

```

(sig : Signal)
(anchor sampleT : Cycle)
(hIsDI : DirectInput sig)
(hOutOfReset :
AfterAllReceiversOutOfReset C sig anchor)
(hWindow : DirectInputAnchorWindow C sig anchor sampleT)
(hLe : anchor ≤ sampleT)
(hNoReset :
NoRelevantResetBetween C sig anchor sampleT) :
InputValue C.env sig anchor =
InputValue C.env sig sampleT :=
hDI.stableWithinResetEpoch
sig
anchor
sampleT
hIsDI
hOutOfReset
hLe
hNoReset

```

Bucket-location rule:

For sequential processes, the logical direct-input  $\text{Read}(\text{sig}, \text{val})$  /  $\text{Write}(\text{sig}, \text{val})$  event belongs only to the E2 anchor bucket  $\beta_{\text{anchor}}$ . A physical late sample, if modeled, is an  $\varepsilon$ -step in a later bucket. The late  $\varepsilon$ -step creates no new  $\Sigma$ -event.  $\text{lateSamplingPreservesLogicalAnchor}$  is used only to justify the value recorded in the logical sequential anchor bucket.

This sequential anchoring rule is not applied to combinational-process signal IO. Combinational processes do not have  $\text{pred\_S}$  /  $\text{succ\_S}$  synchronization anchors. Their observable direct-input Read/Write labels, when any direct input participates in combinational signal IO, are included only inside an emitted R2C fixed-point signal-observation unit for the relevant combinational component. The unit is emitted according to the v42 R2C unit-generation convention: first defined observation, later observable-slice value changes, or an explicitly modeled observer/sampler requirement. Intermediate physical samples, delta-cycle reads/writes, redundant re-evaluations, transient values before the fixed point, and R2C stutter cycles that emit no unit are  $\varepsilon$ /stutter and are not separately  $\Sigma$ -visible.

-- These anchor predicates are derived from the per-process DynamicUniverse  
-- fields  $\text{pred\_S}$  and  $\text{succ\_S}$ . For a sequential read item  $x$ ,

-- SequentialSignalReadAnchoredAt C (LogicalReadEvent sig val) anchor means  
 -- that x is the corresponding read frontier item and pred\_S x = some anchor.  
 -- For a sequential write item x,  
 -- SequentialSignalWriteAnchoredAt C (LogicalWriteEvent sig val) anchor means  
 -- that x is the corresponding write frontier item and succ\_S x = some anchor.  
 -- Thus the anchor used below is the event's actual E2/R2 anchor, not an  
 -- unconstrained universally quantified synchronization point.

structure DirectInputBucketWF (C : ComparisonInstance) where

sequential\_read\_recorded\_only\_at\_pred\_anchor :

$\forall$  sig val anchor sampleT  $\beta$ ,

DirectInput sig  $\rightarrow$

SequentialSignalReadEvent (LogicalReadEvent sig val)  $\rightarrow$

LogicalReadEvent sig val  $\in$  BucketSigmaEvents C  $\beta \rightarrow$

SequentialSignalReadAnchoredAt C (LogicalReadEvent sig val) anchor  $\rightarrow$

DirectInputAnchorWindow C sig anchor sampleT  $\rightarrow$

$\beta = \text{AnchorBucket C anchor}$

The late-sample value agreement field is reset-epoch local. It is derived from stableWithinResetEpoch and may be used only when no relevant dynamic reset occurs between the logical anchor cycle and the physical sample cycle.

sequential\_write\_recorded\_only\_at\_succ\_anchor :

$\forall$  sig val anchor  $\beta$ ,

DirectInput sig  $\rightarrow$

SequentialSignalWriteEvent (LogicalWriteEvent sig val)  $\rightarrow$

LogicalWriteEvent sig val  $\in$  BucketSigmaEvents C  $\beta \rightarrow$

SequentialSignalWriteAnchoredAt C (LogicalWriteEvent sig val) anchor  $\rightarrow$

$\beta = \text{AnchorBucket C anchor}$

late\_sample\_is\_epsilon\_only :

$\forall$  sig anchor sampleT,

DirectInputAnchorWindow C sig anchor sampleT  $\rightarrow$

PhysicalLateSample C sig sampleT  $\rightarrow$

EpsilonStepOnly C sig sampleT  $\wedge$

$\neg \exists$  val  $\beta$ ,

PhysicalSampleEvent sig val sampleT  $\in$  BucketSigmaEvents C  $\beta$

`late_sample_value_matches_anchor :`  
 $\forall \text{ sig anchor sampleT,}$   
`DirectInput sig  $\rightarrow$`   
`AfterAllReceiversOutOfReset C sig anchor  $\rightarrow$`   
`DirectInputAnchorWindow C sig anchor sampleT  $\rightarrow$`   
`anchor  $\leq$  sampleT  $\rightarrow$`   
`NoRelevantResetBetween C sig anchor sampleT  $\rightarrow$`   
`InputValue C.env sig anchor = InputValue C.env sig sampleT`

This explicitly reconciles the bucket model with direct-input late sampling: CompatibleP and CompatibleSys compare the logical  $\Sigma$ -visible direct-input Read/Write events at the appropriate logical unit — the E2 anchor bucket for sequential-process signal IO, or the emitted component-local R2C fixed-point signal-observation unit for combinational-process signal IO — never the later physical sample micro-step.

Combinational-process signal IO uses R2C, not the direct-input late-sampling contract. The R2C rule is therefore a general combinational-process well-formedness obligation. It applies to every combinational-process signal Read/Write event in the compared  $\Sigma$  trace, regardless of whether the signal is marked DirectInput.

`structure R2CObservationUnitWF (C : ComparisonInstance) where`  
`selected_observable_slice_defined :`  
 $\forall \text{ comp,}$   
`R2CComponent C comp  $\rightarrow$`   
 $\exists \text{ obs, ObservableSlice C comp obs}$

`unit_emitted_only_at_canonical_occurrence :`  
 $\forall \text{ comp unit t,}$   
`R2CUnitOfComponent C comp unit  $\rightarrow$`   
`R2CUnitEmittedAt C unit t  $\rightarrow$`   
`FirstDefinedR2CObservation C comp t  $\vee$`   
`ObservableSliceChangedSincePreviousUnit C comp t  $\vee$`   
`ExplicitObserverRequiresFreshR2CUnit C comp t  $\vee$`   
`CycleLockedR2CComponent C comp`

`unit_uses_unique_fixed_point :`  
 $\forall \text{ comp unit t } \mu,$   
`R2CUnitOfComponent C comp unit  $\rightarrow$`



$R2CUnitEmittedAt\ C\ unit\ t \rightarrow$   
 $CombFixedPointAtComponentCycle\ C\ comp\ t\ \mu \rightarrow$   
 $\forall e,$   
 $e \in unit.events \rightarrow$   
 $EventReadsWritesStore\ e\ \mu$

$unit\_not\_split :$   
 $\forall comp\ unit\ e_1\ e_2,$   
 $R2CUnitOfComponent\ C\ comp\ unit \rightarrow$   
 $e_1 \in unit.events \rightarrow$   
 $e_2 \in unit.events \rightarrow$   
 $SameObservableUnit\ C\ .ref\ unit\ e_1\ e_2 \wedge SameObservableUnit\ C\ .post\ unit\ e_1\ e_2$

$stutter\_cycles\_emit\_no\_unit :$   
 $\forall comp\ t,$   
 $R2CStutterCycle\ C\ comp\ t \rightarrow$   
 $\neg \exists unit, R2CUnitOfComponent\ C\ comp\ unit \wedge R2CUnitEmittedAt\ C\ unit\ t$

$matched\_by\_component\_local\_ordinal :$   
 $\forall comp\ k\ u\_ref\ u\_post,$   
 $KthR2CUnit\ C\ .ref\ comp\ k\ u\_ref \rightarrow$   
 $KthR2CUnit\ C\ .post\ comp\ k\ u\_post \rightarrow$   
 $R2CUnitValuesAgree\ C\ comp\ u\_ref\ u\_post$

$multi\_input\_component\_has\_anchor\_or\_confluence\_proof :$   
 $\forall comp,$   
 $MultInputR2CComponent\ C\ comp \rightarrow$   
 $SingleExplicitAnchoringDomain\ C\ comp\ V$   
 $AllAdmissibleInterleavingsInduceSameR2CUnitSequence\ C\ comp\ V$   
 $R2CComponentHiddenBehindExplicitSynchronizationOrTransactor\ C\ comp$

## 6. Temporal logic layer for B2/B3/B4/B5/B6

B2 and B3 should not be opaque Prop fields. They are temporal assumptions over infinite executions.

The document defines B2 as weak fairness over enabled actions, but v42 clarifies the enabled/selected/committed split. For message operations, ChannelEnabled means that

the abstract boundary permits the operation to be selected; it does not by itself mean the transfer has already committed in that cycle. B2 therefore says: if an action remains continuously ChannelEnabled / eligible from some cycle onward, the scheduler or arbiter eventually selects it, and the selected action commits in the selected cycle. B3 remains the P\_push(c) / P\_pop(c) channel-progress obligation; it may be needed to discharge full/empty/rendezvous-blocked conditions before a pending transfer becomes ChannelEnabled.

Define temporal operators:

```
def AlwaysFrom
  (τ : InfiniteExecutionBucketed Σ)
  (t : Nat)
  (P : Nat → Prop) : Prop :=
  ∀ n, t ≤ n → P n
```

```
def EventuallyFrom
  (τ : InfiniteExecutionBucketed Σ)
  (t : Nat)
  (P : Nat → Prop) : Prop :=
  ∃ n, t ≤ n ∧ P n
```

Concrete fairness actions:

```
-- B2 fairness ranges only over scheduler-selectable operations.
-- For message operations, this means persistent blocking Push and Pop boundary
-- requests. Non-blocking PushNB/PopNB polls are not FairAction cases: failed NB
-- polls are ε-steps, and successful NB polls contribute the ordinary committed
-- Push/Pop Σ-event without creating a persistent scheduler-fairness obligation.
-- For synchronization, fairness ranges only over real synchronization calls
-- such as wait(), ac_wait, ac_sync, or SyncChannel calls. It does not range over
-- the auxiliary process-boundary anchors START_P and FINISH_P.
```

inductive FairAction

```
| push (q : DynMsgInstance) (h : q.kind = MsgKind.Push)
| pop (q : DynMsgInstance) (h : q.kind = MsgKind.Pop)
| sync (s : SyncCall)
```

Weak fairness:

Weak fairness is a side-parametric temporal assumption over the FairAction type above. For message operations, FairAction contains only blocking Push and blocking Pop dynamic instances. For synchronization operations, FairAction

contains only real SyncCall instances, not arbitrary SyncAnchor values. Thus START\_P and FINISH\_P may still be used as formal R2/E2 signal anchors when they are realized in the dynamic universe, but they are not scheduler-fairness actions. Non-blocking PushNB/PopNB instances are excluded at the type level; they are handled through the NB execution/witness machinery and through ordinary committed Push/Pop  $\Sigma$ -events when successful.

Weak fairness applies to any valid infinite execution of the selected comparison side. In particular, it is not a post-only premise: the reference-side execution used by Exec\_ref and the post-side execution used by RTL\_B both carry the same WF/B2 obligation whenever that side is being used in a liveness/progress theorem.

For message actions, FairActionEnabled delegates to ChannelEnabledSide and therefore denotes abstract-boundary eligibility for selection, not an already committed transfer. Selection is a separate scheduler/arbitrator fact, and the B2 commit-soundness clause below states that a selected enabled message action, or selected enabled rendezvous pair, commits in that selected cycle.

```
def FairActionCommitsAt
(C : ComparisonInstance)
(side : Side)
( $\tau$  : InfiniteExecutionBucketed C.Sigma)
(a : FairAction)
(n : Cycle) : Prop :=
FairActionSelected C side  $\tau$  a n  $\wedge$ 
SelectedFairActionCommitsInCycle C side  $\tau$  a n
```

```
def WeakFairness
(C : ComparisonInstance)
(side : Side)
( $\tau$  : InfiniteExecutionBucketed C.Sigma) : Prop :=
 $\forall$  a t,
AlwaysFrom  $\tau$  t (fun n => FairActionEnabled C side  $\tau$  a n)  $\rightarrow$ 
EventuallyFrom  $\tau$  t (fun n => FairActionCommitsAt C side  $\tau$  a n)
```

B3 channel-progress predicates should be defined in terms of continuous endpoint-level ready/valid requests, with operation attribution supplied

separately. They should not require that the complementary endpoint's continuous request be witnessed by one persistent dynamic message-operation instance.

A blocking request is persistent from cycle  $t$  through its commit cycle inclusive when the same dynamic operation instance  $q$  keeps its `BoundaryReq` asserted throughout that interval. For Push operations, persistence also includes payload stability while the request is outstanding, including the commit cycle itself:

```
def PersistentRequestFrom
  (C : ComparisonInstance)
  (side : Side)
  ( $\tau$  : InfiniteExecutionBucketed C.Sigma)
  (q : DynMsgInstance)
  (t : Nat) : Prop :=
  Blocking q  $\wedge$ 
   $\forall n$ ,
  t  $\leq$  n  $\rightarrow$ 
   $\neg$  CommitsBefore C side  $\tau$  q n  $\rightarrow$ 
  BoundaryReqAt C side  $\tau$  q n  $\wedge$ 
  (q.kind = MsgKind.Push  $\rightarrow$ 
  PayloadAt C side  $\tau$  q n = PayloadAt C side  $\tau$  q t)
```

`PersistentRequestFrom` remains the right predicate for a single outstanding blocking Push or Pop. It is deliberately not used to collapse a sequence of executed non-blocking polls into one dynamic instance. Repeated PushNB/PopNB polls are distinct source-level dynamic instances, even if the concrete endpoint request signal remains asserted across adjacent cycles.

inductive EndpointDirection

| push  
| pop

```
def DirectionOf
  (q : DynMsgInstance) : EndpointDirection :=
  match q.kind with
  | MsgKind.Push => EndpointDirection.push
  | MsgKind.PushNB => EndpointDirection.push
  | MsgKind.Pop => EndpointDirection.pop
  | MsgKind.PopNB => EndpointDirection.pop
```

```
def IsPushKind
  (k : MsgKind) : Prop :=
  k = MsgKind.Push  $\vee$  k = MsgKind.PushNB
```

```

def IsPopKind
(k : MsgKind) : Prop :=
k = MsgKind.Pop ∨ k = MsgKind.PopNB

def EndpointRequestAssertedAt
(C : ComparisonInstance)
(side : Side)
(τ : InfiniteExecutionBucketed C.Sigma)
(c : Channel)
(dir : EndpointDirection)
(n : Nat) : Prop :=
match dir with
| EndpointDirection.push => EndpointValidAssertedAt C side τ c n
| EndpointDirection.pop  => EndpointReadyAssertedAt C side τ c n

```

```

structure EndpointRequestWitness
(C : ComparisonInstance)
(side : Side)
(τ : InfiniteExecutionBucketed C.Sigma)
(c : Channel)
(dir : EndpointDirection)
(n : Nat) where
endpoint_asserted :
EndpointRequestAssertedAt C side τ c dir n
source_attributed :
∃ q,
q.channel = c ∧
DirectionOf q = dir ∧
BoundaryReqAt C side τ q n
payload_present_if_push :
dir = EndpointDirection.push →
∃ v, EndpointPayloadAt C side τ c n = some v

```

-- Endpoint request continuity is not a forever-after-t predicate. It is scoped  
-- to the interval before the relevant discharge event D. After D occurs, the  
-- endpoint may drop the request or change payload for the next operation.

```

def EndpointRequestHoldsUntilDischarged
(C : ComparisonInstance)
(side : Side)
(τ : InfiniteExecutionBucketed C.Sigma)
(c : Channel)
(dir : EndpointDirection)
(t : Nat)

```

```

(D : Nat → Prop) : Prop :=
  HoldsUntilDischarged τ t
  (fun n =>
    EndpointRequestWitness C side τ c dir n)
  D

```

-- Push payload stability is also scoped only to the outstanding interval. A  
 -- successful transfer may be followed by a deasserted request or a different  
 -- payload for a later operation.

```

def PushPayloadStableUntilDischarged
  (C : ComparisonInstance)
  (side : Side)
  (τ : InfiniteExecutionBucketed C.Sigma)
  (c : Channel)
  (t : Nat)
  (D : Nat → Prop) : Prop :=
  ∃ v,
  HoldsUntilDischarged τ t
  (fun n =>
    EndpointPayloadAt C side τ c n = some v)
  D

```

```

def ContinuousPushRequestUntilDischarged
  (C : ComparisonInstance)
  (side : Side)
  (τ : InfiniteExecutionBucketed C.Sigma)
  (c : Channel)
  (t : Nat)
  (D : Nat → Prop) : Prop :=
  EndpointRequestHoldsUntilDischarged C side τ c EndpointDirection.push t D ∧
  PushPayloadStableUntilDischarged C side τ c t D

```

```

def ContinuousPopRequestUntilDischarged
  (C : ComparisonInstance)
  (side : Side)
  (τ : InfiniteExecutionBucketed C.Sigma)
  (c : Channel)
  (t : Nat)
  (D : Nat → Prop) : Prop :=
  EndpointRequestHoldsUntilDischarged C side τ c EndpointDirection.pop t D

```

```

def BufferedChannel
  (C : ComparisonInstance)
  (c : Channel) : Prop :=

```

$0 < \text{capacity } C \ c$

def RendezvousChannel  
(C : ComparisonInstance)  
(c : Channel) : Prop :=  
capacity C c = 0

def SomePopCommitOn  
(C : ComparisonInstance)  
(side : Side)  
( $\tau$  : InfiniteExecutionBucketed C.Sigma)  
(c : Channel)  
(n : Nat) : Prop :=  
 $\exists$  qPop,  
qPop.channel = c  $\wedge$   
IsPopKind qPop.kind  $\wedge$   
CommitAtCycle C side  $\tau$  qPop n

def SomePushCommitOn  
(C : ComparisonInstance)  
(side : Side)  
( $\tau$  : InfiniteExecutionBucketed C.Sigma)  
(c : Channel)  
(n : Nat) : Prop :=  
 $\exists$  qPush,  
qPush.channel = c  $\wedge$   
IsPushKind qPush.kind  $\wedge$   
CommitAtCycle C side  $\tau$  qPush n

def MatchingTransferCommitOnChannel  
(C : ComparisonInstance)  
(side : Side)  
( $\tau$  : InfiniteExecutionBucketed C.Sigma)  
(c : Channel)  
(n : Nat) : Prop :=  
 $\exists$  qPush qPop,  
qPush.channel = c  $\wedge$   
qPop.channel = c  $\wedge$   
IsPushKind qPush.kind  $\wedge$   
IsPopKind qPop.kind  $\wedge$   
MatchingTransferCommitOn C side  $\tau$  qPush qPop n

def BufferedFullBothEndpointsRequesting  
(C : ComparisonInstance)

```

(side : Side)
( $\tau$  : InfiniteExecutionBucketed C.Sigma)
(c : Channel)
(n : Nat) : Prop :=
BufferedChannel C c  $\wedge$ 
occAt C side  $\tau$  c n = capacity C c  $\wedge$ 
EndpointRequestAssertedAt C side  $\tau$  c EndpointDirection.push n  $\wedge$ 
EndpointRequestAssertedAt C side  $\tau$  c EndpointDirection.pop n

```

```

def BufferedEmptyBothEndpointsRequesting
(C : ComparisonInstance)
(side : Side)
( $\tau$  : InfiniteExecutionBucketed C.Sigma)
(c : Channel)
(n : Nat) : Prop :=
BufferedChannel C c  $\wedge$ 
occAt C side  $\tau$  c n = 0  $\wedge$ 
EndpointRequestAssertedAt C side  $\tau$  c EndpointDirection.pop n  $\wedge$ 
EndpointRequestAssertedAt C side  $\tau$  c EndpointDirection.push n

```

```

def RendezvousBothEndpointsRequesting
(C : ComparisonInstance)
(side : Side)
( $\tau$  : InfiniteExecutionBucketed C.Sigma)
(c : Channel)
(n : Nat) : Prop :=
RendezvousChannel C c  $\wedge$ 
EndpointRequestAssertedAt C side  $\tau$  c EndpointDirection.push n  $\wedge$ 
EndpointRequestAssertedAt C side  $\tau$  c EndpointDirection.pop n

```

-- P holds continuously from t until the first discharge cycle.  
-- This avoids the vacuity that would result from requiring the stalled condition  
-- to remain true forever after the discharge has already occurred.

```

def HoldsUntilDischarged
( $\tau$  : InfiniteExecutionBucketed  $\Sigma$ )
(t : Nat)
(P : Nat  $\rightarrow$  Prop)
(D : Nat  $\rightarrow$  Prop) : Prop :=
 $\forall n, t \leq n \rightarrow (\forall k, t \leq k \rightarrow k < n \rightarrow \neg D k) \rightarrow P n$ 

```

```

def EventuallyDischargedFrom
( $\tau$  : InfiniteExecutionBucketed  $\Sigma$ )
(t : Nat)
(D : Nat  $\rightarrow$  Prop) : Prop :=

```



$\exists n, t \leq n \wedge D n$

Then:

```
def P_push
(C : ComparisonInstance)
(side : Side)
( $\tau$  : InfiniteExecutionBucketed C.Sigma)
(c : Channel) : Prop :=
(BufferedChannel C c  $\rightarrow$ 
 $\forall t$ ,
BufferedFullBothEndpointsRequesting C side  $\tau$  c t  $\rightarrow$ 
ContinuousPushRequestUntilDischarged C side  $\tau$  c t
(fun n =>
SomePopCommitOn C side  $\tau$  c n)  $\rightarrow$ 
ContinuousPopRequestUntilDischarged C side  $\tau$  c t
(fun n =>
SomePopCommitOn C side  $\tau$  c n)  $\rightarrow$ 
EventuallyDischargedFrom  $\tau$  t
(fun n =>
SomePopCommitOn C side  $\tau$  c n))
 $\wedge$ 
(RendezvousChannel C c  $\rightarrow$ 
 $\forall t$ ,
RendezvousBothEndpointsRequesting C side  $\tau$  c t  $\rightarrow$ 
ContinuousPushRequestUntilDischarged C side  $\tau$  c t
(fun n =>
MatchingTransferCommitOnChannel C side  $\tau$  c n)  $\rightarrow$ 
ContinuousPopRequestUntilDischarged C side  $\tau$  c t
(fun n =>
MatchingTransferCommitOnChannel C side  $\tau$  c n)  $\rightarrow$ 
EventuallyDischargedFrom  $\tau$  t
(fun n =>
MatchingTransferCommitOnChannel C side  $\tau$  c n))
```

```
def P_pop
(C : ComparisonInstance)
(side : Side)
( $\tau$  : InfiniteExecutionBucketed C.Sigma)
(c : Channel) : Prop :=
(BufferedChannel C c  $\rightarrow$ 
 $\forall t$ ,
BufferedEmptyBothEndpointsRequesting C side  $\tau$  c t  $\rightarrow$ 
ContinuousPopRequestUntilDischarged C side  $\tau$  c t
(fun n =>
```

SomePushCommitOn C side  $\tau$  c n)  $\rightarrow$   
 ContinuousPushRequestUntilDischarged C side  $\tau$  c t  
 (fun n =>  
 SomePushCommitOn C side  $\tau$  c n)  $\rightarrow$   
 EventuallyDischargedFrom  $\tau$  t  
 (fun n =>  
 SomePushCommitOn C side  $\tau$  c n))  
 $\wedge$   
 (RendezvousChannel C c  $\rightarrow$   
 $\forall t$ ,  
 RendezvousBothEndpointsRequesting C side  $\tau$  c t  $\rightarrow$   
 ContinuousPopRequestUntilDischarged C side  $\tau$  c t  
 (fun n =>  
 MatchingTransferCommitOnChannel C side  $\tau$  c n)  $\rightarrow$   
 ContinuousPushRequestUntilDischarged C side  $\tau$  c t  
 (fun n =>  
 MatchingTransferCommitOnChannel C side  $\tau$  c n)  $\rightarrow$   
 EventuallyDischargedFrom  $\tau$  t  
 (fun n =>  
 MatchingTransferCommitOnChannel C side  $\tau$  c n))

The removed HoldsUntilDischarged premises are not independent B3 assumptions.  
 For bounded FIFOs, the needed full/empty persistence facts are derived from E4: if  $\text{occAt C side } \tau \text{ c t} = \text{capacity C c}$  and no  $\text{Pop\_c}$  commits before n, then the channel remains full at n; dually, if  $\text{occAt C side } \tau \text{ c t} = 0$  and no  $\text{Push\_c}$  commits before n, then the channel remains empty at n. For rendezvous channels, the corresponding “both endpoints requesting until transfer” fact is exactly the conjunction of the two continuous endpoint-request premises together with  $\text{RendezvousChannel C c}$ .

Thus  $P_{\text{push}}$  and  $P_{\text{pop}}$  state only the real B3 progress obligation: once the appropriate blocked condition holds at t and both endpoint requests continue until the matching discharge event, that discharge event must eventually occur. The occupancy-stability facts needed inside proofs should be supplied as E4 derived lemmas rather than duplicated as extra hypotheses in B3.

This matches the document’s conditional B3 reading: B3 forces progress only when both endpoints continuously request the matching transfer until the corresponding discharge event. A producer stalled on a full channel is not automatically a B3 violation if the consumer is not yet requesting the matching pop, and a consumer stalled on an empty channel is not automatically a B3 violation if the producer is not yet requesting the matching push.

The continuous-request premise is endpoint-level and interval-scoped. For blocking Push/Pop it may be witnessed by one persistent dynamic operation

instance using PersistentRequestFrom through the commit cycle. For non-blocking polling loops it may instead be witnessed by a sequence of distinct executed PushNB/PopNB instances on the same endpoint, provided the endpoint request remains continuously asserted until discharge and each cycle's assertion has BoundaryReq attribution to the appropriate executed dynamic instance.

The payload-stability premise for Push is likewise scoped only until discharge. After the discharge event, the producer may deassert valid or present a different payload for a later operation. Thus B3 tracks the ready/valid progress premise without requiring an endpoint request or payload to remain stable forever and without misidentifying multiple NB polls as one persistent q.

B3 modularity (v42). For a single explicit abstract boundary with ordinary E4 buffered or rendezvous semantics, the corresponding B3 discharge is implied by B2 plus E4 once the complementary endpoint request in the B3 premise remains continuously asserted: in the buffered full case,  $\text{occ\_c} = B(c) > 0$  makes the complementary Pop\_c continuously ChannelEnabled, so B2 forces eventual selection and commit of that Pop\_c, freeing space; in the buffered empty case,  $\text{occ\_c} = 0 < B(c)$  makes the complementary Push\_c continuously ChannelEnabled, so B2 forces eventual selection and commit of that Push\_c, providing data; and in the rendezvous case, the two persistent endpoint requests make the rendezvous pair continuously ChannelEnabled, so B2 supplies eventual rendezvous selection and commit. Lean should prove this derivation explicitly:

theorem B3\_ordinary\_boundary\_from\_B2\_E4

(C : ComparisonInstance)

(side : Side)

( $\tau$  : InfiniteExecutionBucketed C.Sigma)

(c : Channel)

(hExec : InfiniteExecutionOfSide C side  $\tau$ )

(hWF : WeakFairness C side  $\tau$ )

(hE4 : E4OccupancyTraceWF C side  $\tau$ )

(hOrd : OrdinaryE4Boundary C c) :

$P\_push\ C\ side\ \tau\ c \wedge P\_pop\ C\ side\ \tau\ c$

Here OrdinaryE4Boundary C c records that c is a single explicit abstract boundary whose ChannelEnabled predicate is exactly the ordinary E4 buffered or rendezvous availability predicate, with no elaborated staged/bundle/join/epoch-gate structure interposed. The proof uses the E4-derived occupancy-stability lemmas (full stays full until a Pop commits; empty stays empty until a Push commits), the continuous complementary endpoint-request premise of  $P\_push / P\_pop$ , and B2 selection-commit soundness for the continuously enabled complementary action or rendezvous pair.

B3 is nevertheless retained as the explicit named B-rule obligation in BRules and DockAssumptions, exactly as in the v42 document: Lemma I.2, Appendix K, and the

channel-validation methodology consume B3 in that modular per-channel form, and it is the form checked per explicit abstract channel in elaborated  $\text{Sys\_B}^{\text{elab}}$  boundary chains, where the boundary in question may not satisfy OrdinaryE4Boundary. Thus B3 should be read as the modular channel-progress assumption used by the proof and checked against channel realizations, not as a claim that the ordinary single-boundary clauses are logically independent of B2 plus E4. For ordinary instances, B3\_ordinary\_boundary\_from\_B2\_E4 lets the B3 field of BRules be discharged inside Lean rather than assumed separately.

The latest scheduling-rules document treats B1 as a standing determinism assumption for the post-HLS RTL\_B side in the main Appendix G/J/K theorem stack. B1 is not a source determinism assumption: the source-side models  $\text{Sys} / \text{Sys\_B} / \text{Sys\_ref}$  may admit multiple  $\varepsilon$ - or latency-different executions having the same  $\Sigma$ -projection, and witness selection / snooping resolve the required source-side correspondence when needed.

For a fixed comparison instance — including the fixed initial state and fixed reactive environment / external stimulus — B1 requires uniqueness of the labeled post-side  $\Sigma$ -trace, not uniqueness of the full internal  $\varepsilon$ -execution. Internal  $\varepsilon$ -steps,  $\varepsilon$ -depths, and other non-observable scheduling details may differ between valid post-side executions.

```
def PostDeterministicSigmaTrace
(C : ComparisonInstance) : Prop :=
 $\forall \tau_1 \tau_2,$ 
InfinitePostExecution C  $\tau_1 \rightarrow$ 
InfinitePostExecution C  $\tau_2 \rightarrow$ 
SigmaTraceOf C .post  $\tau_1 =$ 
SigmaTraceOf C .post  $\tau_2$ 
```

BRules / AssumptionBundle handling of B1. In the concrete Lean development, the B-rule bundle used by Appendix J and Appendix K theorem instances must carry B1 unconditionally for RTL\_B under the fixed comparison instance. Equivalently, JAssumptions C W and KAssumptions C W must expose a field or projection such as:

```
postDeterministicSigmaTrace : PostDeterministicSigmaTrace C
```

This field is the Lean representative of the standing B1 assumption. It is not a separate source-side determinism premise, and it is not conditional only on the phrase ‘when uniqueness is invoked.’ By contrast, B4 and B5 remain separate assumptions or derived structures used when  $\varepsilon$ -normalization,  $\varepsilon$ -quiescent endpoint reasoning, or uniqueness up to finite  $\varepsilon$ -stutter is needed.

Constructive  $\varepsilon$ -normalization and finite-stutter obligations for B4/B5.

B4 and B5 should not be merely opaque Prop fields when used to mechanize Lemma I.1-style arguments. They are assumptions, but they must expose enough structure to prove that internal  $\varepsilon$ -evolution cannot hide an infinite internal tail before the next  $\Sigma$ -action or stable wait.

This Lean strategy distinguishes two implementation-level measurements that are both silent with respect to  $\Sigma$ :

- intra-cycle  $\varepsilon$ -normalization:  $\varepsilon$ -microsteps inside one CycleBucket, used to reach an  $\varepsilon$ -quiescent cutpoint within a cycle; and
- cycle-advancing silent progress: one or more real clock cycles whose buckets contain no  $\Sigma$ -event for the process, used to model HLS-inserted FSM latency, arbitration/bookkeeping latency, failed non-blocking polls, and pipeline advances.

This distinction is only a Lean-internal decomposition. It is not a redefinition of B4. The main-document B4 premise remains a single finite  $\varepsilon$ -stutter bound for each process P: along any consecutive  $\varepsilon$ -only evolution in which P is internally advancing and is not externally stalled, P must reach either a  $\Sigma$ -action or a stable waiting state within a finite bound  $D\_P$ .

Accordingly, ProcessEpsilonBound and IntraCycleNormalizationBound may remain separate Lean-internal implementation parameters, but neither one is the main-document B4 bound  $D\_P$ . FiniteEpsilonDepth must expose a derived unified B4 consequence with one bound per process, and that derived bound is the Lean representative of the main document's  $D\_P$ .

In this strategy, the identification is:

$$D\_P \equiv \text{MainB4Bound } C \ P.$$

ProcessEpsilonBound  $C \ P$  is only an internal bound on cycle-advancing silent progress, and IntraCycleNormalizationBound  $C \ P$  is only an internal bound on intra-cycle  $\varepsilon$ -normalization. They are useful for proof engineering, but together they must imply the single main-document B4 finite-stutter obligation. Any B5 statement of the form “within  $\max\_P \ D\_P$ ” is therefore represented in Lean using  $\max\_P \ \text{MainB4Bound } C \ P$ , not  $\max\_P \ \text{ProcessEpsilonBound } C \ P$ .

```

def ProcessEpsilonBound
(C : ComparisonInstance)
(P : Process) : Nat :=
pipe_depth C P

def IntraCycleNormalizationBound
(C : ComparisonInstance)
(P : Process) : Nat :=
eps_norm_depth C P
-- Main-document B4 bound D_P.
-- MainB4Bound C P is the single D_P-style finite  $\varepsilon$ -stutter bound exposed to
-- the theorem stack. ProcessEpsilonBound and IntraCycleNormalizationBound are
-- Lean-internal components only; they are not separate candidates for D_P. The
-- exact arithmetic expression is not important, provided MainB4Bound dominates
-- the Lean-internal cycle-progress and bucket-normalization bounds strongly
-- enough to prove the unified B4 consequence.
def MainB4Bound
(C : ComparisonInstance)
(P : Process) : Nat :=
(ProcessEpsilonBound C P + 1) *
(2 * IntraCycleNormalizationBound C P + 1)

def IntraCycleEpsilonStep
(C : ComparisonInstance)
(side : Side)
( $\sigma \sigma'$  : State) : Prop :=
InternalStep C side  $\sigma \sigma' \wedge$ 
SigmaEventsBetween C side  $\sigma \sigma' = \emptyset \wedge$ 
 $\sigma'.cycle = \sigma.cycle$ 

def IntraCycleEpsilonSegment
(C : ComparisonInstance)
(side : Side)
( $\sigma_0 \sigma_1$  : State)
(steps : List EpsStep) : Prop :=

```

EpsilonStepsFromTo C side  $\sigma_0$  steps  $\sigma_1 \wedge$   
 $\forall \sigma \sigma',$   
 AdjacentInEpsilonSteps C side  $\sigma_0$  steps  $\sigma_1 \sigma \sigma' \rightarrow$   
 IntraCycleEpsilonStep C side  $\sigma \sigma'$

def StableWaitOrSigmaEnabled  
 (C : ComparisonInstance)  
 (side : Side)  
 ( $\sigma$  : State) : Prop :=  
 StableWaitCutpoint C side  $\sigma \vee$   
 $\exists e, \text{SigmaEnabledAt C side } e \sigma$

-- Event ownership is side-parametric. In simple cases the ownership predicate  
 -- may reduce to a side-agnostic schema-level owner map over C.Sigma, but the  
 -- side parameter is still carried here so bucket membership, execution  
 -- well-formedness, and process attribution are checked in the same side context.

def SigmaCommitsForProcessInBucket  
 (C : ComparisonInstance)  
 (side : Side)  
 (P : Process)  
 ( $\beta$  : CycleBucket C.Sigma) : Prop :=  
 $\exists e,$   
 $e \in \beta.\text{sigmaEvents} \wedge$   
 EventOwnedByProcess C side P e

def SigmaCommitsForProcessAt  
 (C : ComparisonInstance)  
 (side : Side)  
 (P : Process)  
 ( $\tau$  : InfiniteExecutionBucketed C.Sigma)  
 (n : Cycle) : Prop :=  
 SigmaCommitsForProcessInBucket C side P ( $\tau$  n)

def ProcessSilentCycle  
 (C : ComparisonInstance)  
 (side : Side)  
 (P : Process)

```

( $\beta$  : CycleBucket C.Sigma) : Prop :=
WellFormedCycleBucket C side  $\beta$   $\wedge$ 
ProcessActiveInBucket C side P  $\beta$   $\wedge$ 
 $\neg$  SigmaCommitsForProcessInBucket C side P  $\beta$   $\wedge$ 
InternalProgressForProcess C side P  $\beta$ 

```

```

def SilentCycleSegment
(C : ComparisonInstance)
(side : Side)
(P : Process)
( $\tau$  : InfiniteExecutionBucketed C.Sigma)
( $t_0$   $t_1$  : Cycle) : Prop :=
 $t_0 \leq t_1$   $\wedge$ 
 $\forall n$ ,
 $t_0 \leq n \rightarrow$ 
 $n < t_1 \rightarrow$ 
ProcessSilentCycle C side P ( $\tau$  n)

```

```

def StableWaitOrSigmaEnabledAtCycle
(C : ComparisonInstance)
(side : Side)
(P : Process)
( $\tau$  : InfiniteExecutionBucketed C.Sigma)
(n : Cycle) : Prop :=
StableWaitOrSigmaEnabled C side (( $\tau$  n).endState)  $\vee$ 
 $\exists e$ ,
EventOwnedByProcess C side P e  $\wedge$ 
SigmaEnabledAt C side e (( $\tau$  n).endState)

```

```

structure FiniteEpsilonDepth (C : ComparisonInstance) where
-- MainB4Bound is the Lean representative of the main-document D_P. These two
-- fields record that the Lean-internal component bounds are dominated by that
-- single exposed D_P bound.
main_B4_bound_dominates_cycle_bound :
 $\forall P$ ,
ProcessEpsilonBound C P  $\leq$  MainB4Bound C P

```



main\_B4\_bound\_dominates\_normalization\_bound :

$\forall P,$

$\text{IntraCycleNormalizationBound } C \ P \leq \text{MainB4Bound } C \ P$

-- Cycle-level internal consequence used by Lemma I.1-style arguments.

-- This bound is intentionally stated using the Lean-internal

-- `ProcessEpsilonBound`. It is not the main-document  $D_P$ ; it is one component

-- used to prove the unified  $D_P$  bound below.

silent\_cycle\_bound :

$\forall \text{ side } \tau \ P \ t_0 \ t_1,$

$\text{InfiniteExecutionOfSide } C \ \text{side } \tau \rightarrow$

$\text{SilentCycleSegment } C \ \text{side } P \ \tau \ t_0 \ t_1 \rightarrow$

$(\forall n,$

$t_0 \leq n \rightarrow$

$n < t_1 \rightarrow$

$\neg \text{StableWaitOrSigmaEnabledAtCycle } C \ \text{side } P \ \tau \ n) \rightarrow$

$t_1 - t_0 \leq \text{ProcessEpsilonBound } C \ P$

-- Main-document B4 consequence: all consecutive  $\varepsilon$ -only evolution, including

-- any intra-cycle normalization  $\varepsilon$ -steps represented inside `CycleBucket`, is

-- bounded by the single  $D_P$  bound before the process either commits a  $\Sigma$ -action

-- or reaches stable wait. In this Lean strategy, that  $D_P$  bound is

--  $\text{MainB4Bound } C \ P$ .

unified\_epsilon\_stutter\_bound :

$\forall \text{ side } \tau \ P \ s,$

$\text{InfiniteExecutionOfSide } C \ \text{side } \tau \rightarrow$

$\text{ConsecutiveEpsilonOnlySegment } C \ \text{side } P \ \tau \ s \rightarrow$

$\text{ProcessInternallyAdvancingThroughout } C \ \text{side } P \ \tau \ s \rightarrow$

$\text{NoExternalStallThroughout } C \ \text{side } P \ \tau \ s \rightarrow$

$\text{NoSigmaCommitOrStableWaitBeforeEnd } C \ \text{side } P \ \tau \ s \rightarrow$

$\text{EpsilonStepCount } C \ \text{side } P \ \tau \ s \leq \text{MainB4Bound } C \ P$

normalizes\_or\_reaches\_sigma\_or\_wait :

$\forall \text{ side } \tau \ P \ t_0,$

$\text{InfiniteExecutionOfSide } C \ \text{side } \tau \rightarrow$

$\text{ProcessActiveInBucket } C \ \text{side } P \ (\tau \ t_0) \rightarrow$

$\text{BeginsInternalProgressTowardNextObservable } C \ \text{side } P \ \tau \ t_0 \rightarrow$

$\exists n,$

$t_0 \leq n \wedge$   
 $n \leq t_0 + \text{MainB4Bound } C \ P \wedge$   
 $(\text{SigmaCommitsForProcessAt } C \text{ side } P \ \tau \ n \vee$   
 $\text{StableWaitOrSigmaEnabledAtCycle } C \text{ side } P \ \tau \ n)$

-- Counts only  $\varepsilon$ -steps attributed to process P. A global CycleBucket may contain  
 --  $\varepsilon$ -steps from several processes, so the per-process normalization bound must  
 -- be stated over the P-owned slice of the  $\varepsilon$ -prefix /  $\varepsilon$ -suffix, not over the  
 -- entire global list.

def ProcessEpsilonStepsIn  
 (C : ComparisonInstance)  
 (side : Side)  
 (P : Process)  
 (steps : List EpsStep) : Nat :=  
 CountEpsStepsOwnedByProcess C side P steps

-- Equality of the non-clock state components that may not change across an  
 -- idle quiescent bucket. A bucket may still advance the clock, so B5-style  
 -- idle/quiescent reasoning must not assert  $\beta.\text{endState} = \beta.\text{startState}$  unless  
 -- endState is defined before the tick. However, when the bucket has no  
 --  $\Sigma$ -events and no  $\varepsilon$ -prefix/ $\varepsilon$ -suffix steps, the intra-cycle  $\varepsilon$  position must not  
 -- drift silently.

def NonClockStateEq  
 (C : ComparisonInstance)  
 ( $\sigma_0 \ \sigma_1$  : State) : Prop :=  
 $\sigma_0.\text{store} = \sigma_1.\text{store} \wedge$   
 $\sigma_0.\text{epsPos} = \sigma_1.\text{epsPos}$

-- Finite process universe used to form the main-document B5 maximum.  
 -- This is the set of processes in the selected Sys\_ref / RTL\_B comparison.

def ProcessUniverse  
 (C : ComparisonInstance) : Finset Process :=  
 ProcessesOfComparison C

-- Lean representation of the main-document B5 bound  $\max\_P \ D\_P$ .  
 -- Since  $D\_P$  is represented in this strategy by  $\text{MainB4Bound } C \ P$ ,  
 -- SystemQuiescenceBound C is exactly  $\max\_P \ \text{MainB4Bound } C \ P$  over the finite  
 -- process universe of the selected comparison instance.

```

def SystemQuiescenceBound
(C : ComparisonInstance) : Nat :=
(ProcessUniverse C).sup (fun P => MainB4Bound C P)

structure QuiescentNormalization (C : ComparisonInstance) where
bucket_prefix_bounded :
  ∀ side β P,
  WellFormedCycleBucket C side β →
  ProcessEpsilonStepsIn C side P β.epsPrefix ≤ IntraCycleNormalizationBound C P
bucket_suffix_bounded :
  ∀ side β P,
  WellFormedCycleBucket C side β →
  ProcessEpsilonStepsIn C side P β.epsSuffix ≤ IntraCycleNormalizationBound C P
end_quiescent_sound :
  ∀ side β,
  WellFormedCycleBucket C side β →
  β.endQuiescent →
  QuiescentAt C side β.endState

-- Starting a bucket at an ε-quiescent state only says that the incoming state
-- has no pending internal normalization. It does not forbid Σ-events in the
-- bucket, and it does not forbid ε-suffix normalization caused by those Σ-events.
-- The no-further-ε conclusion is valid only for an idle bucket in which no
-- observable action is enabled and no external/environmental change wakes the
-- system.
idle_quiescent_bucket_has_no_epsilon_steps :
  ∀ side β,
  WellFormedCycleBucket C side β →
  QuiescentAt C side β.startState →
  β.sigmaEvents = ∅ →
  NoObservableActionEnabledAtState C side β.startState →
  NoExternalChangeInBucket C side β →
  β.epsPrefix = [] ∧
  β.epsSuffix = [] ∧
  NonClockStateEq C β.startState β.endState

bounded_system_quiescence :

```

$$\begin{aligned}
& \forall \text{ side } \tau t, \\
& \text{InfiniteExecutionOfSide } C \text{ side } \tau \rightarrow \\
& \text{NoObservableActionEnabledAtState } C \text{ side } ((\tau t).\text{startState}) \rightarrow \\
& \text{NoExternalChangeFrom } C \text{ side } \tau t \rightarrow \\
& \exists n, \\
& t \leq n \wedge \\
& n \leq t + \text{SystemQuiescenceBound } C \wedge \\
& \text{QuiescentAt } C \text{ side } ((\tau n).\text{endState}) \wedge \\
& \forall m, \\
& n \leq m \rightarrow \\
& \text{QuiescentAt } C \text{ side } ((\tau m).\text{startState}) \wedge \\
& \text{QuiescentAt } C \text{ side } ((\tau m).\text{endState}) \wedge \\
& ((\text{NoExternalChangeFrom } C \text{ side } \tau m \wedge \\
& \text{NoObservableActionEnabledAtState } C \text{ side } ((\tau m).\text{startState})) \rightarrow \\
& (\tau m).\text{sigmaEvents} = \emptyset \wedge \\
& (\tau m).\text{epsPrefix} = [] \wedge \\
& (\tau m).\text{epsSuffix} = [])
\end{aligned}$$

The `bounded_system_quiescence` field is the Lean-level semantic assumption corresponding to B5. It is triggered from a single cycle whose start state has no enabled observable action. If no external/environmental change wakes the system, this assumption says that the system reaches the  $\varepsilon$ -fixed point within the explicit main-document bound `max_P D_P`.

In this Lean strategy, `D_P` means `MainB4Bound C P`. Therefore the B5 phrase “within `max_P D_P`” is represented by `SystemQuiescenceBound C`, defined as the finite maximum of `MainB4Bound C P` over the process universe of the selected comparison instance. The fact that each `MainB4Bound C P` is below `SystemQuiescenceBound C`, and the least-upper-bound property of `SystemQuiescenceBound C`, are not B5 assumptions; they are Lean lemmas proved from the `Finset.sup` definition of `SystemQuiescenceBound C`. `ProcessEpsilonBound C P` is only an internal cycle-progress component and is not the `D_P` used by B5.

After that point, later buckets remain idle only while the no-external-change condition continues to hold and no observable action becomes enabled again. Thus B5 is a bounded quiescence-closure premise from an idle start state, not a requirement to prove in advance that no observable action is disabled forever from the original cycle.

The field `idle_quiescent_bucket_has_no_epsilon_steps` is deliberately weaker than the invalid statement “`QuiescentAt  $\beta$ .startState` implies  `$\beta$ .epsPrefix = []` and  `$\beta$ .epsSuffix = []` and  `$\beta$ .endState =  $\beta$ .startState`.” A quiescent start state may still be followed by committed  $\Sigma$ -events, and those  $\Sigma$ -events may update state and create  $\epsilon$ -suffix normalization. Also, a bucket may advance the clock, so the idle fixed-point conclusion uses `NonClockStateEq` rather than literal equality of `startState` and `endState`.

`B4_finiteEpsilonDepth` exposes the main-document B4 consequence through `MainB4Bound`: each process has a single finite `D_P`  $\epsilon$ -stutter bound before a  $\Sigma$ -action or stable wait, and in this Lean strategy that `D_P` bound is exactly `MainB4Bound C P`. `ProcessEpsilonBound` and `IntraCycleNormalizationBound` remain useful Lean-internal components, but they are not separate replacements for B4 and must not be used as the `D_P` in B5. The formalization must prove that the internal decomposition implies the single unified B4 finite-stutter obligation. `QuiescentNormalization` may still use `IntraCycleNormalizationBound` for bucket well-formedness and  $\epsilon$ -quiescent cutpoint construction, but those bounds are applied per process by counting only the  $\epsilon$ -steps owned by that process.

Bundle B-rules:

structure `BRules (C : ComparisonInstance)` where

`B1_postDeterminism :`

`PostDeterministicSigmaTrace C`

-- B2 is not post-only. It applies to any valid infinite execution of the  
 -- selected comparison side when that side participates in a progress/liveness  
 -- theorem.

`B2_weakFairness :`

$\forall$  side  $\tau$ ,

`InfiniteExecutionOfSide C side  $\tau \rightarrow$`

`WeakFairness C side  $\tau$`

-- B3 is likewise side-parametric. The `P_push/P_pop` obligations are required  
 -- for the reference-side `Sys_ref` execution as well as the post-side `RTL_B`  
 -- execution when the theorem uses those executions for liveness/deadlock  
 -- reasoning.

`B3_channelProgress :`

$\forall \text{ side } \tau c,$   
 $\text{InfiniteExecutionOfSide } C \text{ side } \tau \rightarrow$   
 $P\_push \ C \text{ side } \tau c \wedge P\_pop \ C \text{ side } \tau c$

$B4\_finiteEpsilonDepth :$   
 $FiniteEpsilonDepth \ C$   
 $B5\_quiescentNormalization :$   
 $QuiescentNormalization \ C$   
 $B6\_idleStutterTimeDivergence :$   
 $IdleStutterTimeDivergence \ C$

B2/B3 remain temporal assumptions usable by K proofs. B4/B5 also remain assumptions, but they are now constructive assumptions: they provide the finite  $\varepsilon$ -normalization facts needed to bound `epsPrefix` / `epsSuffix` lengths in `CycleBucket` and to prove bounded  $\varepsilon$ -chain lemmas such as Lemma I.1.

## 7. Channel semantics and E4

Define channel state over the selected `Sys_ref` boundary.

The abstract occupancy is cycle-boundary, non-fall-through occupancy. It is indexed by the execution side, the infinite bucketed execution, the channel, and the cycle, not by the full intra-cycle state. It must not be a free bookkeeping function disconnected from the committed trace.

`capacity : C.Sys_ref.Channel  $\rightarrow$  Nat`

-- Number of  $\Sigma$ -visible `Push_c` commits on channel `c` strictly before cycle `t`.

`def PushCommitsBefore`

`(C : ComparisonInstance)`

`(side : Side)`

`( $\tau$  : InfiniteExecutionBucketed C.Sigma)`

`(c : C.Sys_ref.Channel)`

`(t : Cycle) : Nat :=`

`CountSigmaCommitsBefore C side  $\tau$  t (fun e => IsPushCommitOnChannel C side e c)`

-- Number of  $\Sigma$ -visible `Pop_c` commits on channel `c` strictly before cycle `t`.

`def PopCommitsBefore`

`(C : ComparisonInstance)`

`(side : Side)`

`( $\tau$  : InfiniteExecutionBucketed C.Sigma)`

```

(c : C.Sys_ref.Channel)
(t : Cycle) : Nat :=
CountSigmaCommitsBefore C side  $\tau$  t (fun e => IsPopCommitOnChannel C side e c)

-- Initial occupancy at the selected comparison boundary. In the usual reset-empty
-- case this is 0. In accounted-initial-occupancy cases it is supplied explicitly
-- by the comparison instance / elaboration package.
initialOcc :
ComparisonInstance → Side → C.Sys_ref.Channel → Nat

-- Commit-history-derived cycle-boundary occupancy.
-- For B(c)=0 rendezvous channels this value is always 0; rendezvous enablement
-- is governed by same-cycle matching endpoint requests, not stored occupancy.
def occAt
(C : ComparisonInstance)
(side : Side)
( $\tau$  : InfiniteExecutionBucketed C.Sigma)
(c : C.Sys_ref.Channel)
(t : Cycle) : Nat :=
if capacity C c = 0 then
0
else
initialOcc C side c

+ PushCommitsBefore C side  $\tau$  c t

- PopCommitsBefore C side  $\tau$  c t

```

For convenience, one may define a state-indexed abbreviation only by projecting the state to its cycle, and only relative to the execution whose state is being inspected:

```

def occAtState
(C : ComparisonInstance)
(side : Side)
( $\tau$  : InfiniteExecutionBucketed C.Sigma)
(c : C.Sys_ref.Channel)
( $\sigma$  : State) : Nat :=
occAt C side  $\tau$  c  $\sigma$ .cycle

```

This abbreviation is not a second occupancy notion. In particular, occupancy is invariant under  $\varepsilon$ -steps and under same-cycle  $\Sigma$ -events because it is defined from commits strictly before the cycle boundary:

```

theorem occ_invariant_within_cycle
(C : ComparisonInstance)
(side : Side)
( $\tau$  : InfiniteExecutionBucketed C.Sigma)
(c : C.Sys_ref.Channel)
( $\sigma_1 \sigma_2$  : State) :
 $\sigma_1.cycle = \sigma_2.cycle \rightarrow$ 
occAtState C side  $\tau$  c  $\sigma_1 = \text{occAtState C side } \tau$  c  $\sigma_2$ 

```

The definition above is the canonical occupancy used by E4, ChannelEnabled, ChannelDisabled, WFG construction, B3 progress reasoning, and drain-closure reasoning. If the concrete Lean development keeps occAt opaque for proof-engineering reasons, it must instead assume the following well-formedness condition and use it everywhere occAt is consumed:

```

structure E4OccupancyTraceWF
(C : ComparisonInstance)
(side : Side)
( $\tau$  : InfiniteExecutionBucketed C.Sigma) where
occ_matches_committed_prefix :
 $\forall$  c t,
occAt C side  $\tau$  c t =
if capacity C c = 0 then
0
else
initialOcc C side c
+ PushCommitsBefore C side  $\tau$  c t
- PopCommitsBefore C side  $\tau$  c t
initial_occupancy_bounded :
 $\forall$  c,
initialOcc C side c  $\leq$  capacity C c
occupancy_bounded :
 $\forall$  c t,
occAt C side  $\tau$  c t  $\leq$  capacity C c
no_underflow :
 $\forall$  c t,
PopCommitsBefore C side  $\tau$  c t  $\leq$ 
initialOcc C side c + PushCommitsBefore C side  $\tau$  c t

```



rendezvous\_has\_no\_stored\_occupancy :

$\forall c\ t,$

capacity C  $c = 0 \rightarrow$

occAt C side  $\tau\ c\ t = 0$

Reset-empty alignment (A10). The main document's formal model is reset-empty: after each reset boundary at which the compared execution begins or restarts, every buffered channel with  $B(c) > 0$  has logical occupancy 0 at the chosen  $\Sigma$  boundary, and the Appendix J/K occupancy equations and the per-channel Push/Pop precedence lemma are interpreted relative to that reset-empty initial channel state. Document-aligned theorem instances therefore set `initialOcc C side c = 0`, and `A10_reset_empty_fifo / ResetEmptyFifoAssumption C` entails exactly that equation at each compared reset boundary. A nonzero accounted `initialOcc` is an extension outside the main document's empty-FIFO model: per v42, preloaded channel contents must instead be modeled explicitly as post-reset producer actions, or supported by a separately stated initialized-state extension whose obligations are then carried by the comparison instance. The `initialOcc` parameter is kept in `occAt` only so that such an explicitly flagged extension does not require re-deriving the occupancy infrastructure; it must not be used to smuggle preloaded contents into a theorem instance that claims the documented A10 assumption.

-- Domain for one-cycle issue-time request reasoning.

-- This domain includes every executed source-level message-operation instance,  
-- including failed PushNB / PopNB polls that produce no committed  $\Sigma$ -event.

def BoundaryReqIssueDomainSide

(C : ComparisonInstance)

(side : Side)

(q : DynMsgInstance)

( $\sigma$  : State) : Prop :=

$\exists P,$

$q \in QExecSide\ C\ side\ P \wedge$

`MessageOperationInstance` q  $\wedge$

`OwnsMessageEndpoint` P q  $\wedge$

`IssuedAtState` C side q  $\sigma$

-- Domain for persistent operation-attributive request/channel reasoning.

-- This is the domain used by `ChannelEnabled`, `ChannelDisabled`, `Front`, `NextFront`,  
-- automatic flush, and WFG construction. Failed NB-only polls are not included  
-- merely because they executed.

def BoundaryReqPersistentDomainSide

(C : ComparisonInstance)

(side : Side)

(q : DynMsgInstance)

```

( $\sigma$  : State) : Prop :=
 $\exists$  P,
q  $\in$  QOmegaSide C side P  $\wedge$ 
MessageOperationInstance q  $\wedge$ 
OwnsMessageEndpoint P q

-- General BoundaryReq domain.
-- A request may be meaningful either because q is issuing in the current cycle
-- or because q is a persistent message-operation frontier instance.
def BoundaryReqDomainSide
(C : ComparisonInstance)
(side : Side)
(q : DynMsgInstance)
( $\sigma$  : State) : Prop :=
BoundaryReqIssueDomainSide C side q  $\sigma$   $\vee$ 
BoundaryReqPersistentDomainSide C side q  $\sigma$ 

```

Executed failed non-blocking polls may belong to Qexec,P for replay and witness purposes while remaining absent from QOmega,P and  $\Omega_P$ . Such polls do acquire the one-cycle BoundaryReq-at-issue obligation when they issue, via BoundaryReqIssueDomainSide. They do not acquire persistent BoundaryReqPersistentDomainSide, ChannelEnabled, ChannelDisabled, Front, NextFront, automatic-flush, or WFG obligations merely because they executed.

```

BoundaryReqSide :
ComparisonInstance  $\rightarrow$  Side  $\rightarrow$  DynMsgInstance  $\rightarrow$  State  $\rightarrow$  Prop

```

```

-- ChannelEnabledSide is the v42 E4 availability/eligibility predicate at the
-- chosen abstract boundary. It is not itself a Commit predicate. For B(c)>0 it
-- is BoundaryReq conjoined with the cycle-boundary FIFO availability predicate;
-- for B(c)=0 it is BoundaryReq conjoined with the same-cycle complementary
-- endpoint request predicate. Hidden grant/arbitration/back-pressure must not
-- be added as an extra conjunct here unless it is represented as explicit
-- Sys_B^elab boundary structure.
ChannelEnabledSide :
ComparisonInstance  $\rightarrow$  Side  $\rightarrow$  DynMsgInstance  $\rightarrow$  State  $\rightarrow$  Prop

```

```

ChannelSelectedSide :
ComparisonInstance  $\rightarrow$  Side  $\rightarrow$  DynMsgInstance  $\rightarrow$  State  $\rightarrow$  Prop

```

```

def ChannelSelectionCommits
(C : ComparisonInstance)
(side : Side) : Prop :=
 $\forall \tau$  q  $\sigma$  t,

```

StateAtCycleInPrefix C side  $\tau$  t  $\sigma \rightarrow$   
 BoundaryReqPersistentDomainSide C side q  $\sigma \rightarrow$   
 ChannelEnabledSide C side q  $\sigma \rightarrow$   
 ChannelSelectedSide C side q  $\sigma \rightarrow$   
 CommitAtCycle C side  $\tau$  q t

```

def ChannelDisabledSide
(C : ComparisonInstance)
(side : Side)
(q : DynMsgInstance)
( $\sigma$  : State) : Prop :=
BoundaryReqPersistentDomainSide C side q  $\sigma$   $\wedge$ 
BoundaryReqSide C side q  $\sigma$   $\wedge$ 
 $\neg$  ChannelEnabledSide C side q  $\sigma$ 

```

For post-side automatic-flush and issue/commit formulas, use the following abbreviations:

```

def BoundaryReqIssueDomain
(C : ComparisonInstance)
(q : DynMsgInstance)
( $\sigma$  : State) : Prop :=
BoundaryReqIssueDomainSide C .post q  $\sigma$ 

```

```

def BoundaryReqPersistentDomain
(C : ComparisonInstance)
(q : DynMsgInstance)
( $\sigma$  : State) : Prop :=
BoundaryReqPersistentDomainSide C .post q  $\sigma$ 

```

```

def BoundaryReqDomain
(C : ComparisonInstance)
(q : DynMsgInstance)
( $\sigma$  : State) : Prop :=
BoundaryReqDomainSide C .post q  $\sigma$ 

```

```

def BoundaryReq
(C : ComparisonInstance)
(q : DynMsgInstance)
( $\sigma$  : State) : Prop :=
BoundaryReqSide C .post q  $\sigma$ 

```

```

def ChannelEnabled
(C : ComparisonInstance)

```

```

(q : DynMsgInstance)
( $\sigma$  : State) : Prop :=
ChannelEnabledSide C .post q  $\sigma$ 

```

```

def ChannelDisabled
(C : ComparisonInstance)
(q : DynMsgInstance)
( $\sigma$  : State) : Prop :=
ChannelDisabledSide C .post q  $\sigma$ 

```

BoundaryReqSide, ChannelEnabledSide, and ChannelDisabledSide are the canonical side-indexed predicates. The shorter BoundaryReq / ChannelEnabled / ChannelDisabled forms are post-side abbreviations used only where the section is already explicitly about RTL\_B / post-side behavior.

BoundaryReqSide is used under BoundaryReqDomainSide. This includes both one-cycle issue requests, via BoundaryReqIssueDomainSide, and persistent frontier/channel requests, via BoundaryReqPersistentDomainSide.

ChannelEnabledSide and ChannelDisabledSide are used only under BoundaryReqPersistentDomainSide or its post-side abbreviation BoundaryReqPersistentDomain. Thus failed non-blocking polls in  $Q\_exec, P \setminus \Omega\_P$  may satisfy the Appendix C one-cycle request-at-issue rule, but they still do not become message-operation frontier items and do not acquire ChannelEnabled, ChannelDisabled, Front, NextFront, automatic-flush, or WFG obligations merely because they executed.

ChannelEnabled is allowed to inspect the settled boundary-request state  $\sigma$  for persistent message-operation frontier instances, but its capacity/occupancy component must use the commit-history-linked value  $occAt\ C\ side\ \tau\ q.channel\ \sigma.cycle$ . Thus E4 enablement is evaluated against the beginning-of-cycle abstract occupancy, with no same-cycle fall-through through  $\varepsilon$ -settling or same-cycle  $\Sigma$ -events.

Capacity cases:

```

inductive CapacityKind
| rendezvous -- B(c) = 0
| buffered -- B(c) > 0

```

Derived E4 lemmas:

```

BufferedPushEnabled
BufferedPopEnabled

```

RendezvousEnabled  
FIFOOrderPreservation  
NoDropNoDuplicate  
OccupancyUpdatePush  
OccupancyUpdatePop

The occupancy-update lemmas are derived from the commit-history definition of  $\text{occAt}$ , or from  $\text{E4OccupancyTraceWF}$  if  $\text{occAt}$  is kept opaque. They are not independent facts about a free primitive. In particular,  $\text{OccupancyUpdatePush}$  and  $\text{OccupancyUpdatePop}$  must prove that moving from cycle  $t$  to  $t+1$  changes  $\text{occAt}$  exactly according to the  $\text{Push}_c$  /  $\text{Pop}_c$  commits in cycle  $t$ , subject to the non-fall-through rule and the single-transfer-per-cycle endpoint constraints.

$\text{E4}$  should be parameterized by  $\text{Sys\_ref}$ , so in elaborated cases it ranges over explicit elaborated channels, not merely outer channels.

## 8. Appendix Q elaboration interface

Appendix Q should be introduced early. In the updated document, elaboration is not a late proof afterthought: in elaborated cases,  $\text{Sys\_ref} = \text{Sys\_B}^{\text{elab}}(\text{E})$ , and all references to  $\text{B}$ ,  $\text{occ}$ ,  $\text{BoundaryReq}$ ,  $\text{ChannelEnabled/Disabled}$ ,  $\text{Front\_P}$ , and  $\text{WFG}$  are interpreted over the explicit channels of the elaborated model.

As in the ordinary  $\text{E4}$  semantics, elaborated-channel occupancy is cycle-boundary, non-fall-through occupancy. Thus  $\text{occ\_E}(d,t)$  and  $\text{A\_E}(c,t)$  are indexed by cycle  $t$ . State-indexed forms such as  $\text{occ\_E}(d,\sigma)$  or  $\text{A\_E}(c,\sigma)$ , when used informally at  $\varepsilon$ -quiescent cutpoints, are only abbreviations for  $\text{occ\_E}(d,\sigma.\text{cycle})$  and  $\text{A\_E}(c,\sigma.\text{cycle})$ .

$\text{Chan\_outer}$  should be a parameter fixed by the chosen outer  $\text{Sys\_B}$ , not a field chosen internally by the elaboration package.

The channel-level elaboration projection follows Appendix Q's single-valued optional form. A proof-internal elaborated channel is either hidden from the outer DUT/testbench boundary or assigned to one outer channel. Separately, a proof-internal token/event may account for multiple outer  $\Sigma_{\text{DUT}}$ -visible channel effects in bundled or joined cases.

structure ElaborationPackage

```

(Sys_B : System)
(Chan_outer : Type)
( $\Sigma$ _DUT : Type) where
Sys_elab : System
Chan_elab : Type
finite_chan_elab : Fintype Chan_elab
ElabToken : Type
B_E : Chan_elab → Nat
occ_E : Chan_elab → Cycle → Nat

-- Channel-level projection/accounting support.
--  $\Pi$ _elab d = some c means that elaborated channel d belongs to the
-- channel-level fiber for outer channel c.  $\Pi$ _elab d = none corresponds to
-- the  $\perp$  case in Appendix Q: d is proof-internal and hidden from the outer
-- DUT/testbench boundary.
Pi_elab : Chan_elab → Option Chan_outer

-- Token/event-level accounting support.
--  $\Pi$ _token_elab tok is the set of outer channels accounted for by one
-- proof-internal token/event. This may contain multiple outer channels for
-- bundled or joined work items. Token-level multiplicity does not make the
-- channel-level  $\Pi$ _elab relation multi-valued.
Pi_token_elab : ElabToken → Set Chan_outer
tokenOfSigma : Sys_elab.Sigma → Option ElabToken

-- Bucketed outer-projection of the elaborated proof-internal execution.
-- This replaces any flat Trace-level projection. The projection may hide
-- proof-internal staged/bundle/join events, but it must preserve the
-- cycle-bucket structure needed for  $\varepsilon$ -quiescent cutpoints, same-cycle atomic
-- units, and temporal B2/B3 reasoning.
piSigmaBucket :
CycleBucket Sys_elab.Sigma →
CycleBucket  $\Sigma$ _DUT

piSigmaFinite :
ExecutionBucketed Sys_elab.Sigma →
ExecutionBucketed  $\Sigma$ _DUT

```

$\text{piSigmaInfinite} :$   
 $\text{InfiniteExecutionBucketed Sys\_elab.Sigma} \rightarrow$   
 $\text{InfiniteExecutionBucketed } \Sigma\_DUT$

$\text{piSigma\_bucket\_wf} : \text{Prop}$   
 $\text{piSigma\_finite\_agrees\_with\_bucket} : \text{Prop}$   
 $\text{piSigma\_infinite\_agrees\_with\_bucket} : \text{Prop}$

$A\_E : \text{Chan\_outer} \rightarrow \text{Cycle} \rightarrow \text{Nat}$   
 $\text{outer\_bucketed\_execution\_preservation} : \text{Prop}$   
 $\text{occupancy\_preservation} : \text{Prop}$   
 $\text{internal\_refinement} : \text{Prop}$   
 $\text{no\_hidden\_cross\_channel\_disablement} : \text{Prop}$   
 $\text{wfg\_visibility} : \text{Prop}$   
 $\text{combinational\_wellformedness} : \text{Prop}$

For each outer channel  $c$ , the channel-level fiber of the elaboration is the single-valued optional fiber

$\Pi\_elab^{-1}(c) = \{ d : \text{Chan\_elab} \mid \text{Pi\_elab } d = \text{some } c \}.$

This channel-level optional map matches Appendix Q's  $\Pi\_elab$  :

$\text{Chan\_elab}(E) \rightarrow \text{Chan\_outer} \cup \{\perp\}$ . The token-level projection  $\text{Pi\_token\_elab}$  is a Lean-side refinement used to make bundled/joined multi-channel accounting explicit when one proof-internal staged/bundle/join token accounts for several outer  $\Sigma\_DUT$ -visible channel effects. It is not an additional named component of Appendix Q.1; the scheduling-rules document folds this accounting into  $A\_E$  and the corresponding Q.1 accounting clauses. The bucketed projection  $\text{piSigmaBucket}$ , together with  $\text{piSigmaFinite}$  and  $\text{piSigmaInfinite}$ , and the occupancy/accounting map  $A\_E$  must be consistent with the channel-level optional projection  $\text{Pi\_elab}$  and with this Lean-side token-level accounting refinement. In particular, elaboration projection is not a flat event-list projection: it preserves the cycle-bucket decomposition needed for  $\varepsilon$ -quiescent cutpoints, same-cycle atomic units, R2C fixed-point buckets, and B2/B3 temporal reasoning, while allowing proof-internal staged/bundle/join events to be hidden from the outer  $\Sigma\_DUT$  view.

For convenience, the elaboration package may also define state-indexed abbreviations by projecting the state to its cycle:

```
occ_E_atState : Chan_elab → State → Nat
occ_E_atState d σ := occ_E d σ.cycle
```

```
A_E_atState : Chan_outer → State → Nat
A_E_atState c σ := A_E c σ.cycle
```

These abbreviations are not separate occupancy or accounting notions. They are invariant under  $\varepsilon$ -steps and same-cycle  $\Sigma$ -events:

```
occ_E_invariant_within_cycle :
∀ d σ1 σ2,
σ1.cycle = σ2.cycle →
occ_E_atState d σ1 = occ_E_atState d σ2
```

```
A_E_invariant_within_cycle :
∀ c σ1 σ2,
σ1.cycle = σ2.cycle →
A_E_atState c σ1 = A_E_atState c σ2
```

Reference choice:

```
inductive ReferenceKind
| ordinarySysB
| elaborated
```

```
structure SysRefChoice where
outerSysB : System
Chan_outer : Type
kind : ReferenceKind
elab? : Option (ElaborationPackage outerSysB Chan_outer Σ_DUT)
Sys_ref : System
sys_ref_matches_kind : Prop
```

For ordinarySysB, sys\_ref\_matches\_kind entails Sys\_ref = outerSysB.

For elaborated, sys\_ref\_matches\_kind entails that there exists an elaboration package E with elab? = some E and Sys\_ref = E.Sys\_elab.

When SysRefChoice is embedded in a ComparisonInstance C, the field



$C.sys\_ref\_consistent$  additionally requires  $C.Sys\_ref = C.refChoice.Sys\_ref$ . Thus  $SysRefChoice$  records the construction/justification of the selected reference system, while  $ComparisonInstance.Sys\_ref$  is the single reference system used by the theorem statements. There is no freedom for these two references to diverge.

This also handles sync-boundary epoch gates: if proof-relevant withholding exists, it must be represented by explicit gate channels in  $Sys\_B^{elab}(E)$ , not by hidden gating on an otherwise E4-enabled outer channel. The earlier review correctly flagged this as a point that must be explicit in the Lean strategy.

**Locality note.** Although  $ElaborationPackage$  is presented as a single structure parameterizing the theorem instance, its components ( $Sys\_elab$ ,  $Chan\_elab$ ,  $\Pi\_elab$ ,  $\Pi\_token,elab$ ,  $\pi\_Σ,E$ ,  $A\_E$ ) are constructed locally per process and per channel boundary. The channel-level projection  $\Pi\_elab$  is single-valued and optional, matching Appendix Q’s map  $Chan\_elab(E) \rightarrow Chan\_outer \cup \{\perp\}$ . The token-level projection  $\Pi\_token,elab$  is the object allowed to map one proof-internal staged/bundle/join token to multiple outer  $Σ\_DUT$ -visible channels. Each per-process elaboration is determined by that process’s local channel structure (number of inputs, sync-boundary discipline, pipelining configuration) and is independent of the elaboration chosen for any other process. The system-level  $ElaborationPackage$  is the disjoint composition of these per-process pieces. Treating  $E$  as a theorem-instance parameter is the right Lean encoding of “the verifier has chosen per-process elaborations and the proof takes them as given”; it is not a claim that the verifier must reason globally to construct  $E$ .

## 9. R-rules, E-rules, and modeling contracts

The R-rule bundle should include only the document’s actual R-rules plus an optional Lean-side array-rule pointer if desired. It should not include a fictional  $R5\_syncBoundaryNoCrossing$ ; sync-boundary crossing is E5. The document explicitly labels “Messages cannot cross their immediate synchronization boundary” as E5.

However, R5 itself is a real R-rule: R5 (Channel capacity semantics; Sys vs.  $Sys\_B$ ). In the Lean strategy, R5 should record the chosen source/reference-model capacity convention: raw Sys uses rendezvous capacity, while  $Sys\_B / Sys\_ref$  uses the selected capacity  $B(c)$ , with  $B(c)=0$  interpreted as rendezvous and  $B(c)>0$

interpreted as cycle-boundary, non-fall-through bounded FIFO. The concrete enablement and occupancy-update obligations remain the E4 channel-capacity semantics; R5 records which source/reference capacity semantics are being used by the selected comparison instance.

The predicates `RendezvousChannel C c` and `BufferedChannel C c` are definitional abbreviations for the zero-capacity and positive-capacity cases. Therefore the substantive R5 obligation is not the tautological case split  $\text{capacity } C \ c = 0 \rightarrow \text{RendezvousChannel } C \ c$  and  $0 < \text{capacity } C \ c \rightarrow \text{BufferedChannel } C \ c$ . The substantive obligation is that the selected source reference model uses the intended capacity convention for the comparison instance: in the raw Sys case this means rendezvous capacity; in `Sys_B / Sys_B^elab` this means the selected `B(c)` capacities, with E4 supplying the corresponding `ChannelEnabled / ChannelDisabled` and occupancy-update semantics.

```
def R2CCombinationalFixedPointRule
(C : ComparisonInstance) : Prop :=
SequentialOnly C ∨
∃ hR2C : R2C_WF C,
R2CEventsObservedOnlyInStutterCanonicalUnits C hR2C ∧
R2CUnitsMatchedByComponentLocalOrdinal C hR2C
```

```
def R5ChannelCapacitySemantics
(C : ComparisonInstance) : Prop :=
SysRefUsesSelectedCapacitySemantics C
```

```
structure RRULES (C : ComparisonInstance) where
R1_syncOrder : Prop
R2_signalAnchorSequential : Prop
R2C_combinationalFixedPoint :
R2CCombinationalFixedPointRule C
R3_syncSemantics : Prop
R4_messageIssueOrder : Prop
R5_channelCapacitySemantics :
R5ChannelCapacitySemantics C
```

`R4_messageIssueOrder` is deliberately named only as an issue-order rule. It records the source-induced no-reverse discipline for message-operation issue. It does not impose a separate commit-order rule; commit timing is constrained

by E4 channel availability, FIFO/rendezvous semantics, backpressure, and the chosen Sys\_ref / RTL\_B comparison instance.

R2 source well-formedness (RULE 1 / RULE 2). In v42, RULE 1 conformance is a source well-formedness obligation for proof-backed flows, not merely a recommended diagnostic. R2\_signalAnchorSequential therefore carries, in addition to the anchoring equations, the following obligations: every observable sequential-process signal Read/Write has a unique real bounding synchronization anchor on each dynamic path on which it occurs (with START\_P usable only when the modeling convention includes it as a real synchronization event in S, and FINISH\_P usable only for an actually terminating process); no observable Read/Write event is assigned clock time  $\infty$ ; and if a blocking message-passing operation separates a signal read or write from its corresponding synchronization anchor, the theorem instance is admissible only if the model is rewritten so the signal IO is adjacent to its intended anchor, or an explicit source-level well-formedness discharge shows that the signal IO has a unique unconditional real anchor on every dynamic path and cannot observe a value inconsistent with that anchor. In the Lean development these facts are part of C.WellFormed together with the read\_anchor\_defined / write\_anchor\_defined / anchor-realized fields of DynamicUniverse; an unresolved tool warning is not by itself a discharge.

E-rules:

structure ERules (C : ComparisonInstance) where

E1\_sync\_order\_preservation : Prop

E2\_signal\_anchor\_preservation : Prop

E3\_message\_order\_preservation : Prop

E4\_channel\_capacity\_semantics : Prop

E5\_sync\_boundary\_preservation : Prop

Modeling contracts live in the single ModelingContracts structure defined in Section 4.2. That structure includes direct-input contracts, array-access mapping rules, and latency-sensitive local-protocol handling:

structure ModelingContracts (C : ComparisonInstance) where

directInputs : DirectInputContract C

arrays : ArrayAccessMappingRules C

latencySensitiveLocalProtocols : LatencySensitiveLocalProtocolContract C

Section 9 should not introduce a second, narrower definition of ModelingContracts. The same contract bundle is used by the R/E-rule theorem instances, Appendix I/J witness

construction, Appendix L snooping/WC, direct-input late-sampling proofs, external-array proofs, and the cadence-sensitive local-protocol isolation obligations.

## 10. Array access mapping

External arrays are not merely internal Store. The document states that external arrays are mapped onto IO operations by an array access mapping layer, and that array accesses before a synchronization call must commit no later than the sync commit, while accesses after the sync must commit strictly after it. It also permits transformations such as caching, merging, splitting, and reordering when allowed by the mapping layer.

Define:

inductive MemoryEvent

| read (array : ArrayId) (addr : Addr) (value : Value)

| write (array : ArrayId) (addr : Addr) (value : Value)

structure ArrayAccessMapping where

sourceArrayAccess : SourceArrayAccess → List CoreIOEvent

commitTime : SourceArrayAccess → Cycle

fixedLatency : SourceArrayAccess → Nat

conflictFreeReorderAllowed :

SourceArrayAccess → SourceArrayAccess → Prop

transformedEquivalent :

SourceArrayAccess → List SourceArrayAccess → Prop

Rules:

Array accesses are constrained by the same dynamic synchronization intervals as the ordinary source-order frontier, but SourceArrayAccess values are not themselves members of prefS(s) or suffS(s). The array-access mapping layer must therefore provide an explicit bridge from source array accesses to the dynamic frontier/synchronization interval structure.

structure ArraySyncIntervalBridge

(C : ComparisonInstance) where

arrayFrontierItemOf :  
SourceArrayAccess  $\rightarrow$  Option FrontierItem

array\_item\_in\_dynamic\_universe :  
 $\forall a\ x,$   
arrayFrontierItemOf a = some x  $\rightarrow$   
 $\exists P\ W,$   
 $x \in (C.witness.dynamicUniverse\ P\ W).\Omega$

array\_before\_sync\_matches\_prefS :  
 $\forall a\ s\ x,$   
arrayFrontierItemOf a = some x  $\rightarrow$   
ArrayAccessBeforeSync C a s  $\leftrightarrow$   
FrontierItemInPrefS C x s

array\_after\_sync\_matches\_suffS :  
 $\forall a\ s\ x,$   
arrayFrontierItemOf a = some x  $\rightarrow$   
ArrayAccessAfterSync C a s  $\leftrightarrow$   
FrontierItemInSuffS C x s

array\_commit\_cycle\_is\_mapping\_commit\_cycle :  
 $\forall a,$   
ArrayCommitCycle C a = ArrayCommitTime a

-- External/shared array accesses are interpreted through the array access  
-- mapping layer, not as ordinary internal Store updates. The mapping layer  
-- supplies the externally relevant memory operation, the addressed storage  
-- location, the commit cycle at which that storage is accessed, and the  
-- declared memory semantics used to compare Sys\_ref and RTL\_B.

inductive ArrayOpKind  
| read  
| write

structure ExternalStorageLocation where  
storage : ExternalArrayStorage  
addr : ArrayAddress

```

def ArrayMappedOpKind
(C : ComparisonInstance)
(a : ArrayAccess) : ArrayOpKind :=
...

```

```

def ArrayMappedLocation
(C : ComparisonInstance)
(a : ArrayAccess) : ExternalStorageLocation :=
...

```

```

def ArrayWriteValue
(C : ComparisonInstance)
(a : ArrayAccess) : Option Value :=
...

```

```

def ArrayReadValue
(C : ComparisonInstance)
(a : ArrayAccess) : Option Value :=
...

```

```

def SameExternalStorageLocation
(C : ComparisonInstance)
(a1 a2 : ArrayAccess) : Prop :=
ArrayMappedLocation C a1 = ArrayMappedLocation C a2

```

```

-- The mapping layer may declare that a transformation splits, merges, caches,
-- or otherwise rewrites source array accesses. Such transformations are legal
-- only when the mapped external memory behavior is equivalent under the
-- declared memory semantics.

```

```

def MappingEquivalentUnderSharedMemorySemantics
(C : ComparisonInstance)
(a : ArrayAccess)
(transformed : ArrayAccess) : Prop :=
MappingEquivalent a transformed ∧
PreservesExternalMemoryBehavior C a transformed

```

```

-- Serialization order for accesses that touch the same externally visible

```

-- storage location and are not proven conflict-free. This is a semantic order  
 -- over mapped memory commits, not a new source-order rule for ordinary core IO.

```
def SharedStorageSerializedBefore
(C : ComparisonInstance)
(a1 a2 : ArrayAccess) : Prop :=
  SameExternalStorageLocation C a1 a2 ∧
  ArrayCommitTime a1 < ArrayCommitTime a2
```

```
def SharedStorageSameCycleCompatible
(C : ComparisonInstance)
(a1 a2 : ArrayAccess) : Prop :=
  SameExternalStorageLocation C a1 a2 ∧
  ArrayCommitTime a1 = ArrayCommitTime a2 ∧
  DeclaredSameCycleMemoryPolicyAllows C a1 a2
```

```
def SharedStorageConflict
(C : ComparisonInstance)
(a1 a2 : ArrayAccess) : Prop :=
  SameExternalStorageLocation C a1 a2 ∧
  ¬ conflictFreeReorderAllowed a1 a2
```

```
def SharedStorageCoherent
(C : ComparisonInstance) : Prop :=
  ∀ aRead aWrite,
  ArrayMappedOpKind C aRead = ArrayOpKind.read →
  ArrayMappedOpKind C aWrite = ArrayOpKind.write →
  SameExternalStorageLocation C aWrite aRead →
  ArrayCommitTime aWrite < ArrayCommitTime aRead →
  NoInterveningWriteToSameLocation C aWrite aRead →
  ArrayReadValue C aRead = ArrayWriteValue C aWrite
```

```
structure ArrayAccessMappingRules (C : ComparisonInstance) where
  syncIntervalBridge :
  ArraySyncIntervalBridge C
  beforeSyncCommitsNoLater :
  ∀ a s,
  ArrayAccessBeforeSync C a s →
  ArrayCommitTime a ≤ SyncCommitTime s
```

afterSyncCommitsAfter :

$\forall a\ s,$

ArrayAccessAfterSync  $C\ a\ s \rightarrow$

SyncCommitTime  $s < \text{ArrayCommitTime } a$

issueConsistentWithFixedLatency :

$\forall a,$

ArrayCommitTime  $a =$

ArrayIssueTime  $a + \text{ArrayFixedLatency } a$

-- Transformed operations must preserve the memory behavior declared by the

-- array access mapping layer, including any externally visible shared-memory

-- effects.

transformedOpsRespectMapping :

$\forall a\ \text{transformed},$

TransformedFrom  $a\ \text{transformed} \rightarrow$

MappingEquivalentUnderSharedMemorySemantics  $C\ a\ \text{transformed}$

-- Reordering of external array/memory accesses is legal only when the mapping

-- layer declares the reordering conflict-free.

conflictFreeReorderingOnly :

$\forall a_1\ a_2,$

Reordered  $a_1\ a_2 \rightarrow$

conflictFreeReorderAllowed  $a_1\ a_2$

-- If two mapped accesses touch the same externally visible storage location and

-- are not covered by a declared conflict-free reordering proof, the mapping

-- layer must provide a serialization relation or an explicit same-cycle memory

-- policy. This is the shared-storage rule required by the main scheduling

-- document for external/shared memories.

sharedStorageConflictsAreSerializedOrPolicyCompatible :

$\forall a_1\ a_2,$

$a_1 \neq a_2 \rightarrow$

SharedStorageConflict  $C\ a_1\ a_2 \rightarrow$

SharedStorageSerializedBefore  $C\ a_1\ a_2\ \vee$

SharedStorageSerializedBefore  $C\ a_2\ a_1\ \vee$

SharedStorageSameCycleCompatible  $C\ a_1\ a_2$

-- Reads observe the value determined by the declared serialization/coherence



-- semantics of the shared external storage.

sharedStorageCoherence :

SharedStorageCoherent C

For early milestones, it is acceptable to assume:

AllExternalArrayAccessesAlreadyMappedToCoreIO C

but the full proof strategy should include array mapping.

## 11. R2C combinational fixed point

ComparisonInstance now has:

combModel? : Option CombComposition

R2C should use that field when combinational processes are in scope. In v42, combinational Read/Write events are not emitted unconditionally in every observable cycle bucket. Instead, each explicitly modeled combinational process or combinational network component  $C\_comb$  has a selected observable slice  $Obs(C\_comb)$ . The composed combinational network still has a unique  $\varepsilon$ -quiescent fixed-point valuation  $\mu\_t$  in each compared cycle, but a  $\Sigma$ -visible R2C fixed-point signal-observation unit is emitted only at canonical observation occurrences: (1) the first cycle in the compared execution in which the component's observable slice is defined and observed; (2) a later cycle in which the vector of fixed-point values on  $Obs(C\_comb)$  differs from the most recent emitted unit for that component; (3) a cycle in which an explicitly modeled observer/sampler requires a fresh observation unit; or (4) every relevant cycle for an explicitly cycle-locked R2C component.

The Lean model should therefore preserve four facts:

1. R2C fixed-point values are read from the composed valuation  $\mu\_t$ .
2. The composed valuation is obtained from explicit local combinational process components, rather than from an unexplained monolithic global object.
3.  $\Sigma$ -visible R2C units are stutter-canonicalized per component unless the component is explicitly cycle-locked.
4. Cross-model matching is by component identity and local R2C occurrence ordinal, not by equality of global cycle bucket number or by co-bucketing with unrelated events.

Define:

```
structure R2CObservableLabel where
proc : Process
sig : Signal
direction : SignalDirection -- read or write
```

```
structure LocalCombNetwork where
proc : Process
combStepLocal : Store → Store → Prop
reads : Finset Signal
writes : Finset Signal
obsSlice : Finset R2CObservableLabel
localDeterministic : Prop
```

```
structure CombNetwork where
combStep : Store → Store → Prop
deterministic : Prop
```

```
structure CombComposition where
components :
Process → Option LocalCombNetwork
```

```
composed :
CombNetwork
```

```
component_processes_are_combinational :
 $\forall P N,$ 
components P = some N →
CombinationalProcess P
```

```
component_owner_matches_process :
 $\forall P N,$ 
components P = some N →
N.proc = P
```

```
observable_slice_well_formed :
 $\forall P N$  label,
components P = some N →
```

label  $\in$  N.obsSlice  $\rightarrow$   
 label.proc = P  $\wedge$  (label.sig  $\in$  N.reads  $\vee$  label.sig  $\in$  N.writes)

local\_step\_embeds\_in\_composed\_step :  
 $\forall P N \sigma \sigma',$   
 components P = some N  $\rightarrow$   
 N.combStepLocal  $\sigma \sigma' \rightarrow$   
 composed.combStep  $\sigma \sigma'$

composed\_step\_generated\_by\_components :  
 $\forall \sigma \sigma',$   
 composed.combStep  $\sigma \sigma' \rightarrow$   
 $\exists P N,$   
 components P = some N  $\wedge$   
 N.combStepLocal  $\sigma \sigma'$

resolved\_driver\_or\_disjoint\_write\_policy :  
 $\forall P_1 P_2 N_1 N_2 \text{sig},$   
 components P<sub>1</sub> = some N<sub>1</sub>  $\rightarrow$   
 components P<sub>2</sub> = some N<sub>2</sub>  $\rightarrow$   
 sig  $\in$  N<sub>1</sub>.writes  $\rightarrow$   
 sig  $\in$  N<sub>2</sub>.writes  $\rightarrow$   
 P<sub>1</sub>  $\neq$  P<sub>2</sub>  $\rightarrow$   
 ResolvedMultiDriverPolicy composed sig

composed\_deterministic\_from\_components :  
 ( $\forall P N$ , components P = some N  $\rightarrow$  N.localDeterministic)  $\rightarrow$   
 composed.deterministic

The scheduling document's R2C rule requires deterministic  $\varepsilon$ -settling to a unique observable fixed point. It does not globally require structural combinational acyclicity for every admissible combinational network. Pure combinational oscillation, ambiguous unresolved multi-driver behavior, intentional combinational-loop behavior outside a unique fixed-point interpretation, or non-convergent  $\varepsilon$ -behavior are excluded through the unique-fixed-point obligation below unless they are represented by explicit state or synchronization elsewhere.

Structural acyclicity may still be imposed as a separate well-formedness

requirement for specific elaboration/coupler constructions in Appendix Q, but it is not a universal requirement of LocalCombNetwork, CombNetwork, or CombComposition itself.

```
def CombQuiescent (N : CombNetwork) (σ : State) : Prop :=
  ¬ ∃ σ', N.combStep σ.store σ'.store
```

```
def CombFixedPointAtCycle
(N : CombNetwork)
(t : Cycle)
(μ : Store) : Prop :=
  ReachesByCombSteps N t μ ∧
  IsFixedPoint N μ
```

```
def ComponentParticipatesInCycle
(C : ComparisonInstance)
(M : CombComposition)
(P : Process)
(t : Cycle) : Prop :=
  ∃ N,
  M.components P = some N ∧
  CombinationalProcessActiveAt C P t
```

```
def R2CUnitVector
(C : ComparisonInstance)
(M : CombComposition)
(P : Process)
(t : Cycle)
(μ : Store) : Vector Value :=
  ValuesOfObservableSlice C M P t μ
```

```
def R2CUnitEmissionCanonical
(C : ComparisonInstance)
(M : CombComposition)
(P : Process)
(t : Cycle) : Prop :=
  FirstDefinedR2CObservation C M P t v
  R2CUnitVectorChangedSincePreviousEmission C M P t v
```

ExplicitObserverRequiresFreshR2CUnit C M P t v  
CycleLockedR2CComponent C M P

Well-formedness:

structure R2C\_WF (C : ComparisonInstance) where  
composition\_present :  
C.combModel?.isSome

uniqueComposedFixedPoint :

$\forall t M,$   
C.combModel? = some M  $\rightarrow$   
 $\exists ! \mu, \text{CombFixedPointAtCycle } M.\text{composed } t \mu$

component\_events\_use\_composed\_fixed\_point :

$\forall P e t \mu M,$   
C.combModel? = some M  $\rightarrow$   
ComponentParticipatesInCycle C M P t  $\rightarrow$   
CombReadWriteEventOfProcess C P e  $\rightarrow$   
e  $\in$  R2CUnitEventsAt C P t  $\rightarrow$   
CombFixedPointAtCycle M.composed t  $\mu \rightarrow$   
EventReadsWritesStore e  $\mu$

all\_comb\_events\_have\_component :

$\forall P e t M,$   
C.combModel? = some M  $\rightarrow$   
CombReadWriteEventOfProcess C P e  $\rightarrow$   
e  $\in$  R2CUnitEventsAt C P t  $\rightarrow$   
 $\exists N,$   
M.components P = some N

units\_emit\_only\_at\_canonical\_occurrences :

$\forall P t M,$   
C.combModel? = some M  $\rightarrow$   
R2CUnitEmitted C P t  $\rightarrow$   
R2CUnitEmissionCanonical C M P t

stutter\_cycles\_have\_no\_sigma\_unit :

$\forall P t M,$

$C.combModel? = \text{some } M \rightarrow$   
 $\neg R2CUnitEmissionCanonical\ C\ M\ P\ t \rightarrow$   
 $\neg R2CUnitEmitted\ C\ P\ t$

$component\_local\_occurrence\_matching :$   
 $\forall P\ k\ u\_ref\ u\_post,$   
 $KthR2CUnit\ C.\ ref\ P\ k\ u\_ref \rightarrow$   
 $KthR2CUnit\ C.\ post\ P\ k\ u\_post \rightarrow$   
 $R2CUnitValuesAgree\ C\ P\ u\_ref\ u\_post$

$multi\_input\_well\_formed :$   
 $\forall P\ M,$   
 $C.combModel? = \text{some } M \rightarrow$   
 $MultInputR2CComponent\ C\ M\ P \rightarrow$   
 $SingleExplicitAnchoringDomain\ C\ M\ P\ V$   
 $AllAdmissibleInterleavingsInduceSameR2CUnitSequence\ C\ M\ P\ V$   
 $R2CComponentHiddenBehindExplicitSynchronizationOrTransactor\ C\ M\ P$

$def\ R2CEventsObservedOnlyInStutterCanonicalUnits$   
 $(C : ComparisonInstance)$   
 $(hR2C : R2C\_WF\ C) : Prop :=$   
 $\forall P\ e\ t,$   
 $CombReadWriteEventOfProcess\ C\ P\ e \rightarrow$   
 $e \in BucketSigmaEvents\ C\ t \rightarrow$   
 $\exists M\ \mu\ unit,$   
 $C.combModel? = \text{some } M \wedge$   
 $ComponentParticipatesInCycle\ C\ M\ P\ t \wedge$   
 $CombFixedPointAtCycle\ M.composed\ t\ \mu \wedge$   
 $R2CUnitEmitted\ C\ P\ t \wedge$   
 $e \in unit.events \wedge$   
 $EventReadsWritesStore\ e\ \mu$

$def\ R2CUnitsMatchedByComponentLocalOrdinal$   
 $(C : ComparisonInstance)$   
 $(hR2C : R2C\_WF\ C) : Prop :=$   
 $\forall P\ k\ u\_ref\ u\_post,$   
 $KthR2CUnit\ C.\ ref\ P\ k\ u\_ref \rightarrow$   
 $KthR2CUnit\ C.\ post\ P\ k\ u\_post \rightarrow$

R2CUnitValuesAgree C P  $\mu_{ref}$   $\mu_{post}$

This formulation keeps the semantic fixed point global where R2C needs it, but makes the proof obligation modular and stutter-canonical: each combinational process contributes a local component and selected observable slice, the composed network used for  $\mu_t$  is justified by the CombComposition fields, and the emitted observable sequence is compared by component-local occurrence ordinal. Thus R2C fixed-point equality remains a system-level statement over  $\mu_t$ , while conformance remains checkable process-by-process plus an explicit composition/driver-resolution and multi-input anchoring/confluence obligation.

If combinational processes are out of scope for an early milestone, use:

SequentialOnly C

and postpone R2C\_WF.

## 12. Issue/Commit well-formedness

Issue and Commit should be relational plus existence/uniqueness, not necessarily computable total functions.

The Issue/Commit layer is side-parametric and uses the same bucketed execution carrier as the rest of the strategy. It must not introduce a separate flat Trace type.

def SystemOfSide

(C : ComparisonInstance) : Side  $\rightarrow$  System

| Side.ref  $\Rightarrow$  C.Sys\_ref

| Side.post  $\Rightarrow$  C.RTL\_B

abbrev SidePrefix

(C : ComparisonInstance)

(side : Side) : Type :=

Prefix (SystemOfSide C side)

abbrev SideInfiniteExecution

(C : ComparisonInstance)

(side : Side) : Type :=

InfiniteExecutionBucketed C.Sigma

IssueAt :

(C : ComparisonInstance) →

(side : Side) →

SidePrefix C side →

DynMsgInstance →

Cycle →

Prop

CommitAt :

(C : ComparisonInstance) →

(side : Side) →

SidePrefix C side →

DynMsgInstance →

Cycle →

Payload →

Prop

def CommitAtCycle

(C : ComparisonInstance)

(side : Side)

(τ : SidePrefix C side)

(q : DynMsgInstance)

(t : Cycle) : Prop :=

∃ v, CommitAt C side τ q t v

Well-formedness:

def BoundaryRequestAtIssue

(C : ComparisonInstance) : Prop :=

∀ side τ q tIssue σ,

IssuedInPrefix C side τ q →

IssueAt C side τ q tIssue →

StateAtCycleInPrefix C side τ tIssue σ →

BoundaryReqIssueDomainSide C side q σ ∧

BoundaryReqSide C side q σ ∧

((q.kind = MsgKind.Push ∨ q.kind = MsgKind.PushNB) →

PayloadAtState C side q σ = PayloadAtIssue C side τ q tIssue)



```

def BlockingBoundaryRequestPersistence
(C : ComparisonInstance) : Prop :=
  ∀ side τ q tIssue t σ,
  IssuedInPrefix C side τ q →
  IssueAt C side τ q tIssue →
  StateAtCycleInPrefix C side τ t σ →
  tIssue ≤ t →
  ¬ CommitsBefore C side τ q t →
  Blocking q →
  BoundaryReqPersistentDomainSide C side q σ ∧
  BoundaryReqSide C side q σ ∧
  (q.kind = MsgKind.Push →
  PayloadAtState C side q σ = PayloadAtIssue C side τ q tIssue)

```

```

def BoundaryRequestSoundness
(C : ComparisonInstance) : Prop :=
  BoundaryRequestAtIssue C ∧
  BlockingBoundaryRequestPersistence C

```

-- SameEndpointDirection means the same concrete ready/valid endpoint, not  
 -- merely the same Push-vs-Pop direction. In particular, two Push operations on  
 -- different channels are not the same endpoint direction and may be outstanding  
 -- concurrently.

```

def SameEndpointDirection
(q0 q1 : DynMsgInstance) : Prop :=
  q0.channel = q1.channel ∧
  DirectionOf q0 = DirectionOf q1

def StableAttribution
(C : ComparisonInstance) : Prop :=
  ∀ side τ q0 q1 t σ t σnext,
  SameEndpointDirection q0 q1 →
  q0 ≠ q1 →
  Blocking q0 →
  StateAtCycleInPrefix C side τ t σ t →
  StateAtCycleInPrefix C side τ (t + 1) σnext →
  BoundaryReqDomainSide C side q0 σ t →
  BoundaryReqSide C side q0 σ t →

```

BoundaryReqDomainSide C side  $q_1$   $\sigma_{\text{next}} \rightarrow$   
 BoundaryReqSide C side  $q_1$   $\sigma_{\text{next}} \rightarrow$   
 $\neg \text{EndpointCommitAt C side } \tau \ q_0.\text{endpoint } t \rightarrow$   
 $\neg \text{EndpointCommitAt C side } \tau \ q_1.\text{endpoint } t \rightarrow$   
 False

Equivalently, the first premise may be written as  $q_0.\text{endpoint} = q_1.\text{endpoint}$  if the concrete Lean development defines endpoint as the pair  $(q.\text{channel}, \text{DirectionOf } q)$ . The important point is that StableAttribution forbids retagging one still-outstanding request to a different operation on the same physical endpoint; it does not forbid independent blocking pushes or pops on different channels.

BoundaryRequestSoundness has two parts. BoundaryRequestAtIssue applies to every issued source-level message-operation instance  $q$ , including PushNB and PopNB instances that fail and therefore produce no committed  $\Sigma$ -event. Such a non-blocking poll must still assert the appropriate endpoint request in its issue cycle. BlockingBoundaryRequestPersistence is the additional non-withdrawal rule for blocking Push/Pop instances: after issue and before commit, the same dynamic operation instance keeps its request asserted, and a Push keeps its payload stable.

StableAttribution applies when the earlier outstanding request  $q_0$  is a blocking Push/Pop. In that case, attribution cannot change to any distinct same-endpoint operation before a commit on that endpoint direction, whether the later candidate  $q_1$  is blocking or non-blocking. For non-blocking PushNB / PopNB calls, request assertion at the issue cycle does not by itself imply persistence beyond that issue cycle, and continuous assertion of the endpoint request signal does not by itself collapse later executed non-blocking call instances into the same dynamic instance  $q$ . Repeated executed NB polls remain distinct source-level dynamic instances unless the witness explicitly identifies the same still-outstanding request instance.

-- Dynamic source interval between synchronization anchors.  
 -- This is witness-specific: only executed dynamic instances in the selected  
 -- witness universe are classified. Untaken branches and unexecuted loop  
 -- iterations have no SyncIntervalld.

structure SyncInterval where

proc : Process

pred? : Option SyncCall

succ? : Option SyncCall

intervalId : Nat

def SyncIntervalId

(C : ComparisonInstance)

(W : WitnessBundle C)

(q : DynMsgInstance) : Option SyncInterval :=

DynamicSourceIntervalOf C W q

structure SyncIntervalWF

(C : ComparisonInstance)

(W : WitnessBundle C) where

interval\_defined\_for\_frontier\_messages :

$\forall P q,$

$q \in (C.witness.dynamicUniverse P W).Q_{\Omega} \rightarrow$

$\exists I, \text{SyncIntervalId } C W q = \text{some } I$

same\_interval\_iff\_same\_anchors :

$\forall q_0 q_1 I_0 I_1,$

$\text{SyncIntervalId } C W q_0 = \text{some } I_0 \rightarrow$

$\text{SyncIntervalId } C W q_1 = \text{some } I_1 \rightarrow$

$I_0 = I_1 \leftrightarrow$

$I_0.proc = I_1.proc \wedge$

$I_0.pred? = I_1.pred? \wedge$

$I_0.succ? = I_1.succ? \wedge$

$I_0.intervalId = I_1.intervalId$

pref\_suff\_consistent :

$\forall s q,$

$q \in \text{prefS } s \vee q \in \text{suffS } s \rightarrow$

$\exists I, \text{SyncIntervalId } C W q = \text{some } I$

def SameSyncInterval

(C : ComparisonInstance)

```

(W : WitnessBundle C)
(q0 q1 : DynMsgInstance) : Prop :=
  ∃ l,
    SyncIntervalld C W q0 = some l ∧
    SyncIntervalld C W q1 = some l

```

```

def SourceIssuesBackToBackAfterCommit
(C : ComparisonInstance)
(W : WitnessBundle C)
(side : Side)
(τ : SidePrefix C side)
(q0 q1 : DynMsgInstance)
(t : Cycle) : Prop :=
  SameEndpointDirection q0 q1 ∧
  SameSyncInterval C W q0 q1 ∧
  CommitAtCycle C side τ q0 t ∧
  NextSourceMessageOpOnEndpoint C q0 q1 ∧
  IssuedInPrefix C side τ q1 ∧
  SourceControlReachesAndIssuesAt C W side τ q1 (t + 1)

```

```

def BackToBackIssueWF
(C : ComparisonInstance)
(W : WitnessBundle C) : Prop :=
  ∀ side τ q0 q1 t,
    SourceIssuesBackToBackAfterCommit C W side τ q0 q1 t →
    EndpointRequestContinuouslyAsserted C side τ q0.endpoint t (t + 1) →
    IssueAt C side τ q1 (t + 1)

```

The SameSyncInterval component of SourceIssuesBackToBackAfterCommit prevents the back-to-back issuance rule from carrying a continuously asserted endpoint request across an explicit wait(), SyncChannel, ac\_sync, or other synchronization boundary. Cross-boundary issue timing is governed instead by SyncBoundaryIssueDiscipline / E5.

The continuous physical endpoint assertion is not what causes q<sub>1</sub> to issue. It only explains why the boundary signal may remain asserted across the q<sub>0</sub>/q<sub>1</sub> transition without a visible deassertion bubble. The source-level cause is the

separate `SourceControlReachesAndIssuesAt` premise:  $q_1$  has actually been reached and issued as the next same-endpoint operation in the same synchronization interval. Without that premise, a held valid/ready signal could be incorrectly used to infer issuance of a source operation that control flow has not yet reached.

```
def SyncBoundaryIssueDiscipline
(C : ComparisonInstance) : Prop :=
  ∀ side τ q0 q1 s tSync t0 t1,
  q0 ∈ prefS s →
  q1 ∈ suffS s →
  SyncCommitAt C side τ s tSync →
  IssueAt C side τ q0 t0 →
  IssueAt C side τ q1 t1 →
  t0 ≤ tSync ∧ tSync < t1
```

```
def NBDynamicInstanceDiscipline
(C : ComparisonInstance) : Prop :=
  ∀ P q0 q1,
  IsNB q0 →
  IsNB q1 →
  ExecutedInProcess C P q0 →
  ExecutedInProcess C P q1 →
  SameSourceCallSiteButDifferentExecution q0 q1 →
  q0 ≠ q1
```

```
structure IssueCommitWF
(C : ComparisonInstance)
(W : WitnessBundle C) where
  issue_exists :
  ∀ side τ q,
  IssuedInPrefix C side τ q →
  ∃ t, IssueAt C side τ q t
  issue_unique :
  ∀ side τ q t1 t2,
  IssueAt C side τ q t1 →
  IssueAt C side τ q t2 →
```

$t_1 = t_2$   
 commit\_exists\_unique :  
 $\forall$  side  $\tau$  q,  
 CommitsInPrefix C side  $\tau$  q  $\rightarrow$   
 $\exists!$  tp : Cycle  $\times$  Payload,  
 CommitAt C side  $\tau$  q tp.1 tp.2  
 boundary\_request\_soundness :  
 BoundaryRequestSoundness C  
 stable\_attribution :  
 StableAttribution C  
 back\_to\_back\_issue :  
 BackToBackIssueWF C W  
 sync\_boundary\_discipline :  
 SyncBoundaryIssueDiscipline C  
 nb\_dynamic\_instance\_discipline :  
 NBDynamicInstanceDiscipline C

Derived timestamp:  
 noncomputable def IssueTime  
 (hWF : IssueCommitWF C W)  
 (h : IssuedInPrefix C side  $\tau$  q) : Cycle :=  
 Classical.choose (hWF.issue\_exists side  $\tau$  q h)

This preserves Appendix C's use of Issue(q) as an auxiliary timestamp while avoiding an unnecessary computability requirement.

### 13. Automatic flush and sync-boundary policy

Define:  
 Pending P  $\sigma$   
 Front P  $\sigma$   
 NextFront P  $\sigma$   
 BlockingMsgNextFront P  $\sigma$   
 Done P  $\sigma$   
 Remain P  $\sigma$

Automatic flush must be stated over the same dynamic universe and source-order objects used by Appendix I/J/K. In particular, it must not be an opaque Prop bundle.

Section 13 reasons primarily about FrontierItem values, while Section 7's BoundaryReqSide / ChannelEnabledSide / ChannelDisabledSide predicates are operation-attributive predicates over DynMsgInstance values. Therefore Section 13 uses the following item-level wrappers. These wrappers are not new semantic predicates; they merely read a message-operation frontier item through its owning dynamic message instance.

```
def MsgOfFrontierItemSide
(C : ComparisonInstance)
(side : Side)
(x : FrontierItem) : Option DynMsgInstance :=
MsgInstanceOfFrontierItem C x.proc x
```

```
def BoundaryReqItemSide
(C : ComparisonInstance)
(side : Side)
(x : FrontierItem)
( $\sigma$  : State) : Prop :=
 $\exists q,$ 
MsgOfFrontierItemSide C side x = some q  $\wedge$ 
BoundaryReqPersistentDomainSide C side q  $\sigma \wedge$ 
BoundaryReqSide C side q  $\sigma$ 
```

```
def ChannelEnabledItemSide
(C : ComparisonInstance)
(side : Side)
(x : FrontierItem)
( $\sigma$  : State) : Prop :=
 $\exists q,$ 
MsgOfFrontierItemSide C side x = some q  $\wedge$ 
ChannelEnabledSide C side q  $\sigma$ 
```

```
def ChannelDisabledItemSide
(C : ComparisonInstance)
```

```

(side : Side)
(x : FrontierItem)
( $\sigma$  : State) : Prop :=
 $\exists q$ ,
MsgOfFrontierItemSide C side x = some q  $\wedge$ 
ChannelDisabledSide C side q  $\sigma$ 

```

```

def PayloadAtFrontierItemState
(C : ComparisonInstance)
(side : Side)
(x : FrontierItem)
( $\sigma$  : State) : Payload :=
PayloadAtState C side (Classical.choose (MsgOfFrontierItemSide_isSome C side x))  $\sigma$ 

```

```

def FrontierItemIssuedAt
(C : ComparisonInstance)
(side : Side)
( $\tau$  : SidePrefix C side)
(x : FrontierItem)
(t : Cycle) : Prop :=
 $\exists q$ ,
MsgOfFrontierItemSide C side x = some q  $\wedge$ 
IssueAt C side  $\tau$  q t

```

```

def FrontierItemCommitsBeforeCycle
(C : ComparisonInstance)
(side : Side)
( $\tau$  : SidePrefix C side)
(x : FrontierItem)
(t : Cycle) : Prop :=
 $\exists q$ ,
MsgOfFrontierItemSide C side x = some q  $\wedge$ 
CommitsBefore C side  $\tau$  q t

```

```

def FrontierItemCommitsBetween
(C : ComparisonInstance)
(side : Side)
( $\tau$  : SidePrefix C side)

```



$(x : \text{FrontierItem})$   
 $(t_0 t : \text{Cycle}) : \text{Prop} :=$   
 $\exists q,$   
 $\text{MsgOfFrontierItemSide } C \text{ side } x = \text{some } q \wedge$   
 $\text{CommitsBetween } C \text{ side } \tau q t_0 t$

$\text{def FrontierItemCommitsAtCycle}$   
 $(C : \text{ComparisonInstance})$   
 $(\text{side} : \text{Side})$   
 $(\tau : \text{SidePrefix } C \text{ side})$   
 $(x : \text{FrontierItem})$   
 $(t : \text{Cycle}) : \text{Prop} :=$   
 $\exists q,$   
 $\text{MsgOfFrontierItemSide } C \text{ side } x = \text{some } q \wedge$   
 $\text{CommitAtCycle } C \text{ side } \tau q t$

$\text{def FrontierBlocked}$   
 $(C : \text{ComparisonInstance})$   
 $(\text{side} : \text{Side})$   
 $(P : \text{Process})$   
 $(\sigma : \text{State}) : \text{Prop} :=$   
 $\text{Front } C \text{ side } P \sigma \neq \emptyset \wedge$   
 $\forall x,$   
 $x \in \text{Front } C \text{ side } P \sigma \rightarrow$   
 $\text{BlockingMessageFrontierItem } x \wedge$   
 $\text{BoundaryReqItemSide } C \text{ side } x \sigma \wedge$   
 $\text{ChannelDisabledItemSide } C \text{ side } x \sigma$

$\text{def LaterThanBlockedFrontier}$   
 $(C : \text{ComparisonInstance})$   
 $(\text{side} : \text{Side})$   
 $(P : \text{Process})$   
 $(\sigma : \text{State})$   
 $(y : \text{FrontierItem}) : \text{Prop} :=$   
 $\exists x,$   
 $x \in \text{Front } C \text{ side } P \sigma \wedge$   
 $\text{SourceOrderLt } C \text{ side } P x y$

```

def FrontierRemainsBlockedOver
(C : ComparisonInstance)
(side : Side)
(P : Process)
( $\tau$  : SidePrefix C side)
( $t_0$   $t$  : Cycle) : Prop :=
 $\forall$   $n$   $\sigma n$ ,
 $t_0 \leq n \rightarrow$ 
 $n \leq t \rightarrow$ 
StateAtCycleInPrefix C side  $\tau$   $n$   $\sigma n \rightarrow$ 
FrontierBlocked C side P  $\sigma n$ 

```

```

def NoNewLaterWorkWhileBlocked
(C : ComparisonInstance)
(side : Side)
(P : Process)
( $\tau$  : SidePrefix C side) : Prop :=
 $\forall$   $t_0$   $t$   $\sigma_0$   $y$ ,
StateAtCycleInPrefix C side  $\tau$   $t_0$   $\sigma_0 \rightarrow$ 
FrontierBlocked C side P  $\sigma_0 \rightarrow$ 
FrontierRemainsBlockedOver C side P  $\tau$   $t_0$   $t \rightarrow$ 
LaterThanBlockedFrontier C side P  $\sigma_0$   $y \rightarrow$ 
 $\neg$  FrontierItemIssuedAt C side  $\tau$   $y$   $t$ 

```

```

def NoWithdrawalWhileBlocked
(C : ComparisonInstance)
(side : Side)
(P : Process)
( $\tau$  : SidePrefix C side) : Prop :=
 $\forall$   $t_0$   $t$   $\sigma_0$   $\sigma t$   $x$ ,
StateAtCycleInPrefix C side  $\tau$   $t_0$   $\sigma_0 \rightarrow$ 
StateAtCycleInPrefix C side  $\tau$   $t$   $\sigma t \rightarrow$ 
 $t_0 \leq t \rightarrow$ 
 $x \in \text{Front C side P } \sigma_0 \rightarrow$ 
BoundaryReqItemSide C side  $x$   $\sigma_0 \rightarrow$ 
 $\neg$  FrontierItemCommitsBetween C side  $\tau$   $x$   $t_0$   $t \rightarrow$ 
BoundaryReqItemSide C side  $x$   $\sigma t \wedge$ 

```

```

(IsPushFrontierItem x →
PayloadAtFrontierItemState C side x σt =
PayloadAtFrontierItemState C side x σ0)

```

```

def FrontierMinimality
(C : ComparisonInstance)
(side : Side)
(P : Process)
(τ : SidePrefix C side) : Prop :=
∀ t σ x y,
StateAtCycleInPrefix C side τ t σ →
x ∈ Front C side P σ →
y ∈ Remain C side P σ →
SourceOrderLt C side P y x →
False

```

The next predicates separate two obligations that should not be conflated.

First, there is the  $K\varepsilon$  / A11-style WFG-inertness obligation. It applies only to intervals that contain no observable  $\Sigma$ -commit. It says that pure internal /  $\varepsilon$ -only movement while a process remains blocked at the same frontier cannot create a new wait-for cause.

Second, there is the active drain-only automatic-flush discipline. It permits already-issued downstream work to commit while the process is blocked at a minimal frontier. Such drain commits may be observable  $\Sigma$ -events, so they must not be excluded by an `InternalOrEpsilonOnlyBetween` premise. The discipline is instead that no new later source work is issued past the blocked frontier, already asserted requests are not withdrawn before commit, and any observable commit owned by the blocked process during the blocked interval is attributable to an already-issued drainable downstream request.

```

def WFGInertForProcess
(C : ComparisonInstance)
(side : Side)
(P : Process)
(σ0 σt : State) : Prop :=

```

$\forall Q,$   
 $\text{WFGE} C \text{ side } \sigma_0 P Q \Leftrightarrow \text{WFGE} C \text{ side } \sigma t P Q$

def NoNewBlockingFrontierCause  
 (C : ComparisonInstance)  
 (side : Side)  
 (P : Process)  
 ( $\sigma_0 \sigma t$  : State) : Prop :=  
 $\forall x,$   
 $x \in \text{Front } C \text{ side } P \sigma t \rightarrow$   
 BlockingMessageFrontierItem  $x \rightarrow$   
 ChannelDisabledItemSide C side  $x \sigma t \rightarrow$   
 $\exists y,$   
 $y \in \text{Front } C \text{ side } P \sigma_0 \wedge$   
 BlockingMessageFrontierItem  $y \wedge$   
 SameBlockingCause C side  $P y x$

def WFGInertUnderInternalOnlyBetween  
 (C : ComparisonInstance)  
 (side : Side)  
 (P : Process)  
 ( $\tau$  : SidePrefix C side) : Prop :=  
 $\forall t_0 t \sigma_0 \sigma t,$   
 StateAtCycleInPrefix C side  $\tau t_0 \sigma_0 \rightarrow$   
 StateAtCycleInPrefix C side  $\tau t \sigma t \rightarrow$   
 $t_0 \leq t \rightarrow$   
 FrontierBlocked C side  $P \sigma_0 \rightarrow$   
 FrontierRemainsBlockedOver C side  $P \tau t_0 t \rightarrow$   
 InternalOrEpsilonOnlyBetween C side  $\tau t_0 t \rightarrow$   
 WFGInertForProcess C side  $P \sigma_0 \sigma t \wedge$   
 NoNewBlockingFrontierCause C side  $P \sigma_0 \sigma t$

def DraggableDownstreamRequest  
 (C : ComparisonInstance)  
 (side : Side)  
 (P : Process)  
 ( $\sigma_0$  : State)  
 ( $x$  : FrontierItem) : Prop :=

AlreadyIssuedBeforeBlockedCutpoint C side P  $\sigma_0 \times \Lambda$   
 DownstreamOfBlockedFrontier C side P  $\sigma_0 \times \Lambda$   
 BlockingMessageFrontierItem  $x \times \Lambda$   
 BoundaryReqItemSide C side  $x \times \sigma_0$

def DrainOnlyCommitDiscipline  
 (C : ComparisonInstance)  
 (side : Side)  
 (P : Process)  
 ( $\tau$  : SidePrefix C side) : Prop :=  
 $\forall t_0 \ t \ \sigma_0 \ e,$   
 StateAtCycleInPrefix C side  $\tau \ t_0 \ \sigma_0 \rightarrow$   
 $t_0 \leq t \rightarrow$   
 FrontierBlocked C side P  $\sigma_0 \rightarrow$   
 FrontierRemainsBlockedOver C side P  $\tau \ t_0 \ t \rightarrow$   
 $e \in \text{SigmaEventsAtCycleInPrefix C side } \tau \ t \rightarrow$   
 EventOwnedByProcess C side P  $e \rightarrow$   
 ObservableCommitEvent  $e \rightarrow$   
 $\exists x,$   
 DrainableDownstreamRequest C side P  $\sigma_0 \times \Lambda$   
 EventCommitsFrontierItem C side  $x \ e \ \Lambda$   
 $\neg \text{FrontierItemCommitsBeforeCycle C side } \tau \ x \ t_0 \ \Lambda$   
 FrontierItemCommitsAtCycle C side  $\tau \ x \ t$

def DrainOnlyDownstreamProgress  
 (C : ComparisonInstance)  
 (side : Side)  
 (P : Process)  
 ( $\tau$  : SidePrefix C side) : Prop :=  
 WFGInertUnderInternalOnlyBetween C side P  $\tau \ \Lambda$   
 DrainOnlyCommitDiscipline C side P  $\tau$

structure AutomaticFlush  
 (C : ComparisonInstance)  
 (side : Side)  
 (P : Process)  
 ( $\tau$  : SidePrefix C side) where  
 block\_point :

$\forall t \sigma,$

StateAtCycleInPrefix C side  $\tau t \sigma \rightarrow$

FrontierBlocked C side P  $\sigma \rightarrow$

FlushBlockedPoint C side P  $\tau t \sigma$

no\_new\_later\_work :

NoNewLaterWorkWhileBlocked C side P  $\tau$

no\_withdrawal :

NoWithdrawalWhileBlocked C side P  $\tau$

frontier\_minimality :

FrontierMinimality C side P  $\tau$

drain\_only\_downstream\_progress :

DrainOnlyDownstreamProgress C side P  $\tau$

sync\_boundary\_policy :

SyncBoundaryPolicy C side P  $\tau$

The sync-boundary policy must explicitly split the ordinary and elaborated cases.

In the elaborated case, it must also state the content of the withholding rule:

producer-facing availability for capacity belonging only to suff\_S(s) is unavailable while the synchronization call s is pending.

def SyncPendingAt

(C : ComparisonInstance)

(side : Side)

(P : Process)

( $\tau$  : SidePrefix C side)

(s : SyncCall)

(t : Cycle) : Prop :=

SyncIssuedButNotCommitted C side P  $\tau s t$

def FutureEpochCapacityFor

(C : ComparisonInstance)

(side : Side)

(P : Process)

(s : SyncCall)  
 (gateChannel : Channel)  
 (q : DynMsgInstance) : Prop :=  
 q ∈ suffS\_msg C side P s ∧  
 MessageOperationInstance q ∧  
 TransferWouldUseCapacityBehindGate C side P s gateChannel q

def SyncBoundaryGateCorrect  
 (C : ComparisonInstance)  
 (side : Side)  
 (P : Process)  
 (τ : SidePrefix C side)  
 (E : ElaborationPackage C.refChoice.outerSysB C.refChoice.Chan\_outer C.Sigma\_DUT)  
 (gateChannel : Channel)  
 (s : SyncCall) : Prop :=  
 GateChannelOf E gateChannel ∧  
 ∀ t σ q,  
 StateAtCycleInPrefix C side τ t σ →  
 SyncPendingAt C side P τ s t →  
 FutureEpochCapacityFor C side P s gateChannel q →  
 BoundaryReqPersistentDomainSide C side q σ →  
 BoundaryReqSide C side q σ →  
 ChannelDisabledSide C side q σ  
 Here suffS\_msg C side P s is the message-operation-instance projection of  
 suffS(s). The sync-boundary gate rule is stated over DynMsgInstance values  
 because BoundaryReqSide, ChannelEnabledSide, and ChannelDisabledSide are  
 obligations of attributed message-operation instances, not of arbitrary  
 FrontierItem records.

inductive SyncBoundaryPolicy  
 (C : ComparisonInstance)  
 (side : Side)  
 (P : Process)  
 (τ : SidePrefix C side)  
 | no\_withholding\_in\_ordinary\_SysB :  
 C.refChoice.kind = ReferenceKind.ordinarySysB →  
 NoProofRelevantSyncBoundaryWithholding C side P τ →  
 SyncBoundaryPolicy C side P τ

```

| explicit_gate_in_elab :
 $\forall E \text{ gateChannel } s,$ 
 $C.\text{refChoice.kind} = \text{ReferenceKind.elaborated} \rightarrow$ 
 $\text{SyncBoundaryGateCorrect } C \text{ side } P \tau E \text{ gateChannel } s \rightarrow$ 
 $\text{SyncBoundaryPolicy } C \text{ side } P \tau$ 

```

In the ordinary Sys\_B case, the sync-boundary clause does not add hidden gating. It asserts that no proof-relevant sync-boundary withholding is present. If such withholding is needed, the theorem instance must use  $\text{Sys\_B}^{\text{elab}(E)}$ , where the withholding appears as ordinary ChannelDisabled behavior at an explicit sync-boundary / epoch-gate abstract channel.

## 14. Witness selector and Appendix L

The uploaded document's witness selector is a partial strategy over  $\varepsilon$ -quiescent source cutpoints and  $\varepsilon$ -modeled choices whose outcomes may affect later  $\Sigma$ -visible behavior; it must be causally available from the matched RTL prefix and global  $\Sigma$ -causal history.

Avoid circularity by defining it over the partial construction so far:

```

structure WitnessSelector (C : ComparisonInstance) where
  choose :
    ( $\tau_{\text{post\_prefix}} : \text{Prefix } C.\text{RTL\_B}$ )  $\rightarrow$ 
    ( $\tau_{\text{ref\_so\_far}} : \text{Prefix } C.\text{Sys\_ref}$ )  $\rightarrow$ 
    ( $\text{cp} : \text{ChoicePoint } \tau_{\text{ref\_so\_far}}$ )  $\rightarrow$ 
    Option ChoiceOutcome
  causal :
     $\forall \tau_p \tau_r \text{ cp out},$ 
    choose  $\tau_p \tau_r \text{ cp} = \text{some out} \rightarrow$ 
    DependsOnlyOnAvailablePrefix  $\tau_p \tau_r \text{ cp out}$ 
  consistent_with_post_prefix :
     $\forall \tau_p \tau_r \text{ cp out},$ 
    choose  $\tau_p \tau_r \text{ cp} = \text{some out} \rightarrow$ 
    ConsistentWithRTLPrefix  $\tau_p \tau_r \text{ cp out}$ 

```

Non-blocking witness mapping.

Failed PushNB / PopNB polls are executed dynamic instances in  $Q_{\text{exec}}$  but produce no  $\Sigma$ -event. Therefore the Post  $\rightarrow$  Ref construction cannot reconstruct them from the  $\Sigma$ -trace alone. The witness bundle must carry an explicit partial dynamic-instance correspondence for such  $\varepsilon$ -modeled NB instances.



structure NBInstanceMap (C : ComparisonInstance) where

$\mu_Q$  :

DynMsgInstance  $\rightarrow$  Option DynMsgInstance

maps\_only\_executed\_nb :

$\forall q\_post\ q\_ref,$

$\mu_Q\ q\_post = \text{some } q\_ref \rightarrow$

ExecutedInPost C  $q\_post \wedge$

ExecutedInRef C  $q\_ref \wedge$

IsNB  $q\_post \wedge$  IsNB  $q\_ref$

preserves\_owner\_and\_interface :

$\forall q\_post\ q\_ref,$

$\mu_Q\ q\_post = \text{some } q\_ref \rightarrow$

$q\_post.\text{proc} = q\_ref.\text{proc} \wedge$

$q\_post.\text{channel} = q\_ref.\text{channel} \wedge$

$q\_post.\text{kind} = q\_ref.\text{kind}$

preserves\_nb\_outcome :

$\forall q\_post\ q\_ref,$

$\mu_Q\ q\_post = \text{some } q\_ref \rightarrow$

NBOutcomePost C  $q\_post = \text{NBOutcomeRef C } q\_ref$

covers\_non\_cadence\_control\_relevant\_failed\_polls :

$\forall q\_post,$

ExecutedInPost C  $q\_post \rightarrow$

IsFailedNB  $q\_post \rightarrow$

AffectsLaterSigmaOrFrontier C  $q\_post \rightarrow$

$\neg \text{CadenceSensitiveNBInstance C } q\_post \rightarrow$

$\exists q\_ref, \mu_Q\ q\_post = \text{some } q\_ref$

cadence\_sensitive\_failed\_polls\_are\_not\_unhandled\_core :

$\forall q\_post,$

ExecutedInPost C  $q\_post \rightarrow$

IsFailedNB  $q\_post \rightarrow$

CadenceSensitiveNBInstance C  $q\_post \rightarrow$

HandledByLatencySensitiveLocalProtocolContract C  $q\_post$

source\_order\_position\_preserved :

$\forall q\_post\ q\_ref,$

$\mu\_Q\ q\_post = \text{some } q\_ref \rightarrow$

$\text{SourceOrderKeyOf } C\ q\_post = \text{SourceOrderKeyOf } C\ q\_ref$

respects\_source\_order\_slice :

$\forall P\ q_{1\_post}\ q_{2\_post}\ q_{1\_ref}\ q_{2\_ref},$

$\mu\_Q\ q_{1\_post} = \text{some } q_{1\_ref} \rightarrow$

$\mu\_Q\ q_{2\_post} = \text{some } q_{2\_ref} \rightarrow$

$\text{SourceOrderLe } C\ P\ (\text{SourceOrderKeyOf } C\ q_{1\_post})\ (\text{SourceOrderKeyOf } C\ q_{2\_post}) \leftrightarrow$

$\text{SourceOrderLe } C\ P\ (\text{SourceOrderKeyOf } C\ q_{1\_ref})\ (\text{SourceOrderKeyOf } C\ q_{2\_ref})$

structure WitnessBundle (C : ComparisonInstance) where

selector : WitnessSelector C

nbMap : NBInstanceMap C

Witness compatibility:

def WC

(C : ComparisonInstance)

(W : WitnessBundle C) : Prop :=

( $\forall$  constructionPrefix,

WitnessChoicesRespectMatchedPrefix C W.selector constructionPrefix)  $\wedge$

WitnessNBMapRespectsMatchedPrefix C W.nbMap  $\wedge$

WitnessChoicesAndNBMapAgree C W.selector W.nbMap  $\wedge$

LatencySensitiveLocalProtocolWC C W C.contracts.latencySensitiveLocalProtocols

Thus WC is the Lean-side packaging of witness selection, the  $\mu\_Q$  matching obligation, and the modeling-contract obligation that cadence-sensitive failed-poll behavior is not left as an unconstrained hidden timing dependence inside the ordinary source core. When E3/E5 reasoning or the Post  $\rightarrow$  Ref construction refers to failed  $\varepsilon$ -modeled non-blocking call instances in  $Q\_exec, P$ , the required  $\mu\_Q$  correspondence is supplied by  $W.nbMap$  as part of WC for non-cadence-sensitive or already boundary-selected cases. If a failed-poll sequence is cadence-sensitive because externally visible behavior depends on the number or cadence of failures, then WC is applied at the chosen observable boundary of the isolated/snoop-aligned local protocol block, not to the raw internal failed-poll cadence.

Snooping assumptions:

structure SnoopingWrapperAssumptions

(C : ComparisonInstance)

$(W : \text{WitnessBundle } C)$  where  
 $\text{covers\_eps\_choices\_affecting\_boundary} :$   
 $\forall P \text{ choice},$   
 $\text{EpsChoiceAffectsLaterSigmaOrFrontier } C \ P \text{ choice} \rightarrow$   
 $\text{ChoiceCoveredBySnoopingOrWitness } C \ W \ P \text{ choice}$   
  
 $\text{does\_not\_relabel\_boundary\_sigma} :$   
 $\forall \tau,$   
 $\text{SnoopedExecution } C \ W \ \tau \rightarrow$   
 $\text{BoundarySigmaTracePreserved } C \ W \ \tau$   
  
 $\text{selected\_choices\_prefix\_causal} :$   
 $\forall \text{ constructionPrefix},$   
 $\text{WitnessChoicesRespectMatchedPrefix } C \ W.\text{selector constructionPrefix}$   
  
 $\text{failed\_polls\_remain\_epsilon} :$   
 $\forall q,$   
 $\text{IsFailedNB } q \rightarrow$   
 $\neg \exists e, \text{FailedPollContributesCommittedSigmaEvent } C \ q \ e$   
  
 $\text{cadence\_sensitive\_sites\_isolated\_or\_snoop\_aligned} :$   
 $\forall P \text{ loc},$   
 $\text{CadenceSensitiveNBPollingSite } C \ P \text{ loc} \rightarrow$   
 $\text{LocalProtocolSiteHandledByContract } C \ W \ C.\text{contracts.latencySensitiveLocalProtocols } P \text{ loc}$

Lemma L.2:

$\text{theorem L2\_snooping\_establishes\_WC}$   
 $(C : \text{ComparisonInstance})$   
 $(W : \text{WitnessBundle } C)$   
 $(h\text{Snoop} : \text{SnoopingWrapperAssumptions } C \ W) :$   
 $\text{WC } C \ W$

This theorem does not turn failed PushNB/PopNB polls into committed  $\Sigma$ -events. Failed polls remain  $\varepsilon$ -steps. For ordinary NB-CF sites, the identity/arbitrary witness suffices. For non-NB-CF but non-cadence-sensitive sites, the witness selects the  $\varepsilon$ -modeled outcomes needed for the matched prefix. For cadence-sensitive retry/timeout/poll-arbiter sites, the theorem may be used only after the local protocol block has been isolated, snoop-aligned, or given an explicit cycle-accurate model, and the comparison is made at that block's chosen observable boundary.

IdentityOrArbitraryNBWitnessBundle C is the witness bundle used in the NB-CF case. Its selector uses identity choices for  $\varepsilon$ -modeled non-blocking outcomes that are fixed by the matched prefix, and arbitrary choices for non-blocking outcomes that NB-CF proves cannot affect later  $\Sigma$ -visible control flow or NextFront behavior. Its nbMap maps matched executed NB instances to the corresponding same dynamic source-call instance when such a match is relevant, and is arbitrary on failed NB instances that NB-CF proves irrelevant. It is not a discharge for cadence-sensitive retry/timeout/poll-arbiter sites unless those sites have first been shown not to affect later  $\Sigma$ -visible behavior, or have been handled by the latency-sensitive local-protocol contract.

```
def IdentityOrArbitraryNBWitnessBundle
(C : ComparisonInstance) : WitnessBundle C :=
{
  selector := IdentityOrArbitraryNBSelector C,
  nbMap := IdentityOrArbitraryNBMap C
}
```

Corollary L.3:

The document-level WC predicate is global for the selected comparison instance and now includes the latency-sensitive local-protocol contract. Therefore a theorem returning global WC must not be stated from a single-process NB-CF premise alone.

Use a global theorem when proving WC for the full comparison instance:

```
theorem L3_nb_cf_identity_witness
(C : ComparisonInstance)
(hNBCF :  $\forall P, \text{NBCF } C P$ )
(hLocalProtocols :
  LatencySensitiveLocalProtocolWC
  C
  (IdentityOrArbitraryNBWitnessBundle C)
  C.contracts.latencySensitiveLocalProtocols) :
WC C (IdentityOrArbitraryNBWitnessBundle C)
```

If a per-process theorem is useful, state it as a local witness fact rather than as global WC:

```
theorem L3_nb_cf_identity_witness_local
(C : ComparisonInstance)
```

(P : Process)

(hNBCF : NBCF C P) :

LocalWCForProcess C (IdentityOrArbitraryNBWitnessBundle C) P

The global theorem is then obtained by combining the local theorem over all processes with the latency-sensitive local-protocol contract:

theorem L3\_nb\_cf\_identity\_witness\_from\_local

(C : ComparisonInstance)

(hLocal :  $\forall P$ , LocalWCForProcess C (IdentityOrArbitraryNBWitnessBundle C) P)

(hLocalProtocols :

LatencySensitiveLocalProtocolWC

C

(IdentityOrArbitraryNBWitnessBundle C)

C.contracts.latencySensitiveLocalProtocols) :

WC C (IdentityOrArbitraryNBWitnessBundle C)

Corollary L.4 (cadence-sensitive polling handled by isolation). The scheduling-rules document names this fact Corollary L.4; the strategy already enforces its content through the LatencySensitiveLocalProtocolContract and the WC definition, but the Lean development should also expose it under its document name:

theorem L4\_cadence\_sensitive\_isolated\_or\_out\_of\_scope

(C : ComparisonInstance)

(P : Process)

(loc : SourceLoc)

(hCadence : CadenceSensitiveNBPollingSite C P loc) :

HandledByLatencySensitiveLocalProtocolContractAt

C C.contracts.latencySensitiveLocalProtocols P loc  $\vee$

OutsideOrdinaryAppendixIJSrcCoreAt C P loc

That is: if NB-CF fails because the source observes the number or cadence of failed PushNB/PopNB polls — retry counters, timeout counters, or arbitration policies whose later  $\Sigma$ -visible behavior depends on how long a channel remained empty or full — then the site is latency-sensitive at that polling site, and it must be isolated, snoop-aligned, or given an explicit cycle-accurate local protocol model under Appendix D / Appendix L before the ordinary Appendix I/J source-core proof is applied. The left disjunct is exactly the contract obligation `cadence_sensitive_nb_sites_handled`; the right disjunct records that an unhandled site removes the process from the scope of the ordinary trace-compatibility theorem at the chosen observable boundary. Lemma L.2 may then be used only at the chosen observable boundary of the handled block, never against the raw internal failed-poll cadence.

Appendix L Note 1:  $\Sigma$ \_DUT-safety-inert local timing choices.

The preceding L.2/L.3 facts are WC-producing facts for the ordinary proof-internal C.Sigma theorem. Appendix L Note 1 is different: it is a projected-safety principle for the outer DUT/testbench alphabet C.Sigma\_DUT. It says that an internal latency-sensitive interface need not be snooped for the  $\Sigma$ \_DUT-safety fragment when its  $\varepsilon$ -modeled timing choices cannot change which  $\Sigma$ \_DUT-visible observable units occur, cannot change their projected values, and can only permute projected units that are causally independent at the DUT boundary.

This is not a replacement for WC in the full C.Sigma theorem, and it is not an Appendix K deadlock/liveness exemption. A  $\Sigma$ \_DUT-safety-inert interface that is also used in Appendix K reasoning must still satisfy the WFG-inert /  $K_\varepsilon$  obligations, such as A11\_Kepsilon\_WFGInert.

-- A proof-internal observable unit whose projection is visible at the chosen  
-- DUT/testbench boundary.

SigmaDUTVisibleUnit :

ComparisonInstance  $\rightarrow$  ObservableUnit C.Sigma  $\rightarrow$  Prop

-- Required projected Appendix J order at the DUT boundary.  
-- This is the mandated projected  $\prec_J$  relation induced by E1–E5, channel  
-- semantics, synchronization semantics, atomic-unit constraints,  
-- witness-control dependencies that survive projection, and elaboration /  
--  $\sigma$ \_DUT-projection constraints. It is not an incidental order that merely  
-- appears in one admissible schedule.

ProjectedRequiredJOrderDUT :

ComparisonInstance  $\rightarrow$

ObservableUnit C.Sigma  $\rightarrow$

ObservableUnit C.Sigma  $\rightarrow$

Prop

-- Two proof-internal observable units are ordered at the DUT boundary exactly  
-- when their projected  $\Sigma$ \_DUT-visible units are related by the required projected  
-- Appendix J order.

def DUTOrdered

(C : ComparisonInstance)

(u v : ObservableUnit C.Sigma) : Prop :=

SigmaDUTVisibleUnit C u  $\wedge$

SigmaDUTVisibleUnit C v  $\wedge$

ProjectedRequiredJOrderDUT C u v

-- A pair is causally comparable if either order direction is required at the

```

-- projected DUT boundary.
def DUTCausallyComparable
(C : ComparisonInstance)
(u v : ObservableUnit C.Sigma) : Prop :=
DUTOrdered C u v V DUTOrdered C v u

-- The choice may change the projected DUT-visible presence or value label of u.
ChoiceMayAddRemoveOrRelabelSigmaDUTUnit :
ComparisonInstance → Process → EpsModeledChoice C P → ObservableUnit C.Sigma →
Prop

-- The choice may change the relative order of two projected DUT-visible units.
ChoiceMayReorderSigmaDUTUnits :
ComparisonInstance → Process → EpsModeledChoice C P →
ObservableUnit C.Sigma → ObservableUnit C.Sigma → Prop

-- A single  $\varepsilon$ -modeled local timing choice is inert for  $\Sigma_{\text{DUT}}$  safety when it
-- neither changes projected event presence/value nor reorders any projected
-- pair whose order is causally required at the DUT boundary.
def EpsChoiceIsSigmaDUTSafetyInert
(C : ComparisonInstance)
(P : Process)
(choice : EpsModeledChoice C P) : Prop :=
(∀ u,
SigmaDUTVisibleUnit C u →
¬ ChoiceMayAddRemoveOrRelabelSigmaDUTUnit C P choice u) ∧
(∀ u v,
ChoiceMayReorderSigmaDUTUnits C P choice u v →
¬ DUTCausallyComparable C u v)

-- A latency-sensitive interface is  $\Sigma_{\text{DUT}}$ -safety-inert when all of the  $\varepsilon$ -modeled
-- local timing choices owned by the interface are safety-inert in the sense above.
def SigmaDUTSafetyInertInterface
(C : ComparisonInstance)
(I : LatencySensitiveInterface C) : Prop :=
∀ P choice,
ChoiceOwnedByInterface C I P choice →
EpsChoiceIsSigmaDUTSafetyInert C P choice

-- A choice has a  $\Sigma_{\text{DUT}}$ -safety discharge either by ordinary snooping/witness
-- coverage or by belonging to a certified  $\Sigma_{\text{DUT}}$ -safety-inert interface.
def ChoiceHasSigmaDUTSafetyDischarge
(C : ComparisonInstance)
(W : WitnessBundle C)

```

```

(P : Process)
(choice : EpsModeledChoice C P) : Prop :=
ChoiceCoveredBySnoopingOrWitness C W P choice  $\vee$ 
 $\exists$  I,
ChoiceOwnedByInterface C I P choice  $\wedge$ 
SigmaDUTSafetyInertInterface C I

```

-- The projected-safety obligation is deliberately weaker than WC. It covers  
-- only choices that can matter to the  $\Sigma$ \_DUT projected safety theorem.

```

def SigmaDUTSafetyWitnessObligation
(C : ComparisonInstance)
(W : WitnessBundle C) : Prop :=
 $\forall$  P choice,
EpsChoiceCanAffectSigmaDUTProjection C P choice  $\rightarrow$ 
ChoiceHasSigmaDUTSafetyDischarge C W P choice

```

-- Local Appendix L Note 1 corollary: a certified safety-inert interface  
-- discharges only its own choices. This theorem is intentionally local; it does  
-- not claim that an arbitrary identity witness satisfies the global obligation  
-- for every other choice in the design.

```

theorem L5_sigmaDUT_safety_inert_interface
(C : ComparisonInstance)
(W : WitnessBundle C)
(I : LatencySensitiveInterface C)
(hI : SigmaDUTSafetyInertInterface C I) :
 $\forall$  P choice,
ChoiceOwnedByInterface C I P choice  $\rightarrow$ 
ChoiceHasSigmaDUTSafetyDischarge C W P choice

```

-- No global L.5 assembly theorem is stated here. The global obligation  
-- SigmaDUTSafetyWitnessObligation C W is intentionally an assumption of  
-- JSigmaDUTSafetyAssumptions. Individual certified inert-interface choices are  
-- discharged locally by L5\_sigmaDUT\_safety\_inert\_interface; all other  
--  $\Sigma$ \_DUT-relevant choices must be discharged by the ordinary snooping/witness  
-- machinery at the point where the global obligation is established.

The L.5 facts are safety-side counterparts of the WFG-inert  $\varepsilon$  /  $K\varepsilon$  idea, but they  
live in Appendix L/J rather than Appendix K. They are used only for the  
 $\Sigma$ \_DUT-projected safety theorem below. They do not produce WC C W, do not  
prove CompatibleSys over the full proof-internal C.Sigma alphabet, and do not  
discharge A11\_Kepsilon\_WFGInert for deadlock/liveness proofs.



## 15. Concrete ObsSigma, Obs-congruence, and deterministic replay

Do not define ObsSigma as an informal least fixed point. Define it as a concrete record of the source-state slices required for replay. The record is P-local: it contains the P-owned state, frontier, pending, attribution, signal-anchor, direct-input, array-mapping, and R2C component-local observation-unit facts needed to determine later P-local  $\Sigma$ -visible behavior and Appendix I–K proof predicates. It does not independently store ChannelEnabled / ChannelDisabled as P-local facts; enablement is obtained from the P-local request/frontier/BoundaryReq facts together with the chosen Sys\_ref boundary state and E4. For combinational-process signal IO, ObsSigma records the selected component identity, local R2C occurrence ordinal, observable slice, and fixed-point values for emitted R2C units, not every stutter cycle of the combinational network.

```
structure ObsSigma
(C : ComparisonInstance)
(W : WitnessBundle C)
(P : Process)
( $\sigma$  : State) where
control_pos : ControlPos P
relevant_vars : RelevantVarSlice C P  $\sigma$ 
selected_eps_outcomes : SelectedOutcomeSlice C W.selector P  $\sigma$ 
nb_witness_slice : NBWitnessSlice C W.nbMap P  $\sigma$ 
msg_attribution_slice : MsgAttributionSlice C W P  $\sigma$ 
pending_slice : PendingSlice C P  $\sigma$ 
front_slice : FrontSlice C P  $\sigma$ 
nextfront_slice : NextFrontSlice C P  $\sigma$ 
boundary_req_slice : BoundaryReqSlice C P  $\sigma$ 
signal_anchor_slice : SignalAnchorSlice C P  $\sigma$ 
array_mapping_slice : ArrayMappingSlice C P  $\sigma$ 
direct_input_slice : DirectInputSlice C P  $\sigma$ 
r2c_unit_slice : R2CUnitSlice C P  $\sigma$ 
```

Here msg\_attribution\_slice records the  $Q_{\text{exec},P} / Q_{\Omega,P}$  attribution facts needed for E3/E5, Issue/Commit replay, WFG construction, and BoundaryReq ownership. It is broader than nb\_witness\_slice: nb\_witness\_slice records  $\mu_Q$  matching for  $\varepsilon$ -modeled failed NB calls, while msg\_attribution\_slice records the source-level operation attribution of pending and frontier message-operation instances generally. The r2c\_unit\_slice field is the previously missing explicit carrier for the combinational-process facts that the document's Definition I.Obs already requires: for each R2C component owned by P, the selected

component identity, the local R2C occurrence ordinal of the most recent emitted unit, the selected observable slice  $\text{Obs}(\text{C\_comb})$ , and the fixed-point values of emitted units — not every stutter cycle of the combinational network. Without this field,  $\text{IObs\_congruence\_one\_step}$  could not propagate R2C-unit agreement, even though the surrounding prose assumed it.

Cadence-sensitive failed-poll behavior is not represented by adding failed-poll cadence events to  $\text{ObsSigma}$  or to  $\Sigma$ . Failed  $\text{PushNB/PopNB}$  polls remain  $\varepsilon$ -steps with no committed transfer event. If a retry counter, timeout counter, polling arbiter, or other local policy computes externally visible behavior from the number or cadence of failed polls, then that site must already be handled by  $\text{C.contracts.latencySensitiveLocalProtocols}$ . The ordinary source-core  $\text{ObsSigma}$  record is then applied at the chosen observable boundary of the isolated/snoop-aligned/explicitly modeled block.

```
def SigmaRelevantEq C W P  $\sigma_1$   $\sigma_2$  :=
  ObsSigma C W P  $\sigma_1$  = ObsSigma C W P  $\sigma_2$ 
```

One-step Obs-congruence:

```
theorem IObs_congruence_one_step
  (C : ComparisonInstance)
  (W : WitnessBundle C)
  (hB : BRules C)
  (hR : RRules C)
  (hE : ERules C)
  (hIC : IssueCommitWF C W)
  (hContracts : ModelingContracts C)
  (hWC : WC C W)
  (P : Process)
  ( $\sigma_1$   $\sigma_2$   $\sigma_1'$   $\sigma_2'$  : State)
  ( $\ell$  : C.Sigma)
  (hObs : SigmaRelevantEq C W P  $\sigma_1$   $\sigma_2$ )
  (hStep1 : SameLabelCutpointStep C W P  $\sigma_1$   $\ell$   $\sigma_1'$ )
  (hStep2 : SameLabelCutpointStep C W P  $\sigma_2$   $\ell$   $\sigma_2'$ )
  (hBoundary : ChosenObservableBoundaryWF C W P hContracts)
  (hCadence :
    AllCadenceSensitiveNBSitesHandled
    C W P hContracts.latencySensitiveLocalProtocols) :
  SigmaRelevantEq C W P  $\sigma_1'$   $\sigma_2'$ 
```

Proof obligations for `IObs_congruence_one_step` should be decomposed by source construct:

- deterministic local computation: equal relevant variables and control position determine equal successor `ObsSigma` fields;
- anchored signal reads/writes: the fixed reactive stimulus and common label  $\ell$  determine equal latched `Read/Write` facts;
- blocking `Push/Pop` and synchronization: common label  $\ell$  plus equal `pending/frontier/attribution/BoundaryReq` facts determine equal successor `ObsSigma` fields;
- successful `PushNB/PopNB`: the committed `Push_c(v) / Pop_c(v)` label  $\ell$  determines equal data, control, attribution, and frontier updates;
- failed `PushNB/PopNB`: the failed poll is  $\varepsilon$ . If `NB-CF` holds, it cannot affect the next  $\Sigma$ -visible control/frontier/attribution facts. If `WC` selects a non-cadence-sensitive failed outcome, the selected outcome is fixed by `W`. If the site is cadence-sensitive, `hCadence` ensures it is not part of the unisolated ordinary source core; replay is performed at the chosen boundary of the isolated/snoop-aligned/explicit local protocol block;
- other  $\varepsilon$ -only choices: `WC` supplies every outcome that can affect later  $\Sigma$ -visible behavior or Appendix I–K proof predicates.

Deterministic replay:

theorem `IDR_deterministic_replay`

(`C` : `ComparisonInstance`)

(`W` : `WitnessBundle C`)

(`hB` : `BRules C`)

(`hR` : `RRules C`)

(`hE` : `ERules C`)

(`hIC` : `IssueCommitWF C W`)

(`hContracts` : `ModelingContracts C`)

(`hWC` : `WC C W`)

(`P` : `Process`)

( $\rho_1 \rho_2$  : `ExecutionPrefix C.Sigma`)

(`hInit` : `SameInitialSourceState C W P  $\rho_1 \rho_2$` )

(`hSameLabels` : `SamePLocalSigmaLabelSequence C P  $\rho_1 \rho_2$` )

(`hBoundary` : `ChosenObservableBoundaryWF C W P hContracts`)

(`hCadence` :

`AllCadenceSensitiveNBSitesHandled`

`C W P hContracts.latencySensitiveLocalProtocols`) :

$\forall k,$   
 $\text{SigmaRelevantEq } C \ W \ P$   
 $(\text{EpsQuiescentCutpointAfter } C \ W \ P \ \rho_1 \ k)$   
 $(\text{EpsQuiescentCutpointAfter } C \ W \ P \ \rho_2 \ k)$

Proof. By induction on  $k$ . The base case follows from  $\text{hInit}$ . For the successor case,  $\text{hSameLabels}$  gives the common next  $P$ -local  $\Sigma$  label  $\ell$ . Apply  $\text{IObs\_congruence\_one\_step}$  to the two same-label cutpoint steps under the same fixed stimulus, same witness  $W$ , and the same latency-sensitive local-protocol handling contract. Therefore the  $\text{ObsSigma}$  slices agree after  $k+1$  labels.

This same  $\text{ObsSigma}$  object and the  $\text{IObs\_congruence\_one\_step}$  theorem should be reused by Lemma S1, Corollary J.1, and any cutpoint-correspondence theorem that needs agreement of source variables, frontier predicates, signal-anchor facts, direct-input facts, or array-mapping facts. For request attribution,  $\text{BoundaryReq}$ , pending, endpoint-request, and enablement facts, the source-side deterministic-replay slice comes from  $\text{ObsSigma} / \text{I.DR}$ , but the RTL\_B-side cutpoint facts used by Corollary J.1 and Appendix K are supplied by  $\text{J.AbsConf}$  for the same selected  $(\tau_{\text{post}}, \tau_{\text{ref}}, \rho_{\text{post}}, \sigma_{\text{ref}}, W)$  triple. Do not treat  $\text{ObsSigma}$  replay alone as proving RTL\_B-side request-attribution agreement.

## 16. Appendix I per-process construction

Bundle the assumptions:

```

structure IConstructionAssumptions
(C : ComparisonInstance)
(W : WitnessBundle C) where
wellformed : C.WellFormed
bRules : BRules C
rRules : RRules C
eRules : ERules C
issueCommitWF : IssueCommitWF C W
implementation : ImplementationAssumptions C
witnessCompatible : WC C W
contracts : ModelingContracts C

```

Main theorems:

```

theorem I1_bounded_epsilon_closure
(C : ComparisonInstance)

```

```

(hB4 : C.assumptions.B.B4_finiteEpsilonDepth)
(hB5 : C.assumptions.B.B5_quiescentNormalization) :
...

```

```

theorem I2_channel_progress_instance
(C : ComparisonInstance)
(hB3 : C.assumptions.B.B3_channelProgress) :
...

```

```

theorem I3_finite_extension_step
(C : ComparisonInstance)
(W : WitnessBundle C)
(hI : IConstructionAssumptions C W) :
...

```

```

theorem I4_finite_prefix_construction
(C : ComparisonInstance)
(W : WitnessBundle C)
(hI : IConstructionAssumptions C W) :
...

```

```

theorem I5_infinite_limit_construction
(C : ComparisonInstance)
(W : WitnessBundle C)
(hI : IConstructionAssumptions C W) :
...

```

I3 should use IssueAt plus derived IssueTime, not an unconstrained guessed issue cycle. For v42 it must also keep the E4/B2 split explicit: matched channel history and B3 can establish that a pending operation eventually becomes continuously ChannelEnabled; B2 then gives eventual scheduler/arbiter selection, and the selection-commit soundness clause gives the actual Commit in the selected cycle. I3 should not collapse ChannelEnabled into immediate same-cycle Commit.

## 17. Appendix J system-level correspondence

Define execution scopes:

```
def FairProgressSatisfying
```

```
(C : ComparisonInstance)
```

```
(side : Side)
```

```
( $\tau^\infty$  : InfiniteExecutionBucketed C.Sigma) : Prop :=
```

```
InfiniteExecutionOfSide C side  $\tau^\infty$   $\wedge$ 
```

```
WeakFairness C side  $\tau^\infty$   $\wedge$ 
```

```
 $\forall$  c,
```

```
P_push C side  $\tau^\infty$  c  $\wedge$ 
```

```
P_pop C side  $\tau^\infty$  c
```

```
def PostAdmissible
```

```
(C : ComparisonInstance)
```

```
( $\tau^\infty$  : InfiniteExecutionBucketed C.Sigma) : Prop :=
```

```
InfinitePostExecution C  $\tau^\infty$   $\wedge$ 
```

```
FairProgressSatisfying C .post  $\tau^\infty$ 
```

```
def RefAdmissible
```

```
(C : ComparisonInstance)
```

```
( $\tau^\infty$  : InfiniteExecutionBucketed C.Sigma) : Prop :=
```

```
InfiniteRefExecution C  $\tau^\infty$   $\wedge$ 
```

```
FairProgressSatisfying C .ref  $\tau^\infty$ 
```

```
def Exec_post (C : ComparisonInstance) ( $\tau$  : Prefix C.RTL_B) : Prop :=
```

```
 $\exists$   $\tau^\infty$ ,
```

```
Extends  $\tau^\infty$   $\tau$   $\wedge$ 
```

```
PostAdmissible C  $\tau^\infty$ 
```

```
def Exec_ref (C : ComparisonInstance) ( $\tau$  : Prefix C.Sys_ref) : Prop :=
```

```
 $\exists$   $\tau^\infty$ ,
```

```
Extends  $\tau^\infty$   $\tau$   $\wedge$ 
```

```
RefAdmissible C  $\tau^\infty$ 
```

Here FairProgressSatisfying is side-parametric. It includes the B2/B3 fairness/progress assumptions for the selected execution side: weak fairness and the P\_push / P\_pop channel-progress obligations. PostAdmissible is the post-side instance of this condition together with InfinitePostExecution; RefAdmissible is the reference-side instance together with InfiniteRefExecution. Thus Exec\_post and Exec\_ref are defined analogously, and any  $\tau_{\text{ref}}$  produced or consumed by

J.1 / K.1 carries the ref-side B2/B3 progress assumptions needed by the drain-closure and deadlock-preservation arguments.

B6 is the globally fixed-point idle-stutter theorem, not an arbitrary deadlock-cutpoint extension principle. It is the general time-divergence / idle-stutter convention only for a reachable finite prefix/state at which the system has reached a B5/global fixed point: after no further  $\Sigma$ -action or internal  $\varepsilon$ -microstep is enabled under the same stimulus, the global clock may continue by infinite idle-stutter ticks. Each idle-stutter bucket has no  $\sigma$ Events and performs no  $\varepsilon$ -prefix or  $\varepsilon$ -suffix work. It changes no non-clock state; the only permitted state change is the clock/tick bookkeeping associated with the idle cycle.

```
def ClockOnlyIdleStep
(C : ComparisonInstance)
(side : Side)
( $\sigma_0 \sigma_1$  : State) : Prop :=
NonClockStateEq  $\sigma_1 \sigma_0 \wedge$ 
 $\sigma_1$ .cycle =  $\sigma_0$ .cycle + 1  $\wedge$ 
 $\sigma_1$ .epsPos = 0
```

```
def IdleStutterBucket
(C : ComparisonInstance)
(side : Side)
( $\sigma$  : State)
( $\beta$  : CycleBucket C.Sigma) : Prop :=
 $\beta$ .startState = AdvanceClockOnly  $\sigma \beta$ .cycle  $\wedge$ 
 $\beta$ .sigmaEvents =  $\emptyset \wedge$ 
 $\beta$ .epsPrefix = []  $\wedge$ 
 $\beta$ .epsSuffix = []  $\wedge$ 
ClockOnlyIdleStep C side  $\beta$ .startState  $\beta$ .endState  $\wedge$ 
 $\beta$ .endQuiescent
```

```
def FixedPointState
(C : ComparisonInstance)
(side : Side)
( $\sigma$  : State) : Prop :=
EpsQuiescent  $\sigma \wedge$ 
NoObservableActionEnabledAtState C side  $\sigma \wedge$ 
NoInternalEpsilonStepEnabledAtState C side  $\sigma$ 
```

```

theorem B6_extend_reachable_prefix_by_idle_stutter
(C : ComparisonInstance)
(side : Side)
( $\pi$  : Prefix (SystemOfSide C side))
( $\sigma$  : State)
(hReach : TerminalCutpoint  $\pi$   $\sigma$ )
(hFP : FixedPointState C side  $\sigma$ ) :
 $\exists \tau$  : InfiniteExecutionBucketed C.Sigma,
ExtendsPrefix C side  $\tau$   $\pi$   $\wedge$ 
 $\forall n$ ,
 $n \geq \text{PrefixLength } \pi \rightarrow$ 
IdleStutterBucket C side  $\sigma$  ( $\tau$   $n$ ) :=
by
-- General B6 theorem: any reachable finite execution prefix/state at a B5/global
-- fixed point is extendable to an infinite execution by appending idle-stutter
-- ticks under the same stimulus. B6 supplies time divergence only. It does not
-- by itself prove the B2/B3 fairness/progress side conditions required for
-- membership in Exec_post or Exec_ref at an arbitrary partial internal-MP
-- deadlock cutpoint.
sorry

```

Thus any broad prose summary of B6 should be read with this B5/global fixed-point side condition.

The v42 K.1 bridge for a partial internal-MP deadlock is therefore not this B6 idle-stutter theorem. It is the separate fair drain-normalization construction in Section 18: already-issued D-owned non-frontier requests that later become ChannelEnabled may commit as finite drain-only progress; outside-D and environment endpoints continue under the ordinary B2/B3 progress premises; and pure idle-stutter is used only for D-owned internal frontier blockers that remain ChannelDisabled inside the closed WFG component.

```

theorem finish_process_idle_padding
(C : ComparisonInstance)
(side : Side)
(P : Process)
( $\pi$  : Prefix (SystemOfSide C side))

```



$(\sigma : \text{State})$   
 $(h\text{Reach} : \text{TerminalCutpoint } \pi \sigma)$   
 $(h\text{Fin} : \text{ProcessFinishedAt } C \text{ side } P \sigma)$   
 $(h\text{NoMoreP} : \text{NoFurtherProcessStepEnabled } C \text{ side } P \sigma) :$   
 $\exists \tau : \text{InfiniteExecutionBucketed } C.\text{Sigma},$   
 $\text{ExtendsPrefix } C \text{ side } \tau \pi \wedge$   
 $\forall n,$   
 $n \geq \text{PrefixLength } \pi \rightarrow$   
 $\text{ProcessIdleStuttered } C \text{ side } P (\tau n) :=$   
 by  
 -- FINISH\_P remains a realized SyncAnchor for R2/E2 signal anchoring, but  
 -- after the process has terminated, the infinite execution is padded by  
 -- process-local idle stutter. This padding creates no new process-owned  
 --  $\Sigma$ -events, no new BoundaryReq, no new Front/NextFront item, and no new  
 -- WFG edge for P.  
 sorry

The FINISH\_P corollary is process-local: a finished process is padded by  $\varepsilon/\text{idle-stutter}$  while the global system may continue executing other processes. This padding must not be counted as additional B4/B5 internal progress by P, because it performs no enabled internal  $\varepsilon$ -step and emits no  $\Sigma$ -event.

Directed compatibility:

Do not define CompatibleP or CompatibleSys as equality of bucketed executions.

First define the observable units extracted from a bucketed execution. An observable unit may be: a single ordinary  $\Sigma$ -event; an explicitly atomic multi-endpoint unit, such as a rendezvous Push/Pop pair or a SyncChannel/ac\_sync pair; or a component-local R2C fixed-point signal-observation unit emitted under the v42 stutter-canonical R2C convention. Independent same-cycle events are not ordered merely because they share a CycleBucket; any strict required order among observable units is imposed only by RequiredJOrder /  $\prec_J$ . The broader causal-dependency graph  $\rightarrow$  may still be recorded for explanatory information-flow reasoning, but it is not the induction order.

structure ObservableUnit ( $\Sigma : \text{Type}$ ) where  
 events : Finset  $\Sigma$   
 atomic : Prop

-- ObservableUnitWF is the well-formedness predicate for the unit extraction used  
 -- by Appendix J. It is where non-splitting and grouping obligations are checked:  
 -- rendezvous Push/Pop pairs, paired SyncChannel/ac\_sync endpoints, R2C  
 -- fixed-point observations, and any other explicitly atomic same-cycle  
 -- transaction must appear as one ObservableUnit rather than as strict internal  
 -- RequiredJOrder edges. It also records that every  $\Sigma$ -event in the finite prefix  
 -- belongs to exactly one extracted unit and that independent same-cycle events  
 -- are not ordered merely by bucket membership.

```
structure ObservableUnitWF
(C : ComparisonInstance)
(side : Side)
( $\tau$  : Prefix (SystemOfSide C side)) where
covers_prefix_events :
 $\forall e,$ 
SigmaEventInPrefix C side  $\tau e \rightarrow$ 
 $\exists u,$ 
 $u \in \text{ObservableUnitsOf } C \text{ side } \tau \wedge$ 
 $e \in u.\text{events}$ 
events_have_unique_unit :
 $\forall e u v,$ 
 $u \in \text{ObservableUnitsOf } C \text{ side } \tau \rightarrow$ 
 $v \in \text{ObservableUnitsOf } C \text{ side } \tau \rightarrow$ 
 $e \in u.\text{events} \rightarrow$ 
 $e \in v.\text{events} \rightarrow$ 
 $u = v$ 
atomic_rendezvous_grouped :
 $\forall u,$ 
AtomicRendezvousUnit C side  $\tau u \rightarrow$ 
 $u \in \text{ObservableUnitsOf } C \text{ side } \tau \wedge u.\text{atomic}$ 
atomic_sync_grouped :
 $\forall u,$ 
AtomicSyncUnit C side  $\tau u \rightarrow$ 
 $u \in \text{ObservableUnitsOf } C \text{ side } \tau \wedge u.\text{atomic}$ 
r2c_fixed_point_unit_grouped :
 $\forall u,$ 
R2CFixedPointObservationUnit C side  $\tau u \rightarrow$ 
 $u \in \text{ObservableUnitsOf } C \text{ side } \tau \wedge u.\text{atomic}$ 
atomic_units_not_split :
 $\forall u e_1 e_2,$ 
 $u \in \text{ObservableUnitsOf } C \text{ side } \tau \rightarrow$ 
 $u.\text{atomic} \rightarrow$ 
 $e_1 \in u.\text{events} \rightarrow$ 
 $e_2 \in u.\text{events} \rightarrow$ 
SameObservableUnit C side  $\tau$ 
```

```

(ObservableUnit.singleton e1)
(ObservableUnit.singleton e2)
independent_same_cycle_not_ordered_by_bucket :
∀ u v,
u ∈ ObservableUnitsOf C side τ →
v ∈ ObservableUnitsOf C side τ →
u ≠ v →
SameCycleOnly C side τ u v →
¬ RequiredJOrderGenerator C side τ u v ∧
¬ RequiredJOrderGenerator C side τ v u

```

ObservableUnitsOf extracts  $\Sigma$ -visible events from the finite bucketed prefix and groups them into comparison units as follows:

- an ordinary committed message, signal, or synchronization event may form a singleton observable unit;
- explicitly atomic multi-endpoint events, such as rendezvous Push/Pop pairs and SyncChannel/ac\_sync handshakes, form one atomic observable unit and may not be split by CompatibleP or CompatibleSys;
- R2C combinational observations are emitted as component-local fixed-point signal-observation units according to R2C\_WF and are matched by local occurrence ordinal;
- independent same-cycle events are unordered merely by virtue of being in the same CycleBucket;
- same-cycle events that are merely coincident in one implementation are not forced to remain same-cycle in the other implementation unless they are explicitly atomic or otherwise ordered by RequiredJOrder.

Because  $\tau$  is a finite prefix, ObservableUnitsOf returns a Finset. If a Set view is needed elsewhere, it should be obtained by coercion from this Finset, rather than by introducing a second observable-unit universe.

```

def ObservableUnitsOf
(C : ComparisonInstance)
(side : Side)
(τ : Prefix (SystemOfSide C side)) :
Finset (ObservableUnit (C.Sigma)) :=
ExtractSigmaUnitsFinset
C side τ
(AtomicRendezvousUnits C side τ)
(AtomicSyncUnits C side τ)
(R2CFixedPointObservationUnits C side τ)
(IndependentSameCycleUnits C side τ)

```

```

def ObservableUnitsOfSet

```

(C : ComparisonInstance)

(side : Side)

( $\tau$  : Prefix (SystemOfSide C side)) :

Set (ObservableUnit (C.Sigma)) :=

fun u => u ∈ ObservableUnitsOf C side  $\tau$

CausalDependency is an information-flow graph. It is useful for explaining that a Pop observes the corresponding Push value, that signal handshakes communicate information, or that a witness-selected  $\varepsilon$  choice affects later visible behavior. It is not itself the strict partial order used by Theorem J.1, and the Lean development should not assume the unquotiented CausalDependency graph is acyclic. CausalDependency is not part of the core mechanized theorem spine for Theorem J.1, Corollary J.1, Appendix K, or Appendix R. The core proof order is RequiredJOrder /  $\leq_J$  over observable units, together with ObservableUnitWF, E1–E5 legality/alignment, witness compatibility, cutpoint correspondence, and the Appendix K WFG/deadlock machinery. If CausalDependency is formalized in Lean, it should be treated as an optional documentation/design-well-formedness layer used to mirror the scheduling-rules document’s explanatory causal-dependency relation, not as a dependency of the main trace-compatibility or liveness-preservation theorems.

def CausalDependency

(C : ComparisonInstance)

(side : Side)

( $\tau$  : Prefix (SystemOfSide C side))

(u v : ObservableUnit (C.Sigma)) : Prop :=

Relation.TransGen

(CausalDependencyGenerator C side  $\tau$ )

u v

-- Intra-process causal trace/source order is part of CausalDependency, because it

-- records source-level information flow and design intent inside one dynamic

-- process execution. It is intentionally non-temporal: by itself it does not

-- assert strict commit-cycle order, and it is not a RequiredJOrder edge unless

-- some explicit required-order generator below also applies.

def IntraProcessCausalTraceOrder

(C : ComparisonInstance)

(side : Side)

( $\tau$  : Prefix (SystemOfSide C side))

(u v : ObservableUnit (C.Sigma)) : Prop :=

SameOwningProcess C side  $\tau$  u v  $\wedge$

SourceTracePrecedes C side  $\tau$  u v  $\wedge$

CausallyRelevantSourceFlow C side  $\tau$  u v

-- Declared shared-storage / memory-serialization dependence is also part of

-- CausalDependency. It records only dependencies supplied by the array-access

-- mapping layer, explicit synchronization, atomic memory semantics, or declared  
 -- serialization/coherence rules for mapped shared storage. It does not turn a  
 -- memory information-flow edge into a proof-order edge unless the selected  
 -- mapping contract also generates an explicit RequiredJOrder edge or atomic unit.

```
def DeclaredMemorySerializationDependency
(C : ComparisonInstance)
(side : Side)
( $\tau$  : Prefix (SystemOfSide C side))
(u v : ObservableUnit (C.Sigma)) : Prop :=
SharedStorageConflict C side  $\tau$  u v  $\wedge$ 
(ArrayMappingSerializes C side  $\tau$  u v  $\vee$ 
ExplicitSyncOrdersMemoryAccess C side  $\tau$  u v  $\vee$ 
DeclaredAtomicMemorySemanticsOrder C side  $\tau$  u v  $\vee$ 
DeclaredCoherenceSerializationOrder C side  $\tau$  u v)
```

```
def CausalDependencyGenerator
(C : ComparisonInstance)
(side : Side)
( $\tau$  : Prefix (SystemOfSide C side))
(u v : ObservableUnit (C.Sigma)) : Prop :=
MessageInformationFlow C side  $\tau$  u v  $\vee$ 
SignalInformationFlow C side  $\tau$  u v  $\vee$ 
SynchronizationInformationFlow C side  $\tau$  u v  $\vee$ 
IntraProcessCausalTraceOrder C side  $\tau$  u v  $\vee$ 
DeclaredMemorySerializationDependency C side  $\tau$  u v  $\vee$ 
WitnessControlDependency C side  $\tau$  u v  $\vee$ 
ElaborationInformationFlow C side  $\tau$  u v
```

RequiredJOrder is the mandated Appendix J order over observable units. It is narrower than CausalDependency. It contains only strict order constraints that CompatibleP / CompatibleSys must preserve as order. It does not contain mere information-flow dependencies, grouping requirements, or source-lexical order that v42 intentionally treats as non-ordering.

The generator must encode the v42 exclusions explicitly:

- E2 signal anchoring contributes equality/alignment of the corresponding anchor cycle and value label; it never contributes an independent strict RequiredJOrder edge.
- E3 same-interface message operations may impose strict same-interface order.

Strict issue/commit attribution is used by the fixed RTL-time witness construction schedule, not as an independent RequiredJOrder /  $\leq_J$  generator.

In particular, issue-attribution facts may help choose the next replay unit, but they do not add observable order among units that v42 otherwise leaves unordered.

- E3 distinct-interface lexical source order is only a no-reverse legality constraint. It is not a positive strict edge between the distinct-interface commits. In particular, two distinct-interface operations that v42 permits to issue/commit in parallel do not create  $u <_J v$  merely because one appears earlier in source text.
- Atomic rendezvous, SyncChannel/ac\_sync handshakes, and R2C fixed-point observations are represented as single ObservableUnit objects. Their non-splitting constraints are checked by ObservableUnitWF and CompatibleSys; they do not add strict internal edges.
- Witness-control dependencies remain in CausalDependency and WC, where they justify selected  $\varepsilon$  outcomes. They are not RequiredJOrder generators.
- Elaboration legality is represented by the explicit Sys\_B^elab channels, E4 channel-realization semantics, and  $\Pi_{\text{elab}} / \text{sigmaDUTProjection}$  projection constraints. It is not a separate RequiredJOrder generator unless it is realized as one of the explicit strict channel or synchronization edges below.

```
def E2SignalAnchorAlignmentConstraint
(C : ComparisonInstance)
(side : Side)
( $\tau$  : Prefix (SystemOfSide C side))
( $u\ v$  : ObservableUnit (C.Sigma)) : Prop :=
SignalAnchorCyclesAndLabelsAligned C side  $\tau\ u\ v$ 
```

```
def E3MessageStrictRequiredOrder
(C : ComparisonInstance)
(side : Side)
( $\tau$  : Prefix (SystemOfSide C side))
( $u\ v$  : ObservableUnit (C.Sigma)) : Prop :=
E3SameInterfaceMessageOrder C side  $\tau\ u\ v$ 
```

```
def E3DistinctInterfaceNoReverseLegality
(C : ComparisonInstance)
(side : Side)
( $\tau$  : Prefix (SystemOfSide C side)) : Prop :=
 $\forall\ u\ v,$ 
DistinctMessageInterfaces C side  $\tau\ u\ v \rightarrow$ 
SourceLexicallyBefore C side  $\tau\ u\ v \rightarrow$ 
 $\neg$  RequiredReverseCommitOrder C side  $\tau\ v\ u$ 
```

-- E5 has two logically different parts.

--

-- The pref\_S(s) side is non-strict legality:

--  $q \in \text{pref\_S}(s)$  gives  $\text{issue}(q) \leq \text{clk}(s)$ , and for committed blocking  $q$

-- the commit cycle of  $q$  is  $\leq \text{clk}(s)$ , with equality allowed.

```

-- It is therefore not a strict RequiredJOrder generator.
--
-- The suff_S(s) side is the strict post-boundary side:
-- q ∈ suff_S(s) gives issue(q) > clk(s), and commit cannot precede issue.
-- Thus the synchronization unit for s is strictly before the observable
-- message unit for q. This is the only E5 contribution to RequiredJOrder.
def E5PrefSyncBoundaryLegality
  (C : ComparisonInstance)
  (side : Side)
  (τ : Prefix (SystemOfSide C side)) : Prop :=
  ∀ s q uSync uMsg,
  SyncUnitOf C side τ s uSync →
  MessageUnitOf C side τ q uMsg →
  q ∈ prefS s →
  IssuedInPrefix C side τ q →
  IssueTimeOf C side τ q ≤ CommitCycleOfUnit C side τ uSync ∧
  (Blocking q →
    CommittedMessageUnit C side τ q uMsg →
    CommitCycleOfUnit C side τ uMsg ≤ CommitCycleOfUnit C side τ uSync)

def E5StrictSuffSyncBoundaryRequiredOrder
  (C : ComparisonInstance)
  (side : Side)
  (τ : Prefix (SystemOfSide C side))
  (u v : ObservableUnit (C.Sigma)) : Prop :=
  ∃ s q,
  SyncUnitOf C side τ s u ∧
  MessageUnitOf C side τ q v ∧
  q ∈ suffS s ∧
  IssuedInPrefix C side τ q ∧
  CommitCycleOfUnit C side τ u < CommitCycleOfUnit C side τ v

-- Keep the old name as the RequiredJOrder generator name, but define it to be
-- exactly the strict successor-side E5 relation.
def E5SyncBoundaryRequiredOrder
  (C : ComparisonInstance)
  (side : Side)
  (τ : Prefix (SystemOfSide C side))
  (u v : ObservableUnit (C.Sigma)) : Prop :=
  E5StrictSuffSyncBoundaryRequiredOrder C side τ u v

-- The actual generator is intentionally the narrow v42 generator below.
def RequiredJOrderGenerator
  (C : ComparisonInstance)

```

```

(side : Side)
( $\tau$  : Prefix (SystemOfSide C side))
( $u\ v$  : ObservableUnit (C.Sigma)) : Prop :=
E1SyncOrder C side  $\tau\ u\ v\ V$ 
E3MessageStrictRequiredOrder C side  $\tau\ u\ v\ V$ 
E4ChannelRequiredOrder C side  $\tau\ u\ v\ V$ 
E5SyncBoundaryRequiredOrder C side  $\tau\ u\ v\ V$ 
SynchronizationUnitRequiredOrder C side  $\tau\ u\ v$ 

```

```

structure RequiredJOrderWF
(C : ComparisonInstance)
(side : Side)
( $\tau$  : Prefix (SystemOfSide C side)) where
-- This is stated as a generator-level exclusion for proof hygiene. A more
-- minimal formalization could prove the same fact as a consequence of the
-- concrete generator definitions, but keeping it in RequiredJOrderWF is
-- harmless and matches the v42 intent that E2 contributes alignment only, not
-- strict observable order.
e2_alignment_only :
 $\forall\ u\ v,$ 
E2SignalAnchorAlignmentConstraint C side  $\tau\ u\ v \rightarrow$ 
 $\neg$  RequiredJOrderGenerator C side  $\tau\ u\ v\ \wedge$ 
 $\neg$  RequiredJOrderGenerator C side  $\tau\ v\ u$ 
distinct_interface_no_positive_source_edge :
 $\forall\ u\ v,$ 
DistinctMessageInterfaces C side  $\tau\ u\ v \rightarrow$ 
SourceLexicallyBefore C side  $\tau\ u\ v \rightarrow$ 
 $\neg$  RequiredJOrderGenerator C side  $\tau\ u\ v$ 
atomicity_is_grouping_not_order :
 $\forall\ u\ v,$ 
AtomicUnitNonSplittingConstraint C side  $\tau\ u\ v \rightarrow$ 
SameObservableUnit C side  $\tau\ u\ v$ 
witness_control_not_order :
 $\forall\ u\ v,$ 
WitnessControlDependency C side  $\tau\ u\ v \rightarrow$ 
 $\neg$  RequiredJOrderGenerator C side  $\tau\ u\ v$ 
elaboration_not_separate_order :
 $\forall\ u\ v,$ 
ElaborationInformationFlow C side  $\tau\ u\ v \rightarrow$ 
(RequiredJOrderGenerator C side  $\tau\ u\ v \rightarrow$ 
E1SyncOrder C side  $\tau\ u\ v\ V$ 
E3MessageStrictRequiredOrder C side  $\tau\ u\ v\ V$ 
E4ChannelRequiredOrder C side  $\tau\ u\ v\ V$ 

```



E5SyncBoundaryRequiredOrder C side  $\tau \cup v \vee$   
 SynchronizationUnitRequiredOrder C side  $\tau \cup v$ )

RequiredJOrderWF is not intended to be a new external tool-flow assumption. It is a definitional sanity theorem about the concrete RequiredJOrderGenerator disjuncts once E1–E5, issue/commit well-formedness, observable-unit grouping, and elaboration well-formedness have been instantiated. The Lean development may keep it as a named structure for proof modularity, but J.1 should obtain it from the following lemma rather than expecting it to appear magically in JAssumptions.

```
theorem RequiredJOrderWF_from_concrete_generators
  (C : ComparisonInstance)
  (W : WitnessBundle C)
  (side : Side)
  ( $\tau$  : Prefix (SystemOfSide C side))
  (hE : ERules C)
  (hIC : IssueCommitWF C W)
  (hOU : ObservableUnitWF C side  $\tau$ )
  (hImpl : ImplementationAssumptions C)
  (hElab : ElaborationProjectionWF C side  $\tau$ ) :
  RequiredJOrderWF C side  $\tau$  :=
  by
    -- Unfold RequiredJOrderGenerator and discharge each WF field by cases on the
    -- concrete v42 generator disjunct:
    -- • E2 contributes alignment only;
    -- • distinct-interface E3 lexical order is no-reverse legality only;
    -- • the pref_S(s) side of E5 is non-strict issue/commit legality only and
    --   contributes no RequiredJOrder edge;
    -- • the suff_S(s) side of E5 is the only E5 strict RequiredJOrder generator;
    -- • atomicity is already enforced by ObservableUnitWF grouping;
    -- • witness-control dependencies remain in WC and CausalDependency only; and
    -- • elaboration information flow contributes order only when realized through
    --   an explicit E1/E3-same-interface/E4-strict-channel/E5-suff/synchronization
    --   generator.
  sorry

def RequiredJOrder
  (C : ComparisonInstance)
  (side : Side)
  ( $\tau$  : Prefix (SystemOfSide C side))
  (u v : ObservableUnit (C.Sigma)) : Prop :=
  Relation.TransGen
  (RequiredJOrderGenerator C side  $\tau$ )
```

u v

theorem RequiredJOrder\_acyclic\_on\_observable\_units

(C : ComparisonInstance)

(side : Side)

( $\tau$  : Prefix (SystemOfSide C side))

(hWF : RequiredJOrderWF C side  $\tau$ ) :

Acyclic (RequiredJOrder C side  $\tau$ ) :=

by

-- Strict RequiredJOrder edges are exactly the v42 strict observable-order  
 -- edges above. Each edge is justified by an explicit synchronization order,  
 -- same-interface message order, FIFO/channel order, rendezvous/channel order,  
 -- the strict successor side of synchronization-boundary separation, or paired  
 -- synchronization/barrier before-after constraints. Atomic multi-endpoint  
 -- transactions are single observable units, E2 signal anchors are alignment  
 -- constraints, distinct-interface lexical order is no-reverse legality rather  
 -- than a positive edge, and witness-control / elaboration information-flow  
 -- dependencies are not in this generator.

--

-- Issue/commit attribution is intentionally not part of RequiredJOrder. It is  
 -- used by RTLTimeNextObservableUnit and the fixed RTL-time ideal-chain  
 -- construction schedule to choose a realizable next unit, but that construction  
 -- schedule adds no observable order.

--

-- The acyclicity proof is the v42 Theorem J.1 argument instantiated per fixed  
 -- legal execution prefix, but with the E5 split made explicit. Assign each  
 -- observable unit  $\gamma$  its commit cycle  $\text{clk}(\gamma)$  in that prefix, with explicitly  
 -- atomic same-cycle units carrying the single common commit cycle of their  
 -- endpoints. Every strict generator edge  $\gamma_1 \prec_J \gamma_2$  implies  $\text{clk}(\gamma_1) < \text{clk}(\gamma_2)$ :

--

-- • E1 synchronization order is strict between successive synchronization units.  
 -- • E3 same-interface message order is strict for the same-interface required  
 -- order generator.  
 -- • E4 channel-order generators are strict for reset-empty, bounded-occupancy,  
 -- non-fall-through, and explicit rendezvous/channel ordering constraints.  
 -- • E5 contributes a strict edge only from the synchronization unit  $s$  to a  
 -- committed message unit for  $q \in \text{suff\_S}(s)$ . This edge is strict because  
 --  $\text{issue}(q) > \text{clk}(s)$  and  $\text{commit}(q)$  cannot occur before  $\text{issue}(q)$ .  
 -- • Paired synchronization/barrier before-after constraints are strict by their  
 -- own generator definitions.

--

-- The  $\text{pref\_S}(s)$  half of E5 is deliberately absent from RequiredJOrder. For  
 --  $q \in \text{pref\_S}(s)$ , E5 gives  $\text{issue}(q) \leq \text{clk}(s)$ , and for a committed blocking  $q$   
 -- the commit cycle of  $q$  is  $\leq \text{clk}(s)$ , with equality allowed. Those facts are

-- E5 legality / boundary well-formedness facts used by issue/commit,  
 -- frontier, cutpoint, and compatibility proofs; they are not strict  $\prec_J$  edges  
 -- and are not used in the acyclicity measure.  
 --  
 -- Hence any hypothetical  $\prec_J$  cycle would force  
 --  $\text{clk}(\gamma_0) < \dots < \text{clk}(\gamma_0)$ , which is impossible. The measure for the actual  
 -- strict RequiredJOrder graph is commit cycle alone. Issue timestamps are used  
 -- only to prove that the strict  $\text{suff}_S(s)$  E5 generator is indeed commit-cycle  
 -- increasing; they are not a second component of the acyclicity measure.  
 -- E2 alignment, atomic grouping, distinct-interface no-reverse legality,  
 --  $\text{pref}_S(s)$  E5 legality, witness-control dependencies, and elaboration  
 -- information flow contribute no strict edges, per RequiredJOrderWF.  
 sorry

CompatibleP : ComparisonInstance  $\rightarrow$  Process  $\rightarrow$  Prefix C.Sys\_ref  $\rightarrow$  Prefix C.RTL\_B  $\rightarrow$  Prop

CompatibleP C P  $\tau_{\text{ref}}$   $\tau_{\text{post}}$  means:

the P-local observable units of  $\tau_{\text{ref}}$  and  $\tau_{\text{post}}$  are related by a value-label-preserving bijection;

explicitly atomic units are mapped to explicitly atomic units and are not split;

the mapping preserves the P-local RequiredJOrder constraints induced by E1, E3  
 same-interface message order, E4, the strict  $\text{suff}_S$  side of E5, and the explicit  
 synchronization constraints;

it preserves E2 signal-anchor equality/alignment constraints, but does not treat  
 E2 as an independent source of strict order;

it preserves the  $\text{pref}_S$  side of E5 as issue/commit legality and synchronization-boundary  
 well-formedness, not as a strict RequiredJOrder edge;

it treats distinct-interface source-order no-reverse as legality, not as a strict edge; and  
 it does not require independent P-local units to occupy the same cycle bucket in both  
 executions.

Issue/commit attribution and selected witness/request-attribution facts are preserved  
 through RTLTimeNextObservableUnit, J.Abs, and J.AbsConf obligations; they are not  
 additional RequiredJOrder edges.

CompatibleSys : ComparisonInstance  $\rightarrow$  Prefix C.Sys\_ref  $\rightarrow$  Prefix C.RTL\_B  $\rightarrow$  Prop

CompatibleSys C  $\tau_{\text{ref}}$   $\tau_{\text{post}}$  means:

$\forall P$ , CompatibleP C P  $\tau_{\text{ref}}$   $\tau_{\text{post}}$ ;

all explicit abstract channels in the chosen Sys\_ref / RTL\_B comparison satisfy the  
 applicable E4 channel semantics at the agreed abstract boundaries;

the system-level observable-unit mapping preserves the mandated global RequiredJOrder /  $\prec_J$  constraints and value labels;  
 explicitly atomic multi-endpoint units remain atomic on both sides;  
 E2 anchor alignment, distinct-interface no-reverse legality, and elaboration projection legality are preserved by their own well-formedness/projection obligations rather than by extra strict  $\prec_J$  edges;  
 no equality of CycleBucket lists, no equality of bucket cycle numbers, and no bucket-by-bucket equality of sigmaEvents is required.

Thus the bijection used by CompatibleSys is over ObservableUnitsOf, and its ordering obligation is preservation of the narrowed v42 RequiredJOrder. Independent observable units may be placed in different cycle buckets or grouped differently in Sys\_ref and RTL\_B, provided their  $\Sigma$  labels, values, atomicity requirements, FIFO/signal/synchronization obligations, no-reverse legality, and required  $\prec_J$  constraints are preserved.

The notation  $\tau_{\text{ref},\text{system}} \approx_{\text{sys}} \tau_{\text{post},\text{system}}$  is an abbreviation for  $\text{CompatibleSys } C \tau_{\text{ref}} \tau_{\text{post}}$ . It is a directed compatibility predicate, not a symmetric equivalence relation. The main construction direction is Post  $\rightarrow$  Ref; the reverse direction is restricted to RTL-consistent source prefixes.

J.0:

```
theorem J0_pairwise_composition
  (C : ComparisonInstance)
  (W : WitnessBundle C)
  (hC : C.WellFormed)
  (hB : BRules C)
  (hR : RRules C)
  (hE : ERules C)
  (hIC : IssueCommitWF C W)
  (hContracts : ModelingContracts C)
  ( $\tau_{\text{ref}}$  : Prefix C.Sys_ref)
  ( $\tau_{\text{post}}$  : Prefix C.RTL_B)
  (hProc :  $\forall P, \text{CompatibleP } C P \tau_{\text{ref}} \tau_{\text{post}}$ )
  (hChan : ChannelsWellFormedAndMatched C  $\tau_{\text{ref}} \tau_{\text{post}}$ ) :
  CompatibleSys C  $\tau_{\text{ref}} \tau_{\text{post}}$ 
```

-- Projected process compatibility for the outer DUT/testbench-visible alphabet.  
 -- This is weaker than CompatibleP over C.Sigma: it compares only the

-- C.sigmaDUTProjection image of the process-owned observable units, preserving  
 -- projected values, atomicity that remains visible at  $\Sigma\_DUT$ , and required  
 -- projected RequiredJOrder /  $\prec\_J$  constraints. It may ignore proof-internal units hidden

-- by the projection and may commute projected units that are not causally  
 -- comparable at the DUT boundary.

```
def CompatiblePDUT
(C : ComparisonInstance)
(P : Process)
( $\tau\_ref$  : Prefix C.Sys_ref)
( $\tau\_post$  : Prefix C.RTL_B) : Prop :=
CompatiblePOnProjection C P C.Sigma_DUT C.sigmaDUTProjection  $\tau\_ref$   $\tau\_post$ 
```

-- Projected system compatibility for the outer DUT/testbench-visible alphabet.  
 -- This is the formal target of the Appendix L Note 1 safety branch.

```
def CompatibleSysDUT
(C : ComparisonInstance)
( $\tau\_ref$  : Prefix C.Sys_ref)
( $\tau\_post$  : Prefix C.RTL_B) : Prop :=
CompatibleSysOnProjection C C.Sigma_DUT C.sigmaDUTProjection  $\tau\_ref$   $\tau\_post$ 
```

-- Projected pairwise composition. This is the  $\Sigma\_DUT$  analogue of J.0 and is  
 -- intentionally weaker than J0\_pairwise\_composition: it requires projected  
 -- process compatibility, projected channel/accounting consistency, and projected  
 -- required-order preservation only at C.Sigma\_DUT.

```
theorem J0_pairwise_composition_DUT
(C : ComparisonInstance)
(W : WitnessBundle C)
(hC : C.WellFormed)
(hB : BRules C)
(hR : RRules C)
(hE : ERules C)
(hIC : IssueCommitWF C W)
(hContracts : ModelingContracts C)
( $\tau\_ref$  : Prefix C.Sys_ref)
( $\tau\_post$  : Prefix C.RTL_B)
(hProc :  $\forall P$ , CompatiblePDUT C P  $\tau\_ref$   $\tau\_post$ )
(hChan : ChannelsWellFormedAndMatchedOnProjection
  C C.Sigma_DUT C.sigmaDUTProjection  $\tau\_ref$   $\tau\_post$ ) :
CompatibleSysDUT C  $\tau\_ref$   $\tau\_post$ 
```

J.1:

```
structure JAssumptions
(C : ComparisonInstance)
```

(W : WitnessBundle C) where  
 wellformed : C.WellFormed  
 bRules : BRules C  
 rRules : RRules C  
 eRules : ERules C  
 issueCommitWF : IssueCommitWF C W  
 implementation : ImplementationAssumptions C  
 witnessCompatible : WC C W  
 contracts : ModelingContracts C  
 environment : ReactiveEnvironmentWF C.env  
 -- Elaboration-projection well-formedness consumed by  
 -- RequiredJOrderWF\_from\_concrete\_generators and by the  
 -- J1\_fixed\_RTL\_time\_selected\_witness\_construction proof (hJ.elaborationProjectionWF).  
 -- Trivial (identity projection) for ordinary Sys\_B; supplied by the  
 -- ElaborationPackage piSigma\_\*\_wf projection obligations in Sys\_B^elab cases.  
 elaborationProjectionWF :  
 ∀ side τ, ElaborationProjectionWF C side τ

-- Assumption bundle for the projected  $\Sigma$ \_DUT safety theorem. This intentionally  
 -- omits witnessCompatible : WC C W and replaces it with the weaker  
 -- SigmaDUTSafetyWitnessObligation C W. It is not suitable for the full  
 -- C.Sigma theorem, cutpoint correspondence, or Appendix K liveness/deadlock  
 -- preservation unless those theorems separately import full WC.

structure JSigmaDUTSafetyAssumptions  
 (C : ComparisonInstance)  
 (W : WitnessBundle C) where  
 wellformed : C.WellFormed  
 bRules : BRules C  
 rRules : RRules C  
 eRules : ERules C  
 issueCommitWF : IssueCommitWF C W  
 implementation : ImplementationAssumptions C  
 sigmaDUTSafetyWitness : SigmaDUTSafetyWitnessObligation C W  
 contracts : ModelingContracts C  
 environment : ReactiveEnvironmentWF C.env

-- RTLTimeNextObservableUnit carries the fixed v42 enumeration discipline. It is  
 -- derived from the selected RTL\_B witness prefix after grouping explicitly  
 -- atomic same-cycle transactions into single ObservableUnit objects. It refines  
 -- RTL\_B commit-cycle order and, within one commit cycle, uses only the selected  
 -- witness/replay attribution needed to choose a realizable next unit. This  
 -- predicate is a construction-schedule predicate, not a RequiredJOrder  
 -- generator, and it adds no observable order among units that <\_J leaves  
 -- unordered.

```

def RTLTimeNextObservableUnit
(C : ComparisonInstance)
(W : WitnessBundle C)
( $\tau\_post$  : Prefix C.RTL_B)
(H : Finset (ObservableUnit C.Sigma))
( $\gamma$  : ObservableUnit C.Sigma) : Prop :=
 $\gamma \in \text{ObservableUnitsOf } C .post \tau\_post \wedge$ 
 $\gamma \notin H \wedge$ 
AllRequiredJPredecessorsIn C .post  $\tau\_post$  H  $\gamma \wedge$ 
RefinesSelectedRTLCommitCycleOrder C W  $\tau\_post$  H  $\gamma \wedge$ 
WithinCycleWitnessReplayAttributionCompatible C W  $\tau\_post$  H  $\gamma \wedge$ 
AtomicUnitAlreadyGrouped C .post  $\tau\_post$   $\gamma$ 

```

```

structure RTLTimeIdealChain
(C : ComparisonInstance)
(W : WitnessBundle C)
( $\tau\_post$  : Prefix C.RTL_B) where
N : Nat
idealAt : Nat  $\rightarrow$  Finset (ObservableUnit C.Sigma)
initial_empty :
idealAt 0 =  $\emptyset$ 
terminal_units :
idealAt N = ObservableUnitsOf C .post  $\tau\_post$ 
constant_after_terminal :
 $\forall n,$ 
 $N \leq n \rightarrow$ 
idealAt n = idealAt N
single_unit_step :
 $\forall n,$ 
 $n < N \rightarrow$ 
 $\exists \gamma,$ 
 $\gamma \notin \text{idealAt } n \wedge$ 
RTLTimeNextObservableUnit C W  $\tau\_post$  (idealAt n)  $\gamma \wedge$ 
idealAt (n+1) = idealAt n  $\cup \{\gamma\}$ 
rtl_time_prefix_closed :
 $\forall n \ u \ v,$ 
 $n \leq N \rightarrow$ 
 $v \in \text{idealAt } n \rightarrow$ 
RequiredJOrder C .post  $\tau\_post$  u v  $\rightarrow$ 
u  $\in \text{idealAt } n$ 

```

```

def NextJUnitRealizable
(C : ComparisonInstance)
(W : WitnessBundle C)

```

```

( $\tau_{\text{post}}$  : Prefix C.RTL_B)
( $\tau_{\text{ref\_prefix}}$  : Prefix C.Sys_ref)
( $u$  : ObservableUnit C.Sigma) : Prop :=
SourceOperationFrontierReached C W  $\tau_{\text{ref\_prefix}}$   $u$   $\vee$ 
AlreadyIssuedPendingInReplayState C W  $\tau_{\text{ref\_prefix}}$   $u$ 

-- SelectedJWitness is produced by the Appendix J.1 construction. It is not an
-- unconstrained hypothesis and it is not inferred later from CompatibleSys.
-- The evidence records the fixed RTL-time ideal chain and the replay proof that
-- builds the particular Sys_ref prefix selected for  $\tau_{\text{post}}$  and W.
structure SelectedJWitnessEvidence
(C : ComparisonInstance)
(W : WitnessBundle C)
( $\tau_{\text{post}}$  : Prefix C.RTL_B)
( $\tau_{\text{ref}}$  : Prefix C.Sys_ref) where
chain : RTLTimeIdealChain C W  $\tau_{\text{post}}$ 
terminal_replay :
ReplayPrefixForIdeal C W  $\tau_{\text{post}}$   $\tau_{\text{ref}}$  (chain.idealAt chain.N)
step_realizability :
 $\forall n,$ 
 $n < \text{chain.N} \rightarrow$ 
 $\exists \gamma \tau_{\text{ref\_prefix}},$ 
 $\gamma \notin \text{chain.idealAt } n \wedge$ 
RTLTimeNextObservableUnit C W  $\tau_{\text{post}}$  (chain.idealAt  $n$ )  $\gamma \wedge$ 
chain.idealAt ( $n+1$ ) = chain.idealAt  $n \cup \{\gamma\} \wedge$ 
ReplayPrefixForIdeal C W  $\tau_{\text{post}}$   $\tau_{\text{ref\_prefix}}$  (chain.idealAt  $n$ )  $\wedge$ 
NextJUnitRealizable C W  $\tau_{\text{post}}$   $\tau_{\text{ref\_prefix}}$   $\gamma$ 
exec_ref : Exec_ref C  $\tau_{\text{ref}}$ 
compatible : CompatibleSys C  $\tau_{\text{ref}}$   $\tau_{\text{post}}$ 

def SelectedJWitness
(C : ComparisonInstance)
(W : WitnessBundle C)
( $\tau_{\text{post}}$  : Prefix C.RTL_B)
( $\tau_{\text{ref}}$  : Prefix C.Sys_ref) : Prop :=
Nonempty (SelectedJWitnessEvidence C W  $\tau_{\text{post}}$   $\tau_{\text{ref}}$ )

```

The J.1 induction must follow the fixed RTL-time ideal chain. It must not extend an arbitrary  $<_J$ -minimal observable unit, because v42's proof relies on the fact that the next RTL-time unit is realizable in the replay state: its source operation is either frontier-reached or already issued and pending. Each chain step adds exactly one ObservableUnit  $\gamma$ . If  $\gamma$  is an explicitly atomic rendezvous, SyncChannel/ac\_sync, or R2C fixed-point transaction, it has already been grouped as one ObservableUnit before the chain step is taken. Acyclicity of



RequiredJOrder is still needed to show that each RTL-time prefix is an ideal, but acyclicity alone is not the construction schedule.

theorem J1\_fixed\_RTL\_time\_selected\_witness\_construction

(C : ComparisonInstance)

(W : WitnessBundle C)

(hJ : JAssumptions C W)

( $\tau_{\text{post}}$  : Prefix C.RTL\_B)

(hExec : Exec\_post C  $\tau_{\text{post}}$ ) :

$\exists$  chain : RTLTimeIdealChain C W  $\tau_{\text{post}}$ ,

$\forall$  n,

$n < \text{chain.N} \rightarrow$

$\exists \gamma \tau_{\text{ref\_prefix}},$

$\gamma \notin \text{chain.idealAt } n \wedge$

RTLTimeNextObservableUnit C W  $\tau_{\text{post}}$  (chain.idealAt n)  $\gamma \wedge$

chain.idealAt (n+1) = chain.idealAt n  $\cup \{\gamma\} \wedge$

ReplayPrefixForIdeal C W  $\tau_{\text{post}}$   $\tau_{\text{ref\_prefix}}$  (chain.idealAt n)  $\wedge$

NextJUnitRealizable C W  $\tau_{\text{post}}$   $\tau_{\text{ref\_prefix}}$   $\gamma :=$

by

-- First obtain observable-unit well-formedness for the finite post prefix.

have hOU\_post :

ObservableUnitWF C .post  $\tau_{\text{post}} :=$

ObservableUnitWF\_of\_exec\_post C W hJ  $\tau_{\text{post}}$  hExec

-- RequiredJOrderWF is derived from the concrete v42 generator definitions; it

-- is not an additional JAssumptions field.

have hWF\_post :

RequiredJOrderWF C .post  $\tau_{\text{post}} :=$

RequiredJOrderWF\_from\_concrete\_generators

C W .post  $\tau_{\text{post}}$

hJ.eRules

hJ.issueCommitWF

hOU\_post

hJ.implementation

hJ.elaborationProjectionWF

have hAcyc\_post :

Acyclic (RequiredJOrder C .post  $\tau_{\text{post}}$ ) :=

RequiredJOrder\_acyclic\_on\_observable\_units

C .post  $\tau_{\text{post}}$  hWF\_post

-- Construct the fixed RTL-time enumeration from the selected RTL\_B witness

-- prefix, after grouping explicitly atomic same-cycle transactions as single

-- ObservableUnit objects. The enumeration refines commit-cycle order and uses

-- within-cycle witness/replay attribution only to choose a realizable next unit.  
 -- The enumeration itself adds no observable order.  
 --  
 -- At each  $n < \text{chain.N}$ , add exactly one next unit  $\gamma$ . The replay-realizability  
 -- case analysis is performed against  $\text{chain.idealAt } n$ , so each within-cycle unit  
 -- is checked after all earlier enumerated same-cycle units have already been  
 -- incorporated into the replay prefix.  
 sorry

```

theorem J1_post_to_ref
(C : ComparisonInstance)
(W : WitnessBundle C)
(hJ : JAssumptions C W)
( $\tau_{\text{post}}$  : Prefix C.RTL_B)
(hExec : Exec_post C  $\tau_{\text{post}}$ ) :
 $\exists \tau_{\text{ref}}$ ,
SelectedJWitness C W  $\tau_{\text{post}}$   $\tau_{\text{ref}}$   $\wedge$ 
Exec_ref C  $\tau_{\text{ref}}$   $\wedge$ 
CompatibleSys C  $\tau_{\text{ref}}$   $\tau_{\text{post}}$ 
-- Appendix L Note 1 projected safety theorem.
-- This theorem proves only the  $\Sigma_{\text{DUT}}$ -visible safety compatibility fragment.
-- Its proof follows the J.1 construction, but whenever an  $\varepsilon$ -modeled local timing
-- choice would normally require WC only to fix proof-internal timing, the proof
-- may instead use  $\text{hJ.sigmaDUTSafetyWitness}$  to show that the choice either is
-- covered by ordinary snooping/witness machinery or is certified
--  $\Sigma_{\text{DUT}}$ -safety-inert. Such inert choices can add/remove/relabel no projected
--  $\Sigma_{\text{DUT}}$  unit and can reorder only DUT-causally-independent projected units,
-- which CompatiblePDUT / CompatibleSysDUT already permit to commute.
theorem J1_post_to_ref_sigmaDUT_safety
(C : ComparisonInstance)
(W : WitnessBundle C)
(hJ : JSigmaDUTSafetyAssumptions C W)
( $\tau_{\text{post}}$  : Prefix C.RTL_B)
(hExec : Exec_post C  $\tau_{\text{post}}$ ) :
 $\exists \tau_{\text{ref}}$ ,
Exec_ref C  $\tau_{\text{ref}}$   $\wedge$ 
CompatibleSysDUT C  $\tau_{\text{ref}}$   $\tau_{\text{post}}$ 

```

Do not use  $\text{J1\_post\_to\_ref\_sigmaDUT\_safety}$  as a substitute for  $\text{J1\_post\_to\_ref}$ .  
 The full theorem still requires  $\text{JAssumptions } C \ W$ , including  $\text{WC } C \ W$ , and concludes  
 $\text{CompatibleSys}$  over the proof-internal alphabet  $C.\text{Sigma}$ . The projected theorem  
 is only the Appendix L Note 1 safety branch at the outer  $C.\text{Sigma}_{\text{DUT}}$  boundary.

Because hJ includes the standing RTL\_B B1 field, RTLConsistentSourcePrefix may be read in the main theorem stack as matching the selected RTL\_B execution prefix and, equivalently by B1, a prefix of the unique RTL\_B  $\Sigma$ -trace for the fixed stimulus. The restriction remains a source-side restriction: arbitrary Sys\_ref executions are not forced to have RTL\_B matches merely because B1 holds on the post side.

Restricted reverse:

```
theorem J1_ref_to_post_restricted
(C : ComparisonInstance)
(W : WitnessBundle C)
(hJ : JAssumptions C W)
( $\tau_{ref}$  : Prefix C.Sys_ref)
(hExec : Exec_ref C  $\tau_{ref}$ )
(hConsistent : RTLConsistentSourcePrefix C W  $\tau_{ref}$ ) :
 $\exists \tau_{post}$ ,
Exec_post C  $\tau_{post} \wedge$ 
CompatibleSys C  $\tau_{ref} \tau_{post}$ 
```

Cutpoint abstraction and correspondence:

```
-- J.Abs is an RTL_B cutpoint-to-source-slice abstraction relation, not a
-- universal bit-level projection from arbitrary RTL registers to source states.
-- It is parameterized by the selected Appendix I/L witness and by the chosen
-- Sys_ref abstract boundaries.
```

```
def JAbs
(C : ComparisonInstance)
(W : WitnessBundle C)
( $\rho_{post}$  : State)
( $\sigma_{ref}$  : State) : Prop :=
ObservablePrefixAgreement C W  $\rho_{post} \sigma_{ref} \wedge$ 
PerProcessSourceFrontierSliceAgreement C W  $\rho_{post} \sigma_{ref} \wedge$ 
SourceVisibleValueSliceAgreement C W  $\rho_{post} \sigma_{ref} \wedge$ 
AbstractChannelStateSliceAgreement C W  $\rho_{post} \sigma_{ref} \wedge$ 
RendezvousEndpointRequestSliceAgreement C W  $\rho_{post} \sigma_{ref} \wedge$ 
BlockingRequestAttributionPendingSliceAgreement C W  $\rho_{post} \sigma_{ref} \wedge$ 
EnablementSliceAgreement C W  $\rho_{post} \sigma_{ref}$ 
```

```
-- J.AbsConf is a compliance obligation for the selected RTL_B implementation,
-- abstract boundary, witness, and normalized cutpoints. It is required whenever
-- Corollary J.1 is used to transfer Pending_P, BoundaryReq, rendezvous endpoint
-- request, ChannelEnabled/ChannelDisabled, Front_P, or WFG facts. The selected
```

```

--  $\tau_{\text{ref}}$  used here is exactly the one marked by  $J1_{\text{post\_to\_ref}}$  through
-- SelectedJWitness; J.AbsConf is not applied to an arbitrary compatible prefix.
structure JAbsConf
(C : ComparisonInstance)
(W : WitnessBundle C) where
holds_for_selected_triple :
 $\forall \tau_{\text{post}} \tau_{\text{ref}} \rho_{\text{post}} \sigma_{\text{ref}},$ 
SelectedJWitness C W  $\tau_{\text{post}} \tau_{\text{ref}} \rightarrow$ 
TerminalCutpoint  $\tau_{\text{post}} \rho_{\text{post}} \rightarrow$ 
TerminalCutpoint  $\tau_{\text{ref}} \sigma_{\text{ref}} \rightarrow$ 
EpsilonNormalizedPair C  $\tau_{\text{ref}} \tau_{\text{post}} \rightarrow$ 
JAbs C W  $\rho_{\text{post}} \sigma_{\text{ref}}$ 
-- "Preserved" here is preservation modulo exactly the changes already
-- accounted for by Issue/Commit linkage, E4 channel-realization semantics, and
-- the explicit boundary / sync-boundary rules. It must not be read as saying
-- that normalization can arbitrarily change request, pending, or attribution
-- facts while still satisfying J.AbsConf.
normalization_preserves_abs_slice :
 $\forall \tau_{\text{post}} \rho_0 \rho_{\text{post}},$ 
EpsilonNormalizationSuffix C .post  $\tau_{\text{post}} \rho_0 \rho_{\text{post}} \rightarrow$ 
SigmaHistoryUnchanged C .post  $\rho_0 \rho_{\text{post}} \wedge$ 
RequestPendingAttributionSlicePreservedModuloIssueCommitE4Boundary
C W  $\rho_0 \rho_{\text{post}}$ 

-- CutpointCorrespondence is the K/J bridge relation between corresponding
-- well-formed  $\varepsilon$ -quiescent cutpoints. The post-side cutpoint is named  $\rho_{\text{post}}$  to
-- emphasize that it is an RTL_B microstate related to  $\sigma_{\text{ref}}$  by J.Abs, not a raw
-- source-state cutpoint. J.Abs is the single source of truth for the transferred
-- source slices. Legacy fields such as NextFront agreement, buffered-channel
-- state agreement, raw rendezvous endpoint request agreement, and pending /
-- enablement agreement should be derived by projection lemmas from hCorr.j_abs,
-- rather than duplicated as independent fields that could drift.
structure CutpointCorrespondence
(C : ComparisonInstance)
(W : WitnessBundle C)
( $\sigma_{\text{ref}}$  : State)
( $\rho_{\text{post}}$  : State) where
j_abs : JAbs C W  $\rho_{\text{post}} \sigma_{\text{ref}}$ 

```

```

ref_quiescent : EpsQuiescent  $\sigma_{\text{ref}}$ 
post_quiescent : EpsQuiescent  $\rho_{\text{post}}$ 
ref_wellformed : WellFormedCutpoint C .ref  $\sigma_{\text{ref}}$ 
post_wellformed : WellFormedCutpoint C .post  $\rho_{\text{post}}$ 
direct_input_anchor_agreement : DirectInputAnchorAgreement C W  $\sigma_{\text{ref}}$   $\rho_{\text{post}}$ 
array_mapping_agreement : ArrayMappingAgreement C W  $\sigma_{\text{ref}}$   $\rho_{\text{post}}$ 
r2c_fixed_point_agreement : R2CFixedPointAgreement C W  $\sigma_{\text{ref}}$   $\rho_{\text{post}}$ 

```

```

theorem cutpoint_nextfront_agreement
(C : ComparisonInstance)
(W : WitnessBundle C)
( $\sigma_{\text{ref}}$   $\rho_{\text{post}}$  : State)
(hCorr : CutpointCorrespondence C W  $\sigma_{\text{ref}}$   $\rho_{\text{post}}$ ) :
NextFrontAgreement C W  $\sigma_{\text{ref}}$   $\rho_{\text{post}}$  :=
by
-- Projection from hCorr.j_abs.PerProcessSourceFrontierSliceAgreement.
sorry

```

```

theorem cutpoint_rendezvous_endpoint_request_agreement
(C : ComparisonInstance)
(W : WitnessBundle C)
( $\sigma_{\text{ref}}$   $\rho_{\text{post}}$  : State)
(hCorr : CutpointCorrespondence C W  $\sigma_{\text{ref}}$   $\rho_{\text{post}}$ ) :
RendezvousEndpointRequestSliceAgreement C W  $\rho_{\text{post}}$   $\sigma_{\text{ref}}$  :=
by
-- Projection from hCorr.j_abs; no separate raw field is maintained.
sorry

```

```

theorem cutpoint_pending_enablement_agreement
(C : ComparisonInstance)
(W : WitnessBundle C)
( $\sigma_{\text{ref}}$   $\rho_{\text{post}}$  : State)
(hCorr : CutpointCorrespondence C W  $\sigma_{\text{ref}}$   $\rho_{\text{post}}$ ) :
BlockingRequestAttributionPendingSliceAgreement C W  $\rho_{\text{post}}$   $\sigma_{\text{ref}}$   $\wedge$ 
EnablementSliceAgreement C W  $\rho_{\text{post}}$   $\sigma_{\text{ref}}$  :=
by
-- Projection from hCorr.j_abs; this is the fact consumed by K0/K2.
sorry

```

```

theorem J1_cutpoint_correspondence
(C : ComparisonInstance)
(W : WitnessBundle C)
(hJ : JAssumptions C W)
(hAbs : JAbsConf C W)
( $\tau_{\text{ref}}$  : Prefix C.Sys_ref)
( $\tau_{\text{post}}$  : Prefix C.RTL_B)
(hCompat : CompatibleSys C  $\tau_{\text{ref}}$   $\tau_{\text{post}}$ )
(hNorm : EpsilonNormalizedPair C  $\tau_{\text{ref}}$   $\tau_{\text{post}}$ )
(hSel : SelectedJWitness C W  $\tau_{\text{post}}$   $\tau_{\text{ref}}$ ) :
 $\exists \sigma_{\text{ref}} \rho_{\text{post}}$ ,
TerminalCutpoint  $\tau_{\text{ref}}$   $\sigma_{\text{ref}}$   $\wedge$ 
TerminalCutpoint  $\tau_{\text{post}}$   $\rho_{\text{post}}$   $\wedge$ 
CutpointCorrespondence C W  $\sigma_{\text{ref}}$   $\rho_{\text{post}}$ 

```

The proof of `J1_cutpoint_correspondence` must populate `CutpointCorrespondence.j_abs` from `hAbs.holds_for_selected_triple` for the same selected  $(\tau_{\text{post}}, \tau_{\text{ref}}, \rho_{\text{post}}, \sigma_{\text{ref}}, W)$  produced by `J1_post_to_ref`. It must populate `CutpointCorrespondence.ref_quiescent` and `CutpointCorrespondence.post_quiescent` from `hNorm` / the terminal normalized cutpoint construction, and must populate `CutpointCorrespondence.ref_wellformed` and `CutpointCorrespondence.post_wellformed` from the terminal cutpoint construction and the side-specific execution well-formedness facts supplied by `hJ.wellformed` and the compatible `Sys_ref` / `RTL_B` executions. Thus any later `K` theorem that receives `hCorr` also receives both the  $\varepsilon$ -quiescence, `WellFormedCutpoint`, and `J.Abs` source-slice facts needed to apply `K0` and to transfer request/pending/endpoint-request facts. Deterministic replay supplies the source-side replay slice; `J.AbsConf` supplies the `RTL_B`-side request, pending, attribution, endpoint-request, and enablement slice. The latter is not inferred from the `P`-local  $\Sigma$ -label sequence alone. Legacy agreement facts, when needed, should be obtained through projection lemmas from `hCorr.j_abs` rather than through duplicate `CutpointCorrespondence` fields.

## 18. Appendix K WFG and liveness

Define WFG over the selected proof-internal `Sys_ref` / `RTL_B` boundaries, not over the outer `Sigma_DUT` projection. In ordinary cases this distinction may collapse to the identity projection. In elaborated cases, however, `Sys_ref` = `Sys_Belab(E)`, and the WFG must see the explicit staged/bundle/join/epoch-gate abstract channels introduced by `E` even when `sigmaDUTProjection` /  `$\Pi_{\text{elab}}$`  hides

their events from the outer DUT/testbench-visible alphabet.

```
def FrontierItemChannelSide
(C : ComparisonInstance)
(side : Side)
(x : FrontierItem)
(c : Channel) : Prop :=
 $\exists$  q,
MsgOfFrontierItemSide C side x = some q  $\wedge$ 
q.channel = c
```

```
def ComplementOwnerItemSide
(C : ComparisonInstance)
(side : Side)
(x : FrontierItem)
(Q : Process) : Prop :=
 $\exists$  q,
MsgOfFrontierItemSide C side x = some q  $\wedge$ 
ComplementOwner C side q = Q
```

```
def InternalChannelOfFrontierItem
(C : ComparisonInstance)
(side : Side)
(x : FrontierItem) : Prop :=
 $\exists$  q,
MsgOfFrontierItemSide C side x = some q  $\wedge$ 
InternalChannel C side q.channel
```

```
def WFGProofInternalBoundaryItem
(C : ComparisonInstance)
(side : Side)
(x : FrontierItem) : Prop :=
 $\exists$  q,
MsgOfFrontierItemSide C side x = some q  $\wedge$ 
q.channel  $\in$  ProofInternalChannels C side  $\wedge$ 
ChannelRepresentedInSigma C.Sigma q.channel  $\wedge$ 
 $\neg$  ChannelVisibilityDefinedOnlyBySigmaDUT C q.channel
```

```

def WFG
(C : ComparisonInstance)
(side : Side)
( $\sigma$  : State) : Digraph Process :=
{ edge := WFGEdge C side  $\sigma$  }

```

```

def WFGEdge
(C : ComparisonInstance)
(side : Side)
( $\sigma$  : State)
(P Q : Process) : Prop :=
 $\exists x,$ 
 $x \in \text{Front } C \text{ side } P \ \sigma \wedge$ 
BlockingMessageFrontierItem  $x \wedge$ 
BoundaryReqItemSide C side  $x \ \sigma \wedge$ 
ChannelDisabledItemSide C side  $x \ \sigma \wedge$ 
ComplementOwnerItemSide C side  $x \ Q \wedge$ 
InternalChannelOfFrontierItem C side  $x \wedge$ 
WFGProofInternalBoundaryItem C side  $x$ 

```

```

theorem WFG_ignores_sigmaDUTProjection
(C : ComparisonInstance)
(side : Side)
( $\sigma$  : State)
(P Q : Process) :
WFGEdge C side  $\sigma \ P \ Q \leftrightarrow$ 
WFGEdgeComputedOverProofInternalSigma C side  $\sigma \ P \ Q :=$ 
by
-- WFG, Front, BoundaryReqItemSide, ChannelEnabledItemSide, and
-- ChannelDisabledItemSide are interpreted over C.Sigma and the chosen
-- Sys_ref / RTL_B abstract channels through the owning dynamic message
-- instance of each message-operation frontier item. sigmaDUTProjection is used
-- only after the proof-internal comparison when projecting to the outer
-- DUT/testbench-visible observation boundary.
sorry

```

Edges are generated from  $\text{Front}_P(\sigma)$ , not arbitrary non-frontier asserted requests. This intentionally makes the WFG sparser than a graph over all asserted



boundary requests: only pending blocking frontier items generate wait-for edges. For message-operation frontier items, `BoundaryReqItemSide` and `ChannelDisabledItemSide` recover the owning `DynMsgInstance` and then apply the operation-attributive `BoundaryReqSide` / `ChannelDisabledSide` predicates. Failed non-blocking polls and already-asserted non-frontier requests may affect state or occupancy through other rules, but they do not themselves create WFG dependency edges.

The Lean development must include the Appendix K.2.2 / Appendix R.1 structural classification theorem as an explicit milestone.  $\text{Front\_P}(\sigma)$  is not an independent primitive for liveness proofs: at  $\varepsilon$ -quiescent cutpoints it is derived from the blocking message-operation slice of `NextFront_P` together with the operation-attributive `BoundaryReq` predicate.

-- Message-operation frontier items whose owning dynamic message instance lies  
 -- in the  $Q_\Omega$  slice for process P. This is a predicate on `FrontierItem`, not a  
 -- direct membership test in  $Q_\Omega$ . The document's  $Q_\Omega, P$  is a set of  
 -- `DynMsgInstance`, while  $\Omega\_msg, P$  / `NextFront_P` are `FrontierItem`-level objects.  
 -- Therefore this predicate relates the two universes through  
 -- `MsgInstanceOfFrontierItem` / `MsgOfFrontierItemSide` rather than by literal set  
 -- intersection.

```
def MessageFrontierItemInQOmega
(C : ComparisonInstance)
(side : Side)
(P : Process)
(x : FrontierItem)
( $\sigma$  : State) : Prop :=
 $\exists q$ ,
MsgOfFrontierItemSide C side x = some q  $\wedge$ 
q  $\in$  QOmegaOfProcess C side P  $\sigma$ 
```

```
def BlockingMsgNextFront
(C : ComparisonInstance)
(side : Side)
(P : Process)
( $\sigma$  : State) : Set FrontierItem :=
{ x |
x  $\in$  NextFront C side P  $\sigma$   $\wedge$ 
```

$\text{MessageFrontierItemInQOmega } C \text{ side } P \times \sigma \wedge$   
 $\text{BlockingMessageFrontierItem } x \wedge$   
 $\text{MessageKindOfFrontierItem } C \text{ side } x \in \{\text{MsgKind.Push, MsgKind.Pop}\}$

This preserves the document's distinction between  $\Omega_{\text{msg},P} / \text{NextFront}_P$ , whose elements are `FrontierItem` objects, and  $Q_{\Omega,P}$ , whose elements are `DynMsgInstance` objects. A message frontier item belongs to the blocking `NextFront` slice only when its owning dynamic message instance is defined and lies in  $Q_{\Omega,P}$ .

```

def FrontFromNext
(C : ComparisonInstance)
(side : Side)
(P : Process)
( $\sigma$  : State) : Set FrontierItem :=
{ x |
x  $\in$  BlockingMsgNextFront C side P  $\sigma$   $\wedge$ 
BoundaryReqItemSide C side x  $\sigma$  }
  
```

theorem K0\_front\_nextfront\_classification

```

(C : ComparisonInstance)
(W : WitnessBundle C)
(side : Side)
(P : Process)
( $\sigma$  : State)
(hK : KAssumptions C W)
(hQ : EpsQuiescent  $\sigma$ )
(hWF : WellFormedCutpoint C side  $\sigma$ ) :
Front C side P  $\sigma$  =
FrontFromNext C side P  $\sigma$  :=
by
  
```

```

-- Appendix K.2.2 / Appendix R.1 proof target, aligned with the v42 expanded
-- proof of Lemma K.0. The forward direction must prove that a pending blocking
-- frontier item is not merely minimal among pending blocking items, but minimal
-- in the full  $\text{Remain}_P(\sigma) / \text{NextFront}_P$  sense. The proof therefore splits on
-- any hypothetical  $y <_P x$ :
-- 1.  $y$  cannot be a preceding synchronization item, because the current
-- interval definition plus R3/E1 puts it in  $\text{Done}_P(\sigma)$ ;
  
```

-- 2.  $y$  cannot be a preceding anchored sequential signal item, because R2/E2  
 -- realizes it at  $\text{pred\_S/succ\_S}$ , nor a preceding combinational item whose  
 -- emitted R2C unit has already been realized;  
 -- 3.  $y$  cannot be an earlier same-interval message item, because R4/E3 gives  
 -- issue order; if it committed it is Done, if it is pending blocking it  
 -- contradicts Front-minimality, if it is successful NB it is Done, and if  
 -- it is failed NB it is not an  $\Omega\_P$  frontier item;  
 -- 4.  $y$  cannot come from an earlier source interval, because the closing sync  
 -- has committed and R3/E5 forces earlier blocking  $\text{pref\_S}$  work to have  
 -- committed no later than that sync.  
 -- The reverse direction uses Appendix C stable attribution: an asserted  
 -- BoundaryReq for  $x \in \text{BlockingMsgNextFront}$  is attributed to its owning  $q\_x$  and  
 -- cannot be retagged to another same-endpoint operation without an intervening  
 -- commit. Younger already-asserted requests may exist as non-frontier protocol  
 -- state under Appendix B no-withdrawal/drain-only rules, but they do not enter  
 --  $\text{Front\_P}(\sigma)$  while an earlier pending blocker is frontier-minimal.  
 sorry

theorem K0\_front\_correspondence\_from\_cutpoint

(C : ComparisonInstance)

(W : WitnessBundle C)

(P : Process)

( $\sigma\_ref$   $\rho\_post$  : State)

(hK : KAssumptions C W)

(hCorr : CutpointCorrespondence C W  $\sigma\_ref$   $\rho\_post$ ) :

Front C .ref P  $\sigma\_ref$  = Front C .post P  $\rho\_post$  :=

by

-- Apply K0\_front\_nextfront\_classification on both sides using hK, with  
 -- hCorr.ref\_quiescent / hCorr.post\_quiescent supplying the  $\varepsilon$ -quiescence  
 -- hypotheses and hCorr.ref\_wellformed / hCorr.post\_wellformed supplying the  
 -- WellFormedCutpoint hypotheses. Then use projection lemmas from hCorr.j\_abs,  
 -- such as cutpoint\_nextfront\_agreement and  
 -- cutpoint\_pending\_enablement\_agreement, for NextFront agreement and  
 -- operation-attributive BoundaryReqItemSide agreement on the common  
 -- BlockingMsgNextFront slice. The document-numbered A1-A11 facts are accessed  
 -- through hK.doc; the Appendix C Issue/Commit attribution premise is accessed  
 -- through hK.issue\_commit\_wf.

sorry

This theorem is intentionally stated for `BlockingMsgNextFront_P`, not for the broader `MsgNextFront_P`. Successful non-blocking message-operation frontier items may contribute committed `Push_c / Pop_c`  $\Sigma$ -labels, and failed non-blocking polls are  $\varepsilon$ -steps, but neither creates a pending blocking WFG cause.

```
def ProcessStalledAtFront
(C : ComparisonInstance)
(side : Side)
(P : Process)
( $\sigma$  : State) : Prop :=
Front C side P  $\sigma \neq \emptyset \wedge$ 
 $\forall x \in \text{Front } C \text{ side } P \sigma,$ 
BlockingMessageFrontierItem x  $\wedge$ 
BoundaryReqItemSide C side x  $\sigma \wedge$ 
ChannelDisabledItemSide C side x  $\sigma$ 
```

-- Conjunction form, unified with `FrontierBlocked` from the Section 13 progress  
-- machinery: every frontier item of `P` is a blocking message-passing operation  
-- whose `BoundaryReq` is asserted and whose channel is disabled. The earlier  
-- implication-chained body (`BlockingMessageFrontierItem x  $\rightarrow$  BoundaryReqItemSide`  
-- `C side x  $\sigma \rightarrow$  ChannelDisabledItemSide C side x  $\sigma$` ) was strictly weaker — it was  
-- vacuously satisfied by a process whose frontier contained a non-message item  
-- or a not-yet-asserted request, which is not “blocked on message passing” in  
-- the v42 K.1 sense and must not be eligible for membership in a deadlocked  
-- set `D`. `ProcessStalledAtFront C side P  $\sigma \leftrightarrow$  FrontierBlocked C side P  $\sigma$`  should  
-- be proved (or the two definitions literally shared) so the WFG construction  
-- and the Section 13 progress lemmas consume the same predicate.

```
def WFGClosed
(C : ComparisonInstance)
(side : Side)
(D : Finset Process)
( $\sigma$  : State) : Prop :=
 $\forall P Q,$ 
 $P \in D \rightarrow$ 
WFGEdge C side  $\sigma$  P Q  $\rightarrow$ 
 $Q \in D$ 
```

```

def WFGTerminalSCC
(C : ComparisonInstance)
(side : Side)
(D : Finset Process)
( $\sigma$  : State) : Prop :=
D.Nonempty  $\wedge$ 
StronglyConnectedSubgraph (WFG C side  $\sigma$ ) D  $\wedge$ 
WFGClosed C side D  $\sigma$ 

-- A drain escape is not necessarily a WFG edge leaving D. It is an enabled
-- non-frontier blocking transfer/drain that would discharge already-issued
-- downstream work for some process in D. Such a drain may be intra-process or
-- may not appear as a graph edge at all, because WFG edges are generated only
-- from current minimal blocking frontier items Front_P( $\sigma$ ).
--
-- Important: "non-frontier" here means "not in the current minimal pending
-- blocking frontier Front_P( $\sigma$ )." It does not mean that the operation instance is
-- absent from the broader dynamic/message-frontier universe  $\Omega\_msg,P / Q\_Q,P$ .
-- Conversely, it should not be characterized as a member of the current
-- BlockingMsgNextFront_P( $\sigma$ ) slice together with BoundaryReq, because K.0 proves
-- that, at  $\varepsilon$ -quiescent well-formed cutpoints, those items are exactly Front_P( $\sigma$ ).
--
-- Downstream automatic-flush work is already-issued downstream protocol state:
-- its owning message-operation frontier item may still be part of the broader
-- witness universe, may still have BoundaryReq asserted, and may commit as
-- drain-only progress, but its item is not currently minimal in Front_P( $\sigma$ ).
def OwnsCurrentFrontierMsgInstance
(C : ComparisonInstance)
(side : Side)
(P : Process)
( $\sigma$  : State)
(q : DynMsgInstance) : Prop :=
 $\exists x$ ,
 $x \in \text{Front } C \text{ side } P \ \sigma \wedge$ 
MsgOfFrontierItemSide C side  $x = \text{some } q$ 

def NonFrontierBlockingDrainCandidate
(C : ComparisonInstance)
(side : Side)
(P : Process)
( $\sigma$  : State)
(q : DynMsgInstance) : Prop :=
 $\exists x$ ,

```

$\text{DrainableDownstreamRequest } C \text{ side } P \sigma \wedge$   
 $\text{MsgOfFrontierItemSide } C \text{ side } x = \text{some } q \wedge$   
 $x \notin \text{Front } C \text{ side } P \sigma$

$\text{def EnabledNonFrontierBlockingDrainForProcess}$   
 $(C : \text{ComparisonInstance})$   
 $(\text{side} : \text{Side})$   
 $(P : \text{Process})$   
 $(\sigma : \text{State}) : \text{Prop} :=$   
 $\exists q,$   
 $\text{OwnsMessageEndpoint } P \ q \wedge$   
 $\text{Blocking } q \wedge$   
 $\text{NonFrontierBlockingDrainCandidate } C \text{ side } P \sigma \ q \wedge$   
 $\text{BoundaryReqPersistentDomainSide } C \text{ side } q \sigma \wedge$   
 $\text{BoundaryReqSide } C \text{ side } q \sigma \wedge$   
 $\text{ChannelEnabledSide } C \text{ side } q \sigma$   
 $\text{def NoEnabledNonFrontierBlockingDrain}$   
 $(C : \text{ComparisonInstance})$   
 $(\text{side} : \text{Side})$   
 $(P : \text{Process})$   
 $(\sigma : \text{State}) : \text{Prop} :=$   
 $\neg \text{EnabledNonFrontierBlockingDrainForProcess } C \text{ side } P \sigma$

$\text{def DrainClosed}$   
 $(C : \text{ComparisonInstance})$   
 $(\text{side} : \text{Side})$   
 $(D : \text{Finset Process})$   
 $(\sigma : \text{State}) : \text{Prop} :=$   
 $\forall P,$   
 $P \in D \rightarrow$   
 $\text{NoEnabledNonFrontierBlockingDrain } C \text{ side } P \sigma$

The K.1 SCC contradiction should use this definition. The closed SCC supplies the internal frontier WFG cycle; DrainClosed separately rules out currently enabled non-frontier blocking drains for processes in D. In v42, however, DrainClosed is not misread as saying that outside-D or environment activity can never make such a non-frontier request ChannelEnabled later. The K.1 Exec bridge therefore uses fair drain-normalization: if a D-owned already-issued non-frontier request later becomes continuously serviceable under B2/B3, allow it to commit as drain-only progress, and use the finite-drain argument to reach a later  $\varepsilon$ -quiescent drain-closed cutpoint when needed.

```

def HasOutgoingWFGEdgeInD
(C : ComparisonInstance)
(side : Side)
(D : Finset Process)
(P : Process)
( $\sigma$  : State) : Prop :=
 $\exists Q,$ 
 $Q \in D \wedge$ 
WFGEdge C side  $\sigma$  P Q

def ExistsClosedDrainClosedWFGCycle
(C : ComparisonInstance)
(side : Side)
( $\sigma$  : State) : Prop :=
 $\exists D,$ 
WFGTerminalSCC C side D  $\sigma \wedge$ 
DrainClosed C side D  $\sigma \wedge$ 
( $\forall P,$ 
 $P \in D \rightarrow$ 
ProcessStalledAtFront C side P  $\sigma$ )  $\wedge$ 
( $\forall P,$ 
 $P \in D \rightarrow$ 
HasOutgoingWFGEdgeInD C side D P  $\sigma$ )

def InternalMPDeadlockAt
(C : ComparisonInstance)
(side : Side)
( $\sigma$  : State) : Prop :=
EpsQuiescent  $\sigma \wedge$ 
ExistsClosedDrainClosedWFGCycle C side  $\sigma$ 

```

The K.1 SCC contradiction should use this definition. The closed SCC supplies the cycle of internal message-passing dependencies; DrainClosed rules out escape by downstream draining; ProcessStalledAtFront ensures every process in D is stalled on its current blocking frontier, not merely carrying a non-frontier asserted request. The additional HasOutgoingWFGEdgeInD condition is the non-vacuity / internal-cause condition: every process in the candidate deadlock set must have at least one WFG edge, generated from an internal proof-visible message-passing frontier item, to another process in the same set. This excludes the spurious

singleton case where a process is merely waiting on an external environment or testbench channel and therefore has no internal WFG dependency edge. For genuine multi-process closed SCCs this condition is normally automatic, but stating it explicitly prevents a length-zero singleton SCC from being misclassified as an internal message-passing deadlock.

Terminated-process convention (v42 K.1). The scheduling-rules document additionally fixes a terminated-process convention: a process that is blocked solely on channels whose unique complementary endpoint process has terminated is not counted as participating in an internal message-passing deadlock; such a configuration is a protocol/termination modeling issue, not a K.1 deadlock. In this strategy that convention is not an extra assumption — it is a corollary of `finish_process_idle_padding` together with the closed-SCC deadlock definition above, and Lean should prove it explicitly:

theorem terminated\_peer\_waiting\_not\_internal\_MP\_deadlock

(C : ComparisonInstance)

(side : Side)

(P : Process)

( $\sigma$  : State)

(hP2P : PointToPointChannels C)

(hWait : BlockedOnlyOnTerminatedPeerChannels C side P  $\sigma$ ) :

$\forall D, P \in D \rightarrow \neg \text{ClosedDrainClosedWFGSet } C \text{ side } D \sigma$

Proof sketch. By A6 point-to-point channels, each channel in P’s blocking frontier has a unique complementary endpoint process Q, and hWait says every such Q has terminated (its `FINISH_Q` anchor has occurred in the current reset epoch). By `finish_process_idle_padding`, a terminated Q is thereafter padded with idle  $\varepsilon$ -steps only: its Front is empty, it asserts no `BoundaryReq`, it commits no message operations, and it generates no new blocking WFG cause. Hence Q fails `ProcessStalledAtFront` — its frontier is empty, violating the first conjunct of the unified conjunction form — so Q is not “blocked on message passing” and cannot belong to any candidate deadlock set D. But every WFG edge out of P generated by P’s blocking frontier targets such a terminated Q. Therefore, for any D containing P, either P’s outgoing WFG edges all leave D, violating the `HasOutgoingWFGEdgeInD / WFGClosed` condition, or D would have to contain a terminated Q, violating the stalled-process condition for Q. Either way no closed drain-closed WFG set containing P exists, so P cannot witness `InternalMPDeadlockAt`.

Scope warning (v42 K.1). This convention is a scoping rule for the deadlock definition only. It does not verify, and must not be read as verifying, intended termination behavior: explicit EOF, done, shutdown, or cancellation protocols, and any design intent that a consumer should observe a producer’s termination, must be modeled explicitly as message-passing



or signal protocol in Sys\_ref and RTL\_B and proved through the ordinary I/J/K obligations, or verified outside Appendix K. A process left waiting forever on a terminated peer is excluded from the K.1 deadlock alarm precisely because the document treats it as a protocol-level modeling defect, not as an internal message-passing deadlock; the exclusion must therefore never be used to argue that such a configuration is acceptable.

K assumptions should preserve the document's A1–A11 as a bundle, and add Lean-side auxiliary contracts without pretending they are doc-numbered assumptions. The document-numbered DockAssumptions structure below is therefore intentionally aligned one-for-one with Appendix K.2.1 A1–A11 of the scheduling-rules document. Lean-only prerequisites needed to instantiate imported Appendix C/J/Q/R2C lemmas are placed in KAssumptions outside the doc field.

The v45 Appendix K update changes A5 substantially. A5 is no longer a scheduler-robust liveness premise over all source scheduling choices. It is the conservative serial-blocking source-liveness premise: for A5 only, a process may not issue a later same-interval blocking Push/Pop until every earlier same-interval blocking Push/Pop of that process has committed. The ordinary Appendix J witness, Corollary J.1 cutpoint correspondence, and RTL\_B implementation may still contain overlapped / liberal issue. Lemma K.2b is the bridge from such a witness-relative overlapped Sys\_ref deadlock to a serial-source deadlock that contradicts A5.

The reference execution classes must be scoped symmetrically. Serial-blocking is a restriction of the ordinary liberal source semantics, not a replacement for the R4/E3/E5/no-reverse issue discipline. Therefore every Exec\_ref\_serial execution must also satisfy the same issue-order well-formedness required of an Exec\_ref\_liberal execution. This can be achieved either by folding the issue-order well-formedness into Exec\_ref / RefAdmissible itself, or, as written below, by making Exec\_ref\_serial explicitly include Exec\_ref\_liberal plus the no-eager restriction. This avoids making A5 quantify over ill-formed serial runs that the source semantics never produces.

```
def A11_Kepsilon_WFGInert
(C : ComparisonInstance) : Prop :=
  ∀ side P τ t0 t σ0 σ t,
  StateAtCycleInPrefix C side τ t0 σ0 →
```

StateAtCycleInPrefix C side  $\tau$  t  $\sigma$  t  $\rightarrow$   
 $t_0 \leq t \rightarrow$   
 InternalOrEpsilonOnlyBetween C side  $\tau$  t<sub>0</sub> t  $\rightarrow$   
 FrontierBlocked C side P  $\sigma_0 \rightarrow$   
 FrontierRemainsBlockedOver C side P  $\tau$  t<sub>0</sub> t  $\rightarrow$   
 WFGInertForProcess C side P  $\sigma_0 \sigma$  t  $\wedge$   
 NoNewBlockingFrontierCause C side P  $\sigma_0 \sigma$

def BarrierWellFormed  
 (C : ComparisonInstance)  
 (W : WitnessBundle C) : Prop :=  
 $\forall$  b,  
 SyncChannelInstance C b  $\rightarrow$   
 DesignatedParticipantsReachSameBarrierCount C W b  $\wedge$   
 NoPermanentBarrierBypass C W b

inductive SourceBlockingIssuePolicy  
 | liberal  
 | serial

-- Completion of an earlier blocking operation before a later same-interval issue.  
 -- In the document-aligned serial policy, serial means the earlier blocking  
 -- operation has committed before the later operation issues; no same-cycle  
 -- commit-then-issue exception is assumed. If a future document revision adds a  
 -- precise intra-bucket commit-before-issue convention, represent it as a  
 -- separately named optional extension theorem, not as part of A5 by default.

def EarlierBlockingOpCompletedBeforeLaterIssue  
 (C : ComparisonInstance)  
 (side : Side)  
 ( $\tau$  : ExecutionBucketed C.Sigma)  
 (qEarlier qLater : DynMsgInstance)  
 (tIssueLater : Cycle) : Prop :=  
 CommitsBefore C side  $\tau$  qEarlier tIssueLater

-- qLater is an eager issue if it is issued before some source-earlier blocking  
 -- Push/Pop in the same dynamic interval has completed in the sense above. Eager  
 -- issue is permitted by the liberal witness policy, but forbidden by the serial A5  
 -- policy.

```

def EagerBlockingIssueAt
(C : ComparisonInstance)
(side : Side)
( $\tau$  : ExecutionBucketed C.Sigma)
(qLater qEarlier : DynMsgInstance)
(t : Cycle) : Prop :=
Blocking qLater  $\wedge$  Blocking qEarlier  $\wedge$ 
SameOwningProcess qEarlier qLater  $\wedge$ 
SameDynamicSourceInterval C side  $\tau$  qEarlier qLater  $\wedge$ 
SourceOrderBefore C side  $\tau$  qEarlier qLater  $\wedge$ 
IssueAtCycle C side  $\tau$  qLater t  $\wedge$ 
 $\neg$  EarlierBlockingOpCompletedBeforeLaterIssue C side  $\tau$  qEarlier qLater t

```

```

structure EagerBlockingIssueEvent
(C : ComparisonInstance)
(side : Side)
( $\rho$  : ExecutionBucketed C.Sigma) where
qLater : DynMsgInstance
qEarlier : DynMsgInstance
tIssue : Cycle
eager : EagerBlockingIssueAt C side  $\rho$  qLater qEarlier tIssue

```

-- The concrete implementation can choose any finite representation of the eager  
-- issue events of the finite prefix  $\rho$ , provided membership is equivalent to the  
-- EagerBlockingIssueAt witness above. This Finset is the measure used by K.2b.

```

def EagerBlockingIssueEvents
(C : ComparisonInstance)
(side : Side)
( $\rho$  : ExecutionBucketed C.Sigma) : Finset (EagerBlockingIssueEvent C side  $\rho$ ) :=
AllEagerBlockingIssueEventsInFinitePrefix C side  $\rho$ 

```

-- IssuedBlockingEndpointSourceOrderStream is ordered by source/endpoint order,  
-- not by issue time. Deferring one eager issue may move an event later in time,  
-- so an issue-time-ordered stream would not be prefix-monotone. The prefix theorem  
-- below is about the source-order endpoint stream: deferral may withhold a suffix  
-- of the endpoint stream, but it never drops an earlier endpoint operation while  
-- retaining a later one.

```

def IssuedBlockingEndpointSourceOrderStream

```

```

(C : ComparisonInstance)
(side : Side)
( $\rho$  : ExecutionBucketed C.Sigma)
(c : Channel)
(dir : EndpointDirection) : List DynMsgInstance :=
BlockingEndpointInstancesIssuedInPrefixOrderedBySourceEndpointOrder C side  $\rho$  c dir

```

```

-- DefersOneEagerIssue is the run-surgery relation used by the one-step reduction.
-- It identifies a particular eager issue event in  $\rho$ , removes or defers that
-- operation to its serial slot, preserves all retained earlier committed transfers
-- and their values, and changes no endpoint stream except by suppressing a
-- source-order suffix on explicit WFG-relevant point-to-point channels. The relation
-- must also record that any new terminal cutpoint  $\sigma'$  used after the edit is reached
-- by finite  $\varepsilon$ /drain normalization from the edited prefix.

```

```

def DefersOneEagerIssue
(C : ComparisonInstance)
(side : Side)
( $\rho$   $\rho'$  : ExecutionBucketed C.Sigma) : Prop :=
 $\exists$  ev : EagerBlockingIssueEvent C side  $\rho$ ,
RemovesOrDelaysExactlyEagerIssue C side  $\rho$   $\rho'$  ev  $\wedge$ 
RetainedCommittedTransferPrefix C side  $\rho'$   $\rho$   $\wedge$ 
 $\forall$  c dir,
ChannelCanContributeWFGEdge C c  $\rightarrow$ 
IsPrefix
(IssuedBlockingEndpointSourceOrderStream C side  $\rho'$  c dir)
(IssuedBlockingEndpointSourceOrderStream C side  $\rho$  c dir)

```

```

def ConservativeSerialBlockingRun
(C : ComparisonInstance)
(side : Side)
( $\tau$  : ExecutionBucketed C.Sigma) : Prop :=
 $\forall$  qLater qEarlier t,
EagerBlockingIssueAt C side  $\tau$  qLater qEarlier t  $\rightarrow$  False

```

```

def LiberalBlockingRun
(C : ComparisonInstance)
(side : Side)
( $\tau$  : ExecutionBucketed C.Sigma) : Prop :=

```

IssueOrderConformsToR4E3E5 C side  $\tau \wedge$   
 NoReverseDistinctInterfaceIssue C side  $\tau$

-- Prefix predicate for all finite prefixes of an infinite execution. The policy  
 -- restrictions must hold on every finite prefix, equivalently on the infinite run;  
 -- they are not checked only on one arbitrarily chosen prefix.

```
def FinitePrefixOfInfinite
(C : ComparisonInstance)
(side : Side)
( $\tau$  : InfiniteExecutionBucketed C.Sigma)
( $\rho$  : ExecutionBucketed C.Sigma) : Prop :=
 $\rho$  IsFinitePrefixOfInfiniteExecution C side  $\tau$ 
```

```
def LiberalBlockingInfiniteRun
(C : ComparisonInstance)
(side : Side)
( $\tau$  : InfiniteExecutionBucketed C.Sigma) : Prop :=
 $\forall \rho$ ,
FinitePrefixOfInfinite C side  $\tau \rho \rightarrow$ 
LiberalBlockingRun C side  $\rho$ 
```

```
def ConservativeSerialBlockingInfiniteRun
(C : ComparisonInstance)
(side : Side)
( $\tau$  : InfiniteExecutionBucketed C.Sigma) : Prop :=
 $\forall \rho$ ,
FinitePrefixOfInfinite C side  $\tau \rho \rightarrow$ 
ConservativeSerialBlockingRun C side  $\rho$ 
```

```
def Exec_ref_liberal
(C : ComparisonInstance)
( $\tau_{ref}$  : InfiniteExecutionBucketed C.Sigma) : Prop :=
Exec_ref C  $\tau_{ref} \wedge$  LiberalBlockingInfiniteRun C .ref  $\tau_{ref}$ 
```

```
def Exec_ref_serial
(C : ComparisonInstance)
( $\tau_{ref}$  : InfiniteExecutionBucketed C.Sigma) : Prop :=
Exec_ref C  $\tau_{ref} \wedge$ 
```

```

LiberalBlockingInfiniteRun C .ref  $\tau\_ref \wedge$ 
ConservativeSerialBlockingInfiniteRun C .ref  $\tau\_ref$ 

```

```

theorem exec_ref_serial_implies_liberal
(C : ComparisonInstance)
( $\tau\_ref$  : InfiniteExecutionBucketed C.Sigma) :
Exec_ref_serial C  $\tau\_ref \rightarrow$  Exec_ref_liberal C  $\tau\_ref$  :=
by
intro h
exact ⟨h.1, h.2.1⟩

```

```

def InternalMPDeadlockFreeOnSerialExecRef
(C : ComparisonInstance) : Prop :=
 $\forall \tau\_ref \sigma\_ref,$ 
Exec_ref_serial C  $\tau\_ref \rightarrow$ 
TerminalCutpoint  $\tau\_ref \sigma\_ref \rightarrow$ 
 $\neg$  InternalMPDeadlockAt C .ref  $\sigma\_ref$ 

```

```

-- Finite measure used by Lemma K.2b. For a finite prefix  $\rho$ , EagerIssueCount is
-- the number of dynamic blocking issue events in  $\rho$  that satisfy
-- EagerBlockingIssueAt. The count is finite because  $\rho$  is finite.

```

```

def EagerIssueCount
(C : ComparisonInstance)
(side : Side)
( $\rho$  : ExecutionBucketed C.Sigma) : Nat :=
Finset.card (EagerBlockingIssueEvents C side  $\rho$ )

```

```

theorem zero_eager_iff_serial_blocking_prefix
(C : ComparisonInstance)
(side : Side)
( $\rho$  : ExecutionBucketed C.Sigma) :
EagerIssueCount C side  $\rho$  = 0  $\leftrightarrow$ 
ConservativeSerialBlockingRun C side  $\rho$  :=
by
-- Unfold EagerIssueCount, EagerBlockingIssueEvents, and
-- ConservativeSerialBlockingRun.
sorry

```

structure DockAssumptions

(C : ComparisonInstance)

(W : WitnessBundle C) where

-- A1. R-rules and B-rules.

-- This is the document-numbered Appendix K assumption that the selected

-- Sys\_ref / RTL\_B comparison satisfies the semantic rule bundles needed by K.

-- E-rules are still required by the imported Appendix J compatibility theorem,

-- but they are not renamed as a document-numbered K assumption here.

A1\_R\_and\_B\_rule\_conformance :

RRules C  $\wedge$  BRules C

-- A2. Finite Pipeline Depth (FPD) / B4 finite  $\epsilon$ -stutter.

-- This exposes the main-document D\_P-style finite  $\epsilon$ -stutter premise.

A2\_finite\_pipeline\_depth\_B4 :

FiniteEpsilonDepth C

-- A3. Weak Fairness (WF) / B2.

A3\_weak\_fairness\_B2 :

$\forall$  side  $\tau$ ,

InfiniteExecutionOfSide C side  $\tau \rightarrow$

WeakFairness C side  $\tau$

-- A4. Channel Progress / B3.

A4\_channel\_progress\_B3 :

$\forall$  side  $\tau$  c,

InfiniteExecutionOfSide C side  $\tau \rightarrow$

P\_push C side  $\tau$  c  $\wedge$  P\_pop C side  $\tau$  c

-- A5. Source-level liveness under conservative serial-blocking source semantics.

-- This is the user-facing Appendix K source liveness premise. It quantifies over

-- serial executions that are also ordinary liberal/R4/E3/E5/no-reverse reference

-- executions; the no-eager condition is an additional restriction, not a license

-- to consider ill-formed reversed-order runs.

A5\_source\_liveness :

InternalMPDeadlockFreeOnSerialExecRef C

-- A6. Point-to-point channels.

-- For Appendix K/K.2b this must hold at every explicit abstract channel where

-- ChannelEnabled / ChannelDisabled and WFG edges are evaluated, including any  
 -- introduced Sys\_B^elab staged, bundle, join, or epoch-gate boundary channel.

A6\_point\_to\_point\_channels :

PointToPointChannels C

-- A7. Determinism / witness selection / cutpoint abstraction.

-- RTL\_B determinism is the B1 part used by Appendix K; WC records the selected  
 -- Sys\_ref witness supplied by the Appendix I/J/L witness machinery; J.AbsConf  
 -- records that the same selected RTL\_B cutpoint, Sys\_ref cutpoint, and witness  
 -- discharge the source-slice abstraction used by Corollary J.1 and Appendix K;  
 -- and the WFG evaluation boundaries used by those edges must be point-to-point.

A7\_determinism\_witness\_selection\_cutpoint\_abstraction :

PostDeterministicSigmaTrace C  $\wedge$

WC C W  $\wedge$

JAbsConf C W  $\wedge$

( $\forall c$ , ChannelCanContributeWFGEdge C c  $\rightarrow$  PointToPointAtWFGEvaluationBoundary C c)

-- A8. Barrier well-formedness (SyncChannels).

-- Barrier stalls caused by mismatched or conditional participation are outside  
 -- Appendix K's internal message-passing deadlock analysis.

A8\_barrier\_well\_formedness :

BarrierWellFormed C W

-- A9. Single-transfer-per-cycle endpoint constraint.

A9\_single\_transfer\_per\_cycle :

SingleTransferPerCycleEndpoints C

-- A10. Reset-empty FIFO assumption.

A10\_reset\_empty\_fifo :

ResetEmptyFifoAssumption C

-- A11. WFG-inert  $\varepsilon$  /  $K\varepsilon$ .

A11\_Kepsilon\_WFG\_inert :

A11\_Kepsilon\_WFGInert C

def KElaborationPrerequisite

(C : ComparisonInstance) : Prop :=

$\forall c$ ,



```

ChannelCanContributeWFGEdge C c →
ChosenAbstractBoundaryOfSysRef C c ∧
PointToPointAtWFGEvaluationBoundary C c ∧
NoHiddenSyncBoundaryGatingAtOrdinaryBufferedBoundary C c ∧
(RequiresSyncBoundaryWithholding C c →
C.refChoice.kind = ReferenceKind.elaborated)

```

Lean-side additions:

```

structure KAssumptions
(C : ComparisonInstance)
(W : WitnessBundle C) where
doc : DockKAssumptions C W

```

```

-- Imported Appendix J compatibility uses E1–E5. Keep this as a Lean-side
-- theorem prerequisite rather than relabeling it as one of Appendix K's
-- document-numbered A1–A11 assumptions.
e_rules_for_imported_J :
ERules C

```

```

-- Appendix C Issue/Commit well-formedness is needed by K.0/K.1-style
-- frontier and request-attribution lemmas, but it is not one of the
-- document-numbered Appendix K A1–A11 assumptions. In particular, the current
-- document's A6 is the point-to-point channel assumption, not Issue/Commit WF.
issue_commit_wf :
IssueCommitWF C W

```

```

-- The K-stack invokes Appendix J / Corollary J.1 to obtain corresponding
-- Sys_ref and RTL_B cutpoints from a post-side deadlock cutpoint. Therefore the
-- non-K-numbered prerequisites required by JAssumptions must also be available
-- to the K theorem instance. These are Lean-side import prerequisites, not
-- additional document-numbered Appendix K assumptions.
wellformed_for_imported_J :
C.WellFormed
implementation_for_imported_J :
ImplementationAssumptions C
environment_for_imported_J :
ReactiveEnvironmentWF C.env
elaboration_projection_wf_for_imported_J :

```

$\forall$  side  $\tau$ , ElaborationProjectionWF C side  $\tau$

contracts : ModelingContracts C

r2c? : Option (R2C\_WF C)

elaboration\_prerequisite :

KElaborationPrerequisite C

-- Derived constructor used whenever a K proof imports Appendix J's

-- post-to-reference or cutpoint-correspondence theorem.

def KAssumptions.toJAssumptions

(C : ComparisonInstance)

(W : WitnessBundle C)

(hK : KAssumptions C W) : JAssumptions C W :=

{

wellformed := hK.wellformed\_for\_imported\_J

bRules := hK.doc.A1\_R\_and\_B\_rule\_conformance.2

rRules := hK.doc.A1\_R\_and\_B\_rule\_conformance.1

eRules := hK.e\_rules\_for\_imported\_J

issueCommitWF := hK.issue\_commit\_wf

implementation := hK.implementation\_for\_imported\_J

witnessCompatible :=

hK.doc.A7\_determinism\_witness\_selection\_cutpoint\_abstraction.2.1

contracts := hK.contracts

environment := hK.environment\_for\_imported\_J

elaborationProjectionWF := hK.elaboration\_projection\_wf\_for\_imported\_J

}

This avoids fictional A12/A13 labels and also avoids silently renumbering the document's A1–A11 assumptions inside Lean. It also makes the Appendix J import path explicit: K proofs use hK.doc for document-numbered A1–A11 facts, and use KAssumptions.toJAssumptions C W hK when they need J1\_post\_to\_ref. When they need J1\_cutpoint\_correspondence they also pass the J.AbsConf discharge for the same selected cutpoint/witness triple; because A7 now has four conjuncts, that proof term should be projected by a named helper such as:

def KAssumptions.jAbsConf

(C : ComparisonInstance)

(W : WitnessBundle C)

(hK : KAssumptions C W) : JAbsConf C W :=  
hK.doc.A7\_determinism\_witness\_selection\_cutpoint\_abstraction.2.2.1

```
def KAssumptions.pointToPointAtWFGBoundary
(C : ComparisonInstance)
(W : WitnessBundle C)
(hK : KAssumptions C W) :
∀ c, ChannelCanContributeWFGEdge C c → PointToPointAtWFGEvaluationBoundary C c :=
hK.doc.A7_determinism_witness_selection_cutpoint_abstraction.2.2.2
```

Standing B1 import for Appendix K. Since Appendix K uses the same main theorem stack as Appendix J, KAssumptions C W should import the standing RTL\_B B1 field from the B-rule / JAssumptions layer. K.1 may still use B4/B5 only at the points where  $\varepsilon$ -normalization, fair drain-normalization,  $\varepsilon$ -quiescent endpoint reasoning, or uniqueness up to finite  $\varepsilon$ -stutter is needed, but it should not re-state B1 as a merely conditional uniqueness premise.

K lemmas:

```
theorem K0_frontier_blocking_classification
(C : ComparisonInstance)
(W : WitnessBundle C)
(hK : KAssumptions C W)
( $\sigma$  : State)
(hQ : EpsQuiescent  $\sigma$ ) :
...
```

```
theorem K0a_rendezvous_endpoint_request_agreement
(C : ComparisonInstance)
(W : WitnessBundle C)
(hCorr : CutpointCorrespondence C W  $\sigma_{ref}$   $\rho_{post}$ )
(c : Channel)
(hB0 : capacity c = 0) :
Req_prod c  $\sigma_{ref}$  = Req_prod c  $\rho_{post}$   $\wedge$ 
Req_cons c  $\sigma_{ref}$  = Req_cons c  $\rho_{post}$ 
```

K0a is not proved from matched observable history alone. For  $B(c)=0$  rendezvous, ChannelEnabled depends on the raw complementary endpoint request in the same  $\varepsilon$ -quiescent cycle, and a raw request may persist as protocol state while not being the current frontier item. The proof therefore projects the required raw

producer-side and consumer-side endpoint-request equality from `hCorr.j_abs`, using the `J.AbsConf` discharge for the same selected  $(\tau_{\text{post}}, \tau_{\text{ref}}, \rho_{\text{post}}, \sigma_{\text{ref}}, W)$  triple.

-- Definitional unfolding lemma.

-- This is useful in proofs, but it is not the substantive Appendix K.1

-- channel-case characterization.

theorem InternalMPDeadlockAt\_unfold\_at\_quiescent

(C : ComparisonInstance)

(W : WitnessBundle C)

(hK : KAssumptions C W)

( $\sigma$  : State)

(hReach : Reachable C.RTL\_B  $\sigma$ )

(hQ : EpsQuiescent  $\sigma$ ) :

InternalMPDeadlockAt C .post  $\sigma \leftrightarrow$

ExistsClosedDrainClosedWFGCycle C .post  $\sigma :=$

by

-- Unfold InternalMPDeadlockAt and use hQ for the  $\varepsilon$ -quiescence conjunct.

sorry

-- Channel-side cause exposed by Appendix K.1.

-- These cases are derived from ChannelDisabledItemSide plus E4 at the chosen

-- abstract Sys\_ref / RTL\_B boundary. For buffered channels, disabled Push means

-- no space and disabled Pop means no data. For rendezvous channels, disabled

-- Push/Pop means the unique complementary endpoint request is absent in the

-- same cycle.

inductive K1BlockedChannelCause

(C : ComparisonInstance)

(side : Side)

( $\sigma$  : State)

(x : FrontierItem)

(q : DynMsgInstance)

(c : Channel) : Prop

| buffered\_push\_full :

q.channel = c  $\rightarrow$

IsPushKind q.kind  $\rightarrow$

0 < capacity C c  $\rightarrow$

occAt c  $\sigma$ .cycle = capacity C c  $\rightarrow$

$K1BlockedChannelCause\ C\ side\ \sigma \times q\ c$   
 $| buffered\_pop\_empty :$   
 $q.channel = c \rightarrow$   
 $IsPopKind\ q.kind \rightarrow$   
 $0 < capacity\ C\ c \rightarrow$   
 $occAt\ c\ \sigma.cycle = 0 \rightarrow$   
 $K1BlockedChannelCause\ C\ side\ \sigma \times q\ c$   
 $| rendezvous\_push\_missing\_pop :$   
 $q.channel = c \rightarrow$   
 $IsPushKind\ q.kind \rightarrow$   
 $capacity\ C\ c = 0 \rightarrow$   
 $\neg ComplementaryPopRequestAt\ C\ side\ c\ \sigma \rightarrow$   
 $K1BlockedChannelCause\ C\ side\ \sigma \times q\ c$   
 $| rendezvous\_pop\_missing\_push :$   
 $q.channel = c \rightarrow$   
 $IsPopKind\ q.kind \rightarrow$   
 $capacity\ C\ c = 0 \rightarrow$   
 $\neg ComplementaryPushRequestAt\ C\ side\ c\ \sigma \rightarrow$   
 $K1BlockedChannelCause\ C\ side\ \sigma \times q\ c$

-- Substantive Appendix K.1 characterization.  
-- This is the theorem corresponding to the main document's Lemma K.1:  
-- every process in the deadlocked SCC is blocked at its current blocking  
-- message frontier, and every such disabled blocking channel operation has one  
-- of the E4 channel-side causes below.

theorem K1\_deadlock\_characterization  
(C : ComparisonInstance)  
(W : WitnessBundle C)  
(hK : KAssumptions C W)  
( $\sigma$  : State)  
(hReach : Reachable C.RTL\_B  $\sigma$ )  
(hDead : InternalMPDeadlockAt C .post  $\sigma$ ) :  
 $\exists D,$   
WFGTerminalSCC C .post D  $\sigma \wedge$   
DrainClosed C .post D  $\sigma \wedge$   
 $\forall P,$   
 $P \in D \rightarrow$   
ProcessStalledAtFront C .post P  $\sigma \wedge$

$\forall x q c,$   
 $x \in \text{Front } C . \text{post } P \sigma \rightarrow$   
 $\text{BlockingMessageFrontierItem } x \rightarrow$   
 $\text{MsgOfFrontierItemSide } C . \text{post } x = \text{some } q \rightarrow$   
 $q.\text{channel} = c \rightarrow$   
 $\text{BoundaryReqItemSide } C . \text{post } x \sigma \rightarrow$   
 $\text{ChannelDisabledItemSide } C . \text{post } x \sigma \rightarrow$   
 $\text{K1BlockedChannelCause } C . \text{post } \sigma x q c :=$   
 by  
 -- Proof outline:  
 -- 1. Unfold InternalMPDeadlockAt using  
 -- InternalMPDeadlockAt\_unfold\_at\_quiescent to obtain an  $\varepsilon$ -quiescent  
 -- closed drain-closed WFG SCC D.  
 -- 2. Use ProcessStalledAtFront to show each  $P \in D$  is stalled on its current  
 -- blocking message frontier, not on an internal  $\varepsilon$ -step.  
 -- 3. Use MsgOfFrontierItemSide to recover the owning dynamic message instance q.  
 -- 4. Use ChannelDisabledItemSide, hK.doc.A1\_R\_and\_B\_rule\_conformance, and  
 -- hK.e\_rules\_for\_imported\_J / E4 channel semantics at  $\sigma.\text{cycle}$ :  
 -- • disabled buffered Push\_c implies  $\text{occAt } c \sigma.\text{cycle} = \text{capacity } C c$ ;  
 -- • disabled buffered Pop\_c implies  $\text{occAt } c \sigma.\text{cycle} = 0$ ;  
 -- • disabled rendezvous Push\_c implies the complementary Pop request is  
 -- absent in  $\sigma$ ;  
 -- • disabled rendezvous Pop\_c implies the complementary Push request is  
 -- absent in  $\sigma$ .  
 -- 5. Use hK.elaboration\_prerequisite to ensure that c is the chosen explicit  
 -- abstract boundary of the Sys\_ref / RTL\_B comparison after any required  
 -- Sys\_B^elab elaboration. Synchronization-boundary ramp-down / epoch-gating  
 -- is therefore represented at explicit abstract channels, not as a hidden  
 -- exception to E4 on an otherwise ordinary buffered channel.  
 sorry

theorem K2\_dependency\_refinement  
 (C : ComparisonInstance)  
 (W : WitnessBundle C)  
 (hK : KAssumptions C W)  
 ( $\sigma_{\text{ref}} \rho_{\text{post}}$  : State)  
 (hCorr : CutpointCorrespondence C W  $\sigma_{\text{ref}} \rho_{\text{post}}$ ) :

WFG C .ref  $\sigma_{\text{ref}}$  = WFG C .post  $\rho_{\text{post}}$

theorem K2a\_drain\_closure\_transfer

(C : ComparisonInstance)

(W : WitnessBundle C)

(hK : KAssumptions C W)

(D : Finset Process)

( $\sigma_{\text{ref}}$   $\rho_{\text{post}}$  : State)

(hCorr : CutpointCorrespondence C W  $\sigma_{\text{ref}}$   $\rho_{\text{post}}$ )

(hDrainPost : DrainClosed C .post D  $\rho_{\text{post}}$ ) :

DrainClosed C .ref D  $\sigma_{\text{ref}}$

K2\_dependency\_refinement and K2a\_drain\_closure\_transfer intentionally do not take separate hQref / hQpost or hWFref / hWFpost parameters. Their cutpoint quiescence facts are part of hCorr, via hCorr.ref\_quiescent and hCorr.post\_quiescent, and their cutpoint well-formedness facts are part of hCorr, via hCorr.ref\_wellformed and hCorr.post\_wellformed. Thus these theorems can invoke K0\_front\_correspondence\_from\_cutpoint without any extra side hypotheses.

K.2b serial-source dominance. The scheduling-rules v45 proof no longer tries to rebuild a fresh serial replay toward the same frontier. Instead, Lean should formalize K.2b as a well-founded reduction on the number of eager blocking issues in a finite liberal witness prefix. The proof edits an existing finite witness-relative deadlocking run into another finite deadlocking run with one fewer eager issue; a minimal-eager deadlocking run therefore has zero eager issues and is a conservative serial-blocking run.

structure DrainClosedDeadlockingPrefix

(C : ComparisonInstance)

(side : Side)

( $\rho$  : ExecutionBucketed C.Sigma)

( $\sigma$  : State) where

terminal : TerminalCutpointOfPrefix C side  $\rho$   $\sigma$

quiescent : EpsQuiescent  $\sigma$

deadlocked : InternalMPDeadlockAt C side  $\sigma$

drain\_closed\_component :

$\exists D, \text{WFGTerminalSCC } C \text{ side } D \ \sigma \wedge \text{DrainClosed } C \text{ side } D \ \sigma$

```

-- Prefix-suppression is not arbitrary deletion. It may remove or defer only a
-- suffix/order-ideal of each endpoint's issued blocking-operation stream. This
-- prefix property is the reason occupancy domination is true; naive "fewer
-- Pushes and fewer Pops" is not monotone.
def EndpointIssuedBlockingStreamPrefix
(C : ComparisonInstance)
(side : Side)
(p' p : ExecutionBucketed C.Sigma)
(c : Channel)
(dir : EndpointDirection) : Prop :=
IsPrefix
(IssuedBlockingEndpointSourceOrderStream C side p' c dir)
(IssuedBlockingEndpointSourceOrderStream C side p c dir)

theorem K2b_endpoint_prefix_suppression
(C : ComparisonInstance)
(W : WitnessBundle C)
(hK : KAssumptions C W)
(p p' : ExecutionBucketed C.Sigma)
(hDef : DefersOneEagerIssue C .ref p p') :
∀ c dir,
ChannelCanContributeWFGEEdge C c →
EndpointIssuedBlockingStreamPrefix C .ref p' p c dir :=
by
-- Use A6 point-to-point ownership, same-endpoint source order, E3 no-reverse
-- issue/commit order, and the construction of DefersOneEagerIssue. The stream is
-- ordered by source/endpoint order, not issue time; otherwise delaying q would
-- invalidate the prefix statement.
sorry

```

```

theorem K2b_joint_value_control_consistency
(C : ComparisonInstance)
(W : WitnessBundle C)
(hK : KAssumptions C W)
(p p' : ExecutionBucketed C.Sigma)
(hPref : ∀ c dir, EndpointIssuedBlockingStreamPrefix C .ref p' p c dir)
(hRetained : RetainedCommittedTransferPrefix C .ref p' p) :

```



JointSourceReplayAgreementUpToRetainedCommits C W  $\rho'$   $\rho$  :=

by

-- This is a simultaneous source/dataflow theorem, not merely a per-process  
-- deterministic replay lemma. Fixed stimulus, fixed NB-CF / witness-selected NB  
-- outcomes, matched retained Push payloads, and matched retained Pop-return  
-- values imply the same operation identities, guards, payloads, and branches up  
-- to the retained commit prefix.

sorry

theorem K2b\_occupancy\_disabledness\_domination

(C : ComparisonInstance)

(W : WitnessBundle C)

(hK : KAssumptions C W)

( $\rho$   $\rho'$  : ExecutionBucketed C.Sigma)

( $\sigma$   $\sigma'$  : State)

(x : FrontierItem)

(hPrefix :  $\forall c \text{ dir},$

ChannelCanContributeWFGEdge C c  $\rightarrow$

EndpointIssuedBlockingStreamPrefix C .ref  $\rho'$   $\rho$  c dir)

(hFront : CorrespondingBlockedFrontierItem C .ref  $\rho$   $\rho'$   $\sigma$   $\sigma'$  x)

(hDisabled : ChannelDisabledItemSide C .ref x  $\sigma$ ) :

ChannelDisabledItemSide C .ref x  $\sigma'$  :=

by

-- Evaluate on the explicit abstract channel where ChannelDisabled is defined.  
-- For a blocked Pop, endpoint-prefix suppression preserves the relevant empty  
-- condition; for a blocked Push, it preserves the relevant full condition; for  
-- rendezvous, suppressing eager issue cannot create the missing unique  
-- complementary endpoint request. Sys\_B^elab cases lift only when the introduced  
-- boundary channel is point-to-point at the WFG evaluation boundary.

sorry

-- This is the monotonicity fact that prevents the eager-reduction proof from  
-- asserting, without proof, that serialization cannot dissolve the deadlock. When  
-- the deferred eager operation commits in the liberal run, excising that early  
-- commit may change channel occupancy. Point-to-point ownership guarantees that  
-- the only party affected through that channel is the unique complementary  
-- endpoint. Therefore the edit cannot unblock a third process or create an  
-- outgoing escape edge from the terminal WFG component; it can only remove

-- progress/help or keep the same blockedness.  
theorem K2b\_no\_third\_party\_enable\_under\_eager\_deferral  
(C : ComparisonInstance)  
(W : WitnessBundle C)  
(hK : KAssumptions C W)  
( $\rho \rho'$  : ExecutionBucketed C.Sigma)  
( $\sigma \sigma'$  : State)  
(D : Finset Process)  
(hDef : DefersOneEagerIssue C .ref  $\rho \rho'$ )  
(hPrefix :  $\forall c \text{ dir}$ ,  
ChannelCanContributeWFGEdge C c  $\rightarrow$   
EndpointIssuedBlockingStreamPrefix C .ref  $\rho' \rho c \text{ dir}$ )  
(hPTP :  $\forall c$ , ChannelCanContributeWFGEdge C c  $\rightarrow$   
PointToPointAtWFGEvaluationBoundary C c)  
(hSCC : WFGTerminalSCC C .ref D  $\sigma$ )  
(hDrain : DrainClosed C .ref D  $\sigma$ )  
(hDisabled :  $\forall P, P \in D \rightarrow \text{FrontierBlocked C .ref P } \sigma$ ) :  
NoThirdPartyEnablementOrOutgoingEscape C .ref  $\rho \rho' \sigma \sigma' D :=$   
by  
-- Split on whether the deferred eager operation is still pending at  $\sigma$  or  
-- committed in  $\rho$ . Pending eager requests are non-frontier protocol state and are  
-- inert under  $K\varepsilon$ . If the eager operation committed, point-to-point plus  
-- endpoint-prefix occupancy domination shows that withholding the early commit  
-- cannot enable any process other than the unique complement, and that complement  
-- is precisely the process accounted for by the WFG edge/SCC argument.  
sorry

-- Once disabledness and no-third-party-enable are known, the edited prefix still  
-- has a terminal internal WFG component after finite  $\varepsilon$ /drain normalization. This  
-- is the finite graph step that should not be hidden inside K2b\_one\_step\_eager\_reduction.

theorem K2b\_terminal\_scc\_persists\_after\_eager\_deferral

(C : ComparisonInstance)  
(W : WitnessBundle C)  
(hK : KAssumptions C W)  
( $\rho \rho'$  : ExecutionBucketed C.Sigma)  
( $\sigma \sigma'$  : State)  
(D : Finset Process)  
(hDef : DefersOneEagerIssue C .ref  $\rho \rho'$ )

$(hOld : \text{WFGTerminalSCC } C .\text{ref } D \ \sigma \wedge \text{DrainClosed } C .\text{ref } D \ \sigma)$   
 $(hNoEscape : \text{NoThirdPartyEnablementOrOutgoingEscape } C .\text{ref } \rho \ \rho' \ \sigma \ \sigma' \ D)$   
 $(hDisabled' : \forall P, P \in D \rightarrow \text{FrontierBlocked } C .\text{ref } P \ \sigma')$   
 $(hDrain' : \text{DrainClosed } C .\text{ref } D \ \sigma') :$   
 $\exists D',$   
 $\text{WFGTerminalSCC } C .\text{ref } D' \ \sigma' \wedge$   
 $\text{DrainClosed } C .\text{ref } D' \ \sigma' \wedge$   
 $(\forall P, P \in D' \rightarrow \text{ProcessStalledAtFront } C .\text{ref } P \ \sigma') \wedge$   
 $(\forall P, P \in D' \rightarrow \text{HasOutgoingWFGEdgeInD } C .\text{ref } D' \ P \ \sigma') :=$   
 by  
 -- Follow WFG edges from blocked processes in D at  $\sigma'$ . Point-to-point ensures the  
 -- complementary owner is unique and no third-party arbitration can be newly  
 -- enabled by removing the eager operation. Finiteness of the process set yields a  
 -- nonempty terminal SCC. Drain closure rules out downstream drain escapes.  
 sorry

theorem K2b\_one\_step\_eager\_reduction

$(C : \text{ComparisonInstance})$   
 $(W : \text{WitnessBundle } C)$   
 $(hK : \text{KAssumptions } C \ W)$   
 $(\rho : \text{ExecutionBucketed } C.\text{Sigma})$   
 $(\sigma : \text{State})$   
 $(hRun : \text{LiberalBlockingRun } C .\text{ref } \rho)$   
 $(hDead : \text{DrainClosedDeadlockingPrefix } C .\text{ref } \rho \ \sigma)$   
 $(hEager : 0 < \text{EagerIssueCount } C .\text{ref } \rho) :$   
 $\exists \rho' \ \sigma',$   
 $\text{LiberalBlockingRun } C .\text{ref } \rho' \wedge$   
 $\text{DrainClosedDeadlockingPrefix } C .\text{ref } \rho' \ \sigma' \wedge$   
 $\text{EagerIssueCount } C .\text{ref } \rho' < \text{EagerIssueCount } C .\text{ref } \rho :=$   
 by  
 -- Choose the temporally last eager blocking issue. If it remains pending at the  
 -- deadlock cutpoint, it is non-frontier protocol state and is inert: removing  
 -- its early assertion changes no commit and no WFG source. If it committed,  
 -- defer that issue to its serial slot. Endpoint-prefix suppression and  
 -- occupancy-disabledness domination show that the relevant blocked frontier  
 -- items remain disabled. K2b\_no\_third\_party\_enable\_under\_eager\_deferral rules out  
 -- newly enabled third-party escape from the component, and  
 -- K2b\_terminal\_scc\_persists\_after\_eager\_deferral extracts the drain-closed

-- terminal SCC after finite  $\varepsilon$ /drain normalization.  
sorry

theorem K2b\_minimal\_eager\_deadlock\_is\_serial  
 (C : ComparisonInstance)  
 (W : WitnessBundle C)  
 (hK : KAssumptions C W)  
 (p : ExecutionBucketed C.Sigma)  
 ( $\sigma$  : State)  
 (hRun : LiberalBlockingRun C .ref p)  
 (hDead : DrainClosedDeadlockingPrefix C .ref p  $\sigma$ )  
 (hMin :  $\forall p' \sigma'$ ,  
 LiberalBlockingRun C .ref p'  $\rightarrow$   
 DrainClosedDeadlockingPrefix C .ref p'  $\sigma' \rightarrow$   
 EagerIssueCount C .ref p  $\leq$  EagerIssueCount C .ref p') :  
 ConservativeSerialBlockingRun C .ref p :=  
 by  
 -- If EagerIssueCount is positive, K2b\_one\_step\_eager\_reduction contradicts  
 -- minimality. Then zero\_eager\_iff\_serial\_blocking\_prefix gives seriality.  
 sorry

theorem K2b\_serial\_source\_dominance\_for\_internal\_deadlock  
 (C : ComparisonInstance)  
 (W : WitnessBundle C)  
 (hK : KAssumptions C W)  
 ( $\tau_{\text{ref}}$  : InfiniteExecutionBucketed C.Sigma)  
 ( $\sigma_{\text{ref}}$  : State)  
 (hExec : Exec\_ref\_liberal C  $\tau_{\text{ref}}$ )  
 (hTerm : TerminalCutpoint  $\tau_{\text{ref}}$   $\sigma_{\text{ref}}$ )  
 (hDead : InternalMPDeadlockAt C .ref  $\sigma_{\text{ref}}$ ) :  
 $\exists \tau_{\text{ser}} \sigma_{\text{ser}}$ ,  
 Exec\_ref\_serial C  $\tau_{\text{ser}}$   $\wedge$   
 TerminalCutpoint  $\tau_{\text{ser}} \sigma_{\text{ser}} \wedge$   
 InternalMPDeadlockAt C .ref  $\sigma_{\text{ser}}$  :=  
 by  
 -- Convert the selected witness execution to a finite terminal prefix p reaching  
 --  $\sigma_{\text{ref}}$ . Among all liberal finite deadlocking prefixes reachable under the same  
 -- fixed stimulus, capacities, NB outcome selection, Sys\_ref boundaries, and

-- ordinary issue-order/no-reverse constraints, choose one with minimal  
 -- EagerIssueCount. The one-step eager-reduction lemma shows the minimum is zero.  
 -- A zero-eager prefix is serial-blocking and still liberal/well-formed; extend it  
 -- to an Exec\_ref\_serial execution using the same fairness/progress and  
 -- drain-closure continuation machinery already used in K.1.

sorry

Implementation note. K.2b is now the single proof bridge that replaces the former scheduler-robust A5 quantification. The RTL\_B / Corollary J.1 witness may still land on a Sys\_ref cutpoint with multiple pending requests for a process, but A5 is consumed only after K.2b converts that witness-relative deadlock into a deadlock reachable under the conservative serial-blocking source semantics.

Exec membership for K.1 should distinguish three cases.

First, for a globally fixed-point cutpoint, B6 idle-stutter may be used to build an infinite execution extension:

```
theorem fixedpoint_cutpoint_in_Exec_post
(C : ComparisonInstance)
(W : WitnessBundle C)
(hK : KAssumptions C W)
(ρ_post : State)
(hReach : Reachable C.RTL_B ρ_post)
(hFP : FixedPointState C .post ρ_post)
(hIdle : C.assumptions.B.B6_idleStutterTimeDivergence) :
∃ τ_post,
TerminalCutpoint τ_post ρ_post ∧
Exec_post C τ_post
```

Second, a partial internal-MP deadlock is not justified by B6 idle-stutter alone, because processes outside D or environment endpoints may still have enabled observable actions. The v42 proof supplies the missing bridge by fair drain-normalization:

```
def DrainNormalizationStep
(C : ComparisonInstance)
(side : Side)
(D : Finset Process)
(σ σ' : State) : Prop :=
∃ P q,
P ∈ D ∧
```

OwnsMessageEndpoint  $P \ q \ \wedge$   
 Blocking  $q \ \wedge$   
 NonFrontierBlockingDrainCandidate C side  $P \ \sigma \ q \ \wedge$   
 BoundaryReqSide C side  $q \ \sigma \ \wedge$   
 ChannelEnabledSide C side  $q \ \sigma \ \wedge$   
 CommitDrainOnly C side  $q \ \sigma \ \sigma' \ \wedge$   
 NoNewWorkCrossesDisabledInternalFrontier C side  $D \ \sigma \ \sigma'$

Thus drain-normalization is intentionally stated over already-issued downstream blocking requests, not over the current  $\text{BlockingMsgNextFront}_P(\sigma)$  slice. At  $\varepsilon$ -quiescent well-formed cutpoints, K.0 identifies  $\text{Front}_P(\sigma)$  with the BoundaryReq-enabled blocking NextFront slice; drain-only progress is for already-issued downstream requests outside that current minimal frontier, while NoNewWorkCrossesDisabledInternalFrontier prevents such commits from advancing new source work past the blocked frontier.

```

def DrainNormalizedDeadlockCutpoint
  (C : ComparisonInstance)
  (side : Side)
  (D : Finset Process)
  ( $\sigma_0 \ \sigma' : \text{State}$ ) : Prop :=
  ReachableByFiniteDrainNormalization C side D  $\sigma_0 \ \sigma' \ \wedge$ 
  EpsQuiescent  $\sigma' \ \wedge$ 
  DrainClosed C side D  $\sigma' \ \wedge$ 
   $\forall P,$ 
   $P \in D \rightarrow \text{ProcessStalledAtFront C side P } \sigma'$ 

```

```

theorem deadlock_cutpoint_in_Exec_post_via_fair_drain_normalization
  (C : ComparisonInstance)
  (W : WitnessBundle C)
  (hK : KAssumptions C W)
  ( $\rho_{\text{post}} : \text{State}$ )
  (hReach : Reachable C.RTL_B  $\rho_{\text{post}}$ )
  (hDead : InternalMPDeadlockAt C .post  $\rho_{\text{post}}$ ) :
   $\exists \tau_{\text{post}} \rho_{\text{post}}' D,$ 
  TerminalCutpoint  $\tau_{\text{post}} \rho_{\text{post}}' \ \wedge$ 
  Exec_post C  $\tau_{\text{post}} \ \wedge$ 

```

DrainNormalizedDeadlockCutpoint C .post D  $\rho_{\text{post}}$   $\rho_{\text{post}}$ '  $\wedge$   
 InternalMPDeadlockAt C .post  $\rho_{\text{post}}$ ' :=  
 by  
 -- Use K1\_deadlock\_characterization to obtain the closed SCC D. B3 cannot force  
 -- progress of D-owned internal frontier blockers without contradicting  
 -- disabled-frontier, drain-closure, or WFG closure; rendezvous blockers are  
 -- handled by absence of the same-cycle complementary endpoint request.  
 -- Environment-facing frontier items are not WFG edges and are handled as  
 -- ordinary fair continuation, not as WFG-closure contradictions. For  
 -- already-issued D-owned non-frontier requests, drain-closure only says they are  
 -- not ChannelEnabled at the current cutpoint; if outside-D or environment  
 -- progress later makes one continuously serviceable, allow it to commit as  
 -- drain-only progress. Automatic flush, finite explicit capacities, finite  
 --  $\varepsilon$ -depth, and no new work across the disabled internal frontier imply only  
 -- finitely many such D-owned drain commits before a new  $\varepsilon$ -quiescent  
 -- drain-closed cutpoint is reached, unless a frontier item commits and the  
 -- proof uses the corresponding reachable cutpoint according to the v45 case  
 -- split. Interleave this finite drain-normalization with fair outside-D and  
 -- environment continuation to obtain Exec\_post.  
 sorry

Third, the main K theorem can be exposed in both execution-scoped and reachable forms. The execution-scoped post-side form is convenient internally; the reachable form is what Appendix K states after applying the fair drain-normalization bridge. The source-side liveness premise used to contradict a reference deadlock is now serial, not liberal/witness-relative.

```

def InternalMPDeadlockFreeOnExecPost
  (C : ComparisonInstance) : Prop :=
   $\forall \tau_{\text{post}} \rho_{\text{post}},$ 
  Exec_post C  $\tau_{\text{post}} \rightarrow$ 
  TerminalCutpoint  $\tau_{\text{post}} \rho_{\text{post}} \rightarrow$ 
   $\neg$  InternalMPDeadlockAt C .post  $\rho_{\text{post}}$ 
  
```

```

theorem K1_liveness_preservation_on_exec
  (C : ComparisonInstance)
  (W : WitnessBundle C)
  (hK : KAssumptions C W) :
  
```

InternalMPDeadlockFreeOnSerialExecRef C  $\rightarrow$   
InternalMPDeadlockFreeOnExecPost C

def InternalMPDeadlockFreeReachablePost  
(C : ComparisonInstance) : Prop :=  
 $\forall$  p\_post,  
Reachable C.RTL\_B p\_post  $\rightarrow$   
EpsQuiescent p\_post  $\rightarrow$   
 $\neg$  InternalMPDeadlockAt C .post p\_post

Main K theorem:

theorem K1\_liveness\_preservation  
(C : ComparisonInstance)  
(W : WitnessBundle C)  
(hK : KAssumptions C W) :  
InternalMPDeadlockFreeOnSerialExecRef C  $\rightarrow$   
InternalMPDeadlockFreeReachablePost C

Proof dependency note. Prove the reachable post-side conclusion by contradiction. Assume there is a reachable  $\varepsilon$ -quiescent p\_post with InternalMPDeadlockAt C .post p\_post. Apply deadlock\_cutpoint\_in\_Exec\_post\_via\_fair\_drain\_normalization to obtain an Exec\_post prefix ending at the relevant  $\varepsilon$ -quiescent drain-normalized cutpoint p\_post'. Import Appendix J using

hJ := KAssumptions.toJAssumptions C W hK,

then apply J1\_post\_to\_ref and destruct the result as  $\langle \tau_{\text{ref}}, h_{\text{Sel}}, h_{\text{ExecRef}}, h_{\text{Compat}} \rangle$ . The hExecRef component gives Exec\_ref C  $\tau_{\text{ref}}$ , but K2b requires the selected reference execution to be an Exec\_ref\_liberal execution. Establish this missing conjunct explicitly, either because RefAdmissible / Exec\_ref has been strengthened to include R4/E3/E5/no-reverse issue-order well-formedness, or by a separate lemma such as:

theorem J1\_ref\_execution\_is\_liberal  
(C : ComparisonInstance)  
(W : WitnessBundle C)  
(hK : KAssumptions C W)  
( $\tau_{\text{ref}}$  : InfiniteExecutionBucketed C.Sigma)



(hSel : SelectedJWitness C W  $\tau_{\text{ref}}$ )  
 (hExecRef : Exec\_ref C  $\tau_{\text{ref}}$ ) :  
 Exec\_ref\_liberal C  $\tau_{\text{ref}}$

Let hExecRefLiberal be the result of this lemma. Then the hSel component is the SelectedJWitness marker produced by the fixed RTL-time J.1 construction, so J1\_cutpoint\_correspondence can be applied to the Exec\_post prefix ending at  $\rho_{\text{post}}$ , passing hSel and the J.AbsConf discharge KAssumptions.jAbsConf C W hK for that same selected  $(\tau_{\text{post}}, \tau_{\text{ref}}, \rho_{\text{post}}, \sigma_{\text{ref}}, W)$  triple. The resulting CutpointCorrespondence C W  $\sigma_{\text{ref}} \rho_{\text{post}}$  includes  $\varepsilon$ -quiescence, well-formedness, and hCorr.j\_abs. Use K2\_dependency\_refinement, K2a\_drain\_closure\_transfer, and K0a\_rendezvous\_endpoint\_request\_agreement to transfer the closed drain-closed internal WFG deadlock from  $\rho_{\text{post}}$  to  $\sigma_{\text{ref}}$  on the selected witness-relative Exec\_ref\_liberal execution. This  $\sigma_{\text{ref}}$  may include overlapped pending requests and is not necessarily reachable under A5's conservative serial-blocking source semantics. Apply K2b\_serial\_source\_dominance\_for\_internal\_deadlock to hExecRefLiberal to obtain  $\tau_{\text{ser}}$  and  $\sigma_{\text{ser}}$  with Exec\_ref\_serial C  $\tau_{\text{ser}}$ , TerminalCutpoint  $\tau_{\text{ser}}$   $\sigma_{\text{ser}}$ , and InternalMPDeadlockAt C .ref  $\sigma_{\text{ser}}$ . This contradicts InternalMPDeadlockFreeOnSerialExecRef C, equivalently hK.doc.A5\_source\_liveness.

This is intentionally one-way. Do not state:

InternalMPDeadlockFreeOnSerialExecRef C  $\leftrightarrow$  InternalMPDeadlockFreeReachablePost C

A5 deduplication. The A5\_source\_liveness field should be definitionally (or by a one-line unfolding lemma) the serial execution-scoped predicate

InternalMPDeadlockFreeOnSerialExecRef C, not the witness-relative predicate

InternalMPDeadlockFreeOnExecRef C used in the older v42 strategy. In v45, A5 asserts that the prescribed source reference model is internally message-passing-deadlock-free under the conservative serial-blocking source semantics and A1–A4. The proof of K.1 still first obtains a selected, witness-relative Sys\_ref deadlock from Corollary J.1; that cutpoint may contain overlapped pending requests and therefore may not itself lie in the serial execution class. K2b\_serial\_source\_dominance\_for\_internal\_deadlock is the required bridge: it converts the witness-relative deadlock into a serial-source deadlock before A5 is invoked. With that identification, K1\_liveness\_preservation and K1\_liveness\_preservation\_on\_exec may consume hK.doc.A5\_source\_liveness in place of an explicit InternalMPDeadlockFreeOnSerialExecRef C premise, and the theorem instance then carries no duplicated source-liveness hypothesis. The explicit-premise form shown above is retained purely for presentation, so the dependence of K.1 on the source-side serial

liveness input remains visible in the theorem statement; the two statements must be interderivable by definitional unfolding only.

## 19. Appendix R facade

Appendix R should be formalized last as aliases/wrappers over earlier theorems.

Use Lean names that avoid the doc's R.2 collision:

```
theorem R_lemma1_front_nextfront_classification := ...
theorem R_theorem1_prefix_matching := ...
theorem R_corollary1_cutpoint_correspondence := ...
theorem R_corollary1a_rendezvous_endpoint_request_agreement := ...
theorem R_lemma2_dependency_refinement := K2_dependency_refinement
theorem R_theorem2_closed_cycle_and_one_way_drain_transfer := ...
```

Do not call cutpoint correspondence `R1_cutpoint_correspondence`; in the doc structure, it corresponds to a corollary, not Theorem R.1.

Appendix R cutpoint correspondence should expose the v42 naming: the post-side terminal cutpoint is  $\rho_{\text{post}}$ , the relation  $\sigma_{\text{ref}} \equiv \rho_{\text{post}}$  is the compact notation for the  $J.\text{Abs}/J.\text{AbsConf}$  source-slice correspondence  $\rho_{\text{post}} \Downarrow_{\{ \text{Abs}, J, w \}} \sigma_{\text{ref}}$ , and it is not equality of arbitrary RTL microstate.

## 20. What remains as hypotheses

These remain theorem parameters or design/tool obligations:

1. R-rules R1, R2, R2C, R3, R4, R5.
2. E-rules E1-E5, with the v42 RequiredJOrder restriction made explicit: E2 is alignment only, distinct-interface E3 source order is no-reverse legality rather than a positive strict edge, atomic multi-endpoint operations are single observable units rather than internal order edges, and E5 contributes strict RequiredJOrder only on its `suff_S` side. The `pref_S` side of E5 remains a non-strict issue/commit legality and sync-boundary well-formedness obligation:  $q \in \text{pref\_S}(s)$  gives  $\text{issue}(q) \leq \text{clk}(s)$ , and for committed blocking  $q$  the commit cycle is  $\leq \text{clk}(s)$ , with equality allowed. The sanity package RequiredJOrderWF is not a new external assumption; it should be proved inside Lean as RequiredJOrderWF\_from\_concrete\_generators once the concrete E1-E5, IssueCommitWF, ObservableUnitWF, R2C unit, and elaboration/projection definitions are instantiated.
3. B-rule semantic obligations B1-B6. B1 is the standing RTL\_B

deterministic- $\Sigma$ -trace assumption `PostDeterministicSigmaTrace C` for the fixed comparison instance; it is not a source-side determinism premise and is not only introduced when uniqueness is invoked. B2/B3 are encoded temporally and B5 is limited to the genuine quiescence-closure assumptions, including `bounded_system_quiescence`. B2 must encode the v42 split: `ChannelEnabled` means eligibility, selection is the fairness event, and the selected enabled action commits in the selected cycle. B3 remains the channel-progress obligation that can discharge persistent full/empty/rendezvous-blocked premises before a pending transfer becomes `ChannelEnabled`. The max and least-upper-bound facts for `SystemQuiescenceBound C` are not hypotheses; they are Lean lemmas proved from the `Finset.sup` definition. Appendix K A5 is now the separate serial-source liveness premise `InternalMPDeadlockFreeOnSerialExecRef C`; the witness-relative reference deadlock produced by Appendix J is connected to this premise only by `K2b_serial_source_dominance_for_internal_deadlock`.

4. Well-formedness of the selected comparison instance `C.WellFormed`.

5. Implementation assumptions such as `I.A0`.

6. Appendix C Issue/Commit well-formedness and request-attribution discipline `IssueCommitWF C W`.

7. Reactive-environment well-formedness `ReactiveEnvironmentWF C.env`.

8. Witness compatibility `WC`, derived from Lemma L.2 or Corollary L.3 where applicable, for the full proof-internal `C.Sigma` theorem.

8a. For the Appendix L Note 1  $\Sigma$ -DUT-projected safety theorem only, the weaker `SigmaDUTSafetyWitnessObligation C W` may replace `WC`. This obligation is discharged choice-by-choice: each  $\Sigma$ -DUT-relevant  $\varepsilon$ -modeled local timing choice is either covered by ordinary snooping/witness machinery or belongs to a certified `SigmaDUTSafetyInertInterface`. This does not discharge full `WC` and does not discharge Appendix K WFG-inertness /  $K\varepsilon$  obligations.

8b. `J.AbsConf C W` for the selected `RTL_B` implementation, abstract boundary, witness, and normalized cutpoints. This is a tool/RTL-flow compliance obligation alongside E4 channel-realization conformance and B-faithfulness; it is not proved by deterministic replay alone.

8c. B-faithfulness of each back-annotated capacity `B(c)`, as a named tool/RTL-flow compliance hypothesis `BFaithful C c` per explicit abstract channel: (i) boundary completeness — every  $\Sigma$ -visible transfer of `c` in `RTL_B` crosses the agreed abstract boundary as a committed `Push_c` or `Pop_c`, with no hidden side entrance or exit; (ii) no hidden capacity and no hidden enablement — the realized boundary stores at most `B(c)` uncommitted-to-`Pop` items, and there is no hidden grant/arbitration/backpressure condition that gates a transfer beyond the documented E4 availability predicate unless it is modeled as explicit `Sys_B^elab` boundary structure; (iii)  $\varepsilon$ -opacity of internal movement — internal movement of data within the realized channel between commit points is  $\varepsilon$  and contributes no  $\Sigma$ -labels and no WFG causes; and (iv) an evidence obligation — each `B(c)` is supported by a trusted back-annotation algorithm, a per-channel certificate, a structural RTL check, an interface monitor, or a formal flow, recorded in the side-condition discharge record of Section 1.1. Like 8b, this is an obligation on the tool/RTL flow and the chosen comparison instance; it is not derived by E4, Proposition J.0, or Theorem J.1, all of which

assume it.

9. Direct-input contracts.

10. Array access mapping rules.

11. R2C fixed-point well-formedness  $R2C\_WF\ C$  when combinational-process signal IO is in scope, including selected  $Obs(C\_comb)$  slices, stutter-canonical unit emission, component-local occurrence-ordinal matching, and either a single anchoring domain or a confluence proof for multi-input R2C components; cycle-locked R2C components must be declared explicitly.

12. Automatic-flush / drain-only blocked-frontier discipline for pipelined or overlapped implementations: while a minimal blocking frontier remains fully disabled, no new later source work may issue past that frontier; already asserted blocking requests are not withdrawn before commit; and any observable work by the blocked process during the interval is attributable only to already-issued downstream drainable work.

13. Elaboration package E, when  $Sys\_ref = Sys\_B^{elab}(E)$ .

14. K elaboration prerequisite  $KElaborationPrerequisite\ C$ : every WFG-relevant channel is represented at the chosen abstract  $Sys\_ref / RTL\_B$  comparison boundary; hidden sync-boundary gating is not allowed at ordinary buffered boundaries; and sync-boundary withholding that affects WFG reasoning is represented through explicit  $Sys\_B^{elab}$  channels.

15. Appendix K structural assumptions not already covered above: point-to-point channels, point-to-point WFG evaluation boundaries for any introduced  $Sys\_B^{elab}$  channels,  $SyncChannel/barrier$  well-formedness, single-transfer-per-cycle endpoint constraints, reset-empty FIFO/accounted initial occupancy, and WFG-inert  $\varepsilon / K\varepsilon$ .

16. Source-level internal message-passing deadlock freedom for K.1 under A5's conservative serial-blocking source semantics:  $InternalMPDeadlockFreeOnSerialExecRef\ C$ .

This premise quantifies only over serial executions that are also ordinary liberal/R4/E3/E5/no-reverse reference executions; it is not a scheduler-robust liveness premise over all source scheduling choices and it is not derived from local R-rules or from  $RequiredJOrderWF$ .

$K2b\_serial\_source\_dominance\_for\_internal\_deadlock$

is the bridge from the witness-relative liberal reference deadlock to this serial A5 premise.

17. Fair drain-normalization for K.1  $Exec\_post$  applicability: a reachable partial internal-MP deadlock cutpoint is not extended by pure idle-stutter unless it is a global fixed point; the proof must construct a fair continuation that allows finite D-owned non-frontier drain-only commits when B2/B3 later makes them continuously serviceable, while preserving the disabled internal frontier used for the WFG contradiction.

18. The side-condition discharge record of Section 1.1. This is audit bookkeeping, not an additional semantic assumption: it records, per comparison instance, how each of WC, NB-CF, the Appendix L snooping selector W,  $Sys\_B^{elab}(E)$ ,  $J.AbsConf$ , and  $K\varepsilon$  is discharged (inapplicable, by syntactic restriction, by tool/RTL evidence, or by witness construction), so a completed proof instance carries an explicit map from the v42/v45 side

conditions to their discharges. Every theorem instance claimed against the document should be accompanied by a fully populated `SideConditionDischargeRecord`.

## 21. What should be proved inside Lean

Lean should prove:

- bucketed execution infrastructure;
- finite `ObservableUnitsOf` extraction for prefixes, together with `ObservableUnitWF` grouping/non-splitting facts for rendezvous, `SyncChannel/ac_sync`, stutter-canonical R2C fixed-point signal-observation units, and any other explicitly atomic same-cycle units;
- temporal lemmas for `AlwaysFrom` / `EventuallyFrom`;
- the elementary `Finset.sup` facts for `SystemQuiescenceBound C`: every `MainB4Bound C P` for  $P \in \text{ProcessUniverse } C$  is bounded by `SystemQuiescenceBound C`, and `SystemQuiescenceBound C` is the least such finite upper bound. These are proved from the definition of `SystemQuiescenceBound C`, not assumed as part of B5.
- channel-capacity lemmas for rendezvous and buffered cases;
- `ChannelEnabled-as-eligibility` lemmas, `FairActionSelected/FairActionCommitsAt` lemmas, and E4 safety lemmas showing that committed transfers are legal at the chosen abstract boundary;
- issue existence/uniqueness and derived `IssueTime`;
- `RequiredJOrderWF_from_concrete_generators`, proving that the concrete v42 `RequiredJOrderGenerator` satisfies the E2-alignment-only, distinct-interface no-positive-edge, atomicity-as-grouping, witness-control-not-order, R2C-unit-as-atomic-observation, and elaboration-not-separate-order restrictions;
- `RequiredJOrder / <_J` acyclicity over observable units, separately from the informational `CausalDependency` graph, proved by the v42 per-execution commit-cycle measure `clk( $\gamma$ )` over the fixed legal prefix;
- `B3_ordinary_boundary_from_B2_E4`, discharging the `B3 P_push / P_pop` obligation from B2 plus E4 at ordinary single explicit abstract boundaries (Section 6), so that for `OrdinaryE4Boundary` instances B3 is a theorem rather than an assumption;
- `terminated_peer_waiting_not_internal_MP_deadlock`, the Appendix K terminated-process convention corollary (Section 18), derived from `finish_process_idle_padding`, A6 point-to-point channels, and the closed-SCC deadlock definition;
- K.2b serial-source dominance lemmas: `EagerBlockingIssueEvents / EagerIssueCount`, `zero_eager_iff_serial_blocking_prefix`, source-order endpoint-prefix suppression, `DefersOneEagerIssue` run surgery, joint value/control consistency over retained committed

transfers, occupancy/disabledness domination at point-to-point WFG evaluation boundaries, no-third-party-enable monotonicity, terminal-SCC persistence after eager deferral, one-step eager reduction, minimal-eager deadlock implies serial, and K2b\_serial\_source\_dominance\_for\_internal\_deadlock;

- direct-input late-sampling preservation from primitive environment contracts;
- array sync-boundary commit-order preservation;
- R2C fixed-point observation lemmas from R2C\_WF, including selected observable slices, stutter-canonical unit generation, component-local occurrence-ordinal matching, cycle-locked components, and multi-input anchoring/confluence obligations;
- automatic-flush frontier lemmas;
- Lemma L.2: snooping establishes WC, including the latency-sensitive local-protocol handling premise for cadence-sensitive retry/timeout/poll-arbiter sites;
- Corollary L.3: NB-CF identity witness, with the global/per-boundary contract obligations needed by the updated WC definition;
- Appendix L Note 1 projected-safety branch. These prove only the  $\Sigma$ \_DUT-visible safety fragment and do not produce WC or Appendix K deadlock/liveness obligations. The global SigmaDUTSafetyWitnessObligation C W remains an explicit projected-safety assumption; certified inert-interface choices are discharged locally by L5\_sigmaDUT\_safety\_inert\_interface, while non-inert choices require ordinary snooping/witness coverage:
  - o SigmaDUTSafetyInertInterface,
  - o SigmaDUTSafetyWitnessObligation,
  - o L5\_sigmaDUT\_safety\_inert\_interface,
  - o CompatiblePDUT,
  - o CompatibleSysDUT,
  - o J0\_pairwise\_composition\_DUT, and
  - o J1\_post\_to\_ref\_sigmaDUT\_safety.
- IObs\_congruence\_one\_step over concrete ObsSigma;
- IDR\_deterministic\_replay, proved as an induction over IObs\_congruence\_one\_step;
- Appendix I finite and infinite witness construction, with I.3 using ChannelEnabled  $\rightarrow$  eventual selection  $\rightarrow$  commit rather than treating enablement as immediate commit;
- Proposition J.0;
- RTLTimeNextObservableUnit and RTLTimeIdealChain, including the single-unit fixed RTL-time enumeration of the selected RTL\_B witness prefix and the proof that every prefix idealAt n is down-closed under RequiredJOrder;
- J1\_fixed\_RTL\_time\_selected\_witness\_construction, producing the selected single-unit RTL-time ideal chain and the per-step replay-realizability facts;
- Theorem J.1 post-to-reference;

- restricted reverse theorem for RTL-consistent source prefixes;
- Corollary J.1 cutpoint correspondence over  $J.Abs/J.AbsConf$ , using  $\rho_{post} \Downarrow_{\{Abs,J,w\}}$   $\sigma_{ref}$  rather than an implicit RTL\_B-to-source projection;
- $J1_{ref\_execution\_is\_liberal}$ , or an equivalent unfolding lemma showing that the selected J.1 reference execution satisfies  $Exec_{ref\_liberal}$  before it is passed to K.2b;
- K.0/K.0a/K.1/K.2/K.2a, with K.0 proving full NextFront minimality and K.0a using raw rendezvous endpoint-request equality from  $J.Abs/J.AbsConf$ ;
- $fixedpoint\_cutpoint\_in\_Exec\_post$  as the B6 idle-stutter special case for globally fixed-point post-side cutpoints;
- $deadlock\_cutpoint\_in\_Exec\_post\_via\_fair\_drain\_normalization$  for the v42/v45 K.1 reachable-deadlock proof;
- execution-scoped deadlock-freedom predicates  
InternalMPDeadlockFreeOnSerialExecRef and InternalMPDeadlockFreeOnExecPost;
- Theorem K.1 one-way liveness preservation, exposed both as an execution-scoped theorem and as the reachable post-side theorem obtained using fair drain-normalization, with Appendix K A5 stated as conservative serial-blocking source liveness and the witness-relative reference deadlock bridged to A5 by  
 $K2b\_serial\_source\_dominance\_for\_internal\_deadlock$ ;
- Appendix R aliases.

Optional non-core formalization. The Lean development may additionally prove CausalDependencyGenerator coverage, showing that the informational relation mirrors the scheduling-rules causal-dependency definition: message information flow, signal information flow, synchronization information flow, intra-process causal trace/source order, declared memory-serialization/shared-storage dependence, witness-control dependence, and elaboration information flow. This optional coverage theorem is not a premise of Theorem J.1, Corollary J.1, Appendix K, or Appendix R. Those theorem instances depend on RequiredJOrder /  $\prec_J$ , ObservableUnitWF, E1–E5, witness/cutpoint correspondence, and the relevant progress/liveness assumptions, not on acyclicity or completeness of the broader CausalDependency graph.

## 22. Development milestones

### **Milestone A — Finite sequential, blocking-only, non-elaborated core**

Scope:

Sys\_ref = Sys\_B;

sequential processes only;

blocking Push/Pop only;

no NB calls;

no snooping;  
 no external arrays;  
 no direct-input late sampling;  
 no combinational R2C;  
 finite bucketed prefixes and finite bucketed execution fragments;  
 R1–R5 and E1–E5 infrastructure for the sequential/blocking subset;  
 E4 channel-capacity legality for rendezvous and buffered channels;  
 Issue/Commit existence, uniqueness, and request-attribution discipline;  
 conditional finite extension lemmas stated with explicit enabledness or  
 progress witnesses, not with B2/B3 fairness/progress conclusions;  
 finite-prefix observable-unit extraction and blocking-only pairwise  
 composition.

Targets:

E4\_channel\_capacity\_lemmas\_blocking

IssueCommitWF\_blocking

I3\_conditional\_finite\_extension\_blocking\_given\_enabledness

J0\_pairwise\_composition\_blocking

Milestone A deliberately does not target Theorem J.1, Corollary J.1, or  
 Theorem K.1. Those results are execution-level / temporal results: they quantify  
 over Exec\_post / Exec\_ref, rely on FairProgressSatisfying infinite executions,  
 and use B6 idle-stutter/time-divergence only for globally fixed-point cutpoints.  
 General partial internal-MP deadlock reasoning is execution-scoped: the relevant  
 deadlock cutpoint is already on an Exec\_post / Exec\_ref execution. These results  
 therefore belong after the temporal B2/B3/B6 machinery is available.

### **Milestone B — Temporal B2/B3/B6 blocking proof**

Add: infinite bucketed executions; AlwaysFrom, EventuallyFrom; concrete FairAction;  
 P\_push, P\_pop; WeakFairness and side-parametric B2; side-parametric B3 channel  
 progress; B6 idle-stutter/time-divergence; FairProgressSatisfying; Exec\_post and  
 Exec\_ref; Appendix I finite-to-infinite witness construction for the blocking-only case;  
 fixedpoint\_cutpoint\_in\_Exec\_post for the globally fixed-point B6 idle-stutter special case;  
 execution-scoped definitions Exec\_ref\_liberal, Exec\_ref\_serial,  
 InternalMPDeadlockFreeOnSerialExecRef, and InternalMPDeadlockFreeOnExecPost; K.1  
 liveness preservation using the serial A5 execution-scoped formulation plus K.2b, not a  
 reconstructed Exec\_post membership theorem for arbitrary reachable partial-deadlock  
 states.



Targets:

RequiredJOrderWF\_from\_concrete\_generators\_blocking  
RTLTimeIdealChain\_single\_unit\_blocking  
J1\_fixed\_RTL\_time\_selected\_witness\_construction\_blocking  
J1\_post\_to\_ref\_blocking  
JAbsConf\_blocking  
J1\_ref\_execution\_is\_liberal\_blocking  
J1\_cutpoint\_correspondence\_blocking  
fixedpoint\_cutpoint\_in\_Exec\_post\_blocking  
InternalMPDeadlockFreeOnSerialExecRef\_blocking  
K2b\_endpoint\_prefix\_suppression\_blocking  
K2b\_no\_third\_party\_enable\_blocking  
K2b\_terminal\_scc\_persists\_blocking  
K2b\_serial\_source\_dominance\_blocking  
K1\_liveness\_preservation\_blocking  
B3\_from\_B2\_E4\_ordinary\_boundary\_blocking  
terminated\_peer\_not\_internal\_MP\_deadlock\_blocking

### **Milestone C — Non-blocking under NB-CF**

Add:

Qexec;

failed NB polls as  $\varepsilon$ ;

successful NB transfers as  $\Sigma$ ;

Corollary L.3.

Target:

J1\_post\_to\_ref\_nb\_cf

### **Milestone D — Witness selector, snooping, and Appendix L Note 1 projected safety**

Add:

partial-prefix WitnessSelector;

Lemma L.2;

SigmaDUTSafetyInertInterface;

SigmaDUTSafetyWitnessObligation;

CompatiblePDUT and CompatibleSysDUT;

projected pairwise composition J0\_pairwise\_composition\_DUT;

Corollary L.5 / J1\_post\_to\_ref\_sigmaDUT\_safety for interfaces whose  $\varepsilon$ -modeled local timing choices cannot add/remove/relabel  $\Sigma$ \_DUT-visible units and can only commute DUT-causally-independent projected units.

Targets:

J1\_post\_to\_ref\_with\_snooping

J1\_post\_to\_ref\_sigmaDUT\_safety

### **Milestone E — Direct inputs**

Add:

- DirectInputContract;
- derived lateSamplingPreservesLogicalAnchor;
- bucket placement for logical anchor vs physical late sample.

Target:

J1\_post\_to\_ref\_direct\_inputs

### **Milestone F — External arrays**

Add:

- MemoryEvent;
- ArrayAccessMapping;
- sync-relative array commit rules.

Target:

J1\_post\_to\_ref\_external\_arrays

### **Milestone G — R2C**

Add:

- combModel?;
- LocalCombNetwork;
- CombNetwork;
- CombComposition;
- R2C\_WF;
- selected observable slices Obs(C\_comb);
- R2C fixed-point signal-observation units;
- stutter-canonical unit emission and explicit cycle-locked component mode;
- component-local occurrence-ordinal matching;
- multi-input anchoring-domain / interleaving-confluence obligations;
- per-process combinational component obligations;
- composed fixed-point observation semantics.

Target:

J1\_post\_to\_ref\_with\_R2C\_v42\_units

### **Milestone H — Elaboration**

Add:

- ElaborationPackage Sys\_B Chan\_outer  $\Sigma$ \_DUT;
- $\Pi_{\text{elab}}$ ;
- $\pi_{\Sigma, E}$ ;
- $A_E$ ;
- explicit-channel interpretation of B, occ, BoundaryReq, ChannelEnabled, Front, and WFG.

Targets:

J1\_post\_to\_ref\_elab

K1\_liveness\_preservation\_elab

### **Milestone I — Appendix R facade**

Add compact R theorem wrappers after the full proof stack is available.

## **23. Expected difficulty**

The hardest pieces remain:

1. Bucketed execution and ideal induction.
2. Temporal B2/B3 over infinite executions.
3. Concrete ObsSigma and I.DR.
4. Appendix J finite-ideal construction.
5. Direct-input late-sampling contracts.
6. External array access mapping.
7. R2C fixed-point semantics.
8. Appendix Q elaboration obligations.
9. Appendix K serial-A5 liveness/deadlock preservation, including K.2b eager-count reduction / serial-source dominance and fixed-point-only Exec\_post extension.

A polished full formalization is plausibly in the 18k–30k Lean line range if direct inputs, arrays, R2C, temporal fairness/progress, non-blocking witness selection, snooping, automatic flush, Issue/Commit attribution, A11/K $\epsilon$ , and elaboration are all included. A core sequential/blocking/non-elaborated proof would be substantially smaller, plausibly 6k–10k lines; an intermediate blocking-plus-temporal-WFG proof would likely fall around 10k–16k lines.