

# Catapult Matchlib Memory Architecture Modeling Methodology

Stuart Swan

Platform Architect

Siemens EDA

18 April 2025

## Introduction

This document introduces some common scenarios and useful guidelines for modeling memories when using Matchlib and Catapult HLS.

In pre-HLS models, memories appear as normal C/C++ arrays, and are often mapped to RAMs during HLS. To meet performance and area goals, RAMs and their associated logic must often be carefully constructed. For example, to meet design throughput requirements, RAMs may often need to have multiple banks so that multiple accesses to the banks can occur concurrently.

The discussion in this document assumes that the RAMs used in the examples support a maximum of 1 read and 1 write operation per clock cycle, which is a typical limit. Note that it is possible to use RAMs with more ports in Catapult (e.g. 2r2w), but there is a significant area cost for such RAMs so they are typically avoided unless required.

The Catapult Matchlib memory architecture modeling methodology described in this document can be used in both the Catapult SystemC and C++ flows. The Catapult SystemC flow examples which accompany this document are examples 36\*, 38\*, 40\*, 41\*, and 42\* in the Matchlib examples kit. The Catapult C++ flow examples which accompany this document are examples 37\*, 43\* and 44\* in the Matchlib examples kit.

## Terms Used in this Document

**implicit memory:** an implicit memory appears as a normal C/C++ array in the pre-HLS model and is mapped to a RAM during HLS.

**explicit memory:** an explicit memory appears as an instantiation of a particular RAM module (`sc_module`) within a SystemC design. Explicit memories are supported in Catapult but their usage is now deprecated. Instead, the implicit memory modeling approach should always be used.

**shared memory:** a shared memory is a memory that is directly accessed by more than one process. The processes may be in the DUT and/or within the testbench.

**banked memory:** a banked memory is a group of related memories (banks) that are constructed to operate like a single larger memory. Each bank has its own dedicated RAM ports, so multiple IO requests may proceed concurrently to separate banks. Done properly, banking increases memory IO throughput while adding minimal or no area or performance penalty compared to implementing the same RAM storage without banking.

interleave: interleave is a banked memory organization where the low order bits of the addresses are used to select specific banks. Catapult has the INTERLEAVE directive which can automatically create banked memories with this organization.

block\_size: block\_size is used to specify a banked memory organization where the high order bits of the addresses are used to select specific banks. Catapult has the BLOCK\_SIZE directive which can automatically create banked memories with this organization.

word\_width: Catapult has the WORD\_WIDTH directive which is used to widen the width of the data that is read or written on each RAM IO operation, compared to the width of the array accesses in the pre-HLS model. Sometimes this can increase design throughput.

### Goals of the Matchlib Memory Architecture Modeling Methodology

1. Meet PPA (power, performance, area) goals without any compromises.
2. Discover and resolve all functional and performance issues related to memories in the pre-HLS model, not in the post-HLS model.
3. The pre-HLS model code should be clean and easy to understand.
4. Automatically emit an error message in both pre-HLS and post-HLS RTL simulations if any invalid indexes are used with an array, including checking on each dimension of a multi-dimensional array.

It is important to understand that memories play a key role in most designs, and that functional and performance issues related to memory organization or surrounding logic is a common problem. This makes goal #2 very important, since finding and fixing problems in the post-HLS RTL can be very time-consuming because the RTL is machine-generated.

Goal #4 is important for real-world designs, since invalid indexes may be presented to memories in various design scenarios before the design is fully debugged, and if these errors are not caught automatically they can be very time-consuming to debug.

### Overview of the Matchlib Memory Architecture Models

The Matchlib memory architecture models are carefully designed to provide best possible QOR during HLS and to provide robust error checking.

- In almost all cases these classes are implemented as RAMs/ROMs in the RTL generated by HLS.
- They are also sometimes mapped to register arrays in the RTL.
- These classes are not intended to be used for other purposes (e.g. modeling arrays within transaction payloads).
- These classes provide optimal QOR even for "non-power of 2" memories.

Because the intent is to map to RAMs/ROMs, these classes should not be read or written to in the reset state of processes which access them. With this restriction, Catapult can generate PSL or SVA assertions into the RTL to perform checking for index violations for these models.

All classes work in both the Catapult SystemC flow as well as the C++ flow.

Normal C arrays should not be used to model RAMs/ROMs within newly developed Catapult models, since they do not have the benefits described above.

The Matchlib memory architecture model classes are:

#### `ac_array_1D`

- Simple non-shared 1D array implemented as a RAM/ROM.
- Provides built-in checking for index violations for all array dimensions in both pre-HLS model and in RTL (via PSL or SVA assertions in RTL).
- Usage is same as normal C array, aside from declaration.
- See Matchlib toolkit example 35\*.

#### `ac_shared_array_1D`

- Simple shared 1D array implemented as a RAM/ROM.
- Provides built-in checking for index violations for all array dimensions in both pre-HLS model and in RTL (via PSL or SVA assertions in RTL).
- Usage is same as normal C array, aside from declaration.
- See Matchlib toolkit example 12\*.

#### `ac_bank_array_2D`, `ac_bank_array_3D`, `ac_bank_array_vary`

- Banked memory where each bank is a separate RAM instance.
- During HLS, right-most index implemented as RAM port, all other indexes become either constants after loop unrolling or a mux tree.
- Should only be used in situations where bank conflicts are clearly not possible.
- Provides best achievable QOR in all cases, including for "non-power of 2" arrays.
- Provides built-in checking for index violations for all array dimensions in both pre-HLS model and in RTL (via PSL or SVA assertions in RTL).
- Usage is same as normal C array, aside from declaration.
- See Matchlib toolkit example 36\*.

#### `ac_shared_bank_array_2D`, `ac_shared_bank_array_3D`, `ac_shared_bank_array_vary`

- Shared banked memory where each bank is a separate RAM instance.
- During HLS, right-most index implemented as RAM port, all other indexes become either constants after loop unrolling or a mux tree.
- Should only be used in situations where bank conflicts are clearly not possible.
- Provides best achievable QOR in all cases, including for "non-power of 2" arrays.
- Provides built-in checking for index violations for all array dimensions in both pre-HLS model and in RTL (via PSL or SVA assertions in RTL).
- Usage is same as normal C array, aside from declaration.
- See Matchlib toolkit example 38\*.

#### `ac_wr_mask_array_1D` / `ac_shared_wr_mask_array_1D`

- Simple 1D array implemented as a RAM/ROM with configurable write mask.
- Each element can have 1 or more slices, any or all slices can be written as per the write mask.

- Provides built-in checking for index violations for all array dimensions in both pre-HLS model and in RTL (via PSL or SVA assertions in RTL).
- Usage is same as normal C array, aside from declaration.
- See Matchlib toolkit example 33\*.

### Scratchpad

- Banked memory where low order bits of address select bank (like Catapult interleave directive).
- Requests routed thru crossbar to proper bank, responses routed thru crossbar from proper bank.
- Emits error message on any bank conflicts.
- No arbitration or queueing, backpressure is not possible.
- Supports memory transaction log file generation.
- See Matchlib toolkit example 41\* and later discussion in this document.

### ArbitratedScratchpad

- Banked memory where low order bits of address select bank (like Catapult interleave directive).
- Requests routed thru crossbar to proper bank, responses routed thru crossbar from proper bank.
- Bank conflicts are allowed.
- Has arbitration and queues to handle potential bank conflicts.
- Incoming requests subject to backpressure due to potential bank conflicts.
- See Matchlib toolkit example 42\* and later discussion in this document.

### extended\_array

- Provides memory transaction log file generation for debugging.
- Inherits from `ac_array_1D`, and can be used in place of it.
- Can also be used in place of `ac_shared_array_1D` for pre-HLS sim only.
- Usage is same as normal C array, aside from declaration.
- Emits error message on uninitialized memory reads (UMR) in pre-HLS simulation only
- See Matchlib example 12\* and "memory\_logging\_and\_debug.pdf" in the Matchlib toolkit doc directory.

### `ac_array_subrange<>`

- Creates a subrange array from an existing 1D array class object.
- The size and offset of the subrange with respect to the base array are specified at runtime by user.
- The subrange array is used like a normal independent array object.
- Automatic checking for index violations occurs in both the pre-HLS and post-HLS simulations wrt the subrange.
- See Matchlib example 39\*.

The following classes are mentioned here for completeness, but should never be used to model RAMs/ROMs:

#### ac\_interleave\_array

- Should never be used, use Scratchpad or ArbitratedScratchpad instead.
- Scratchpad provides predictable usage model during HLS and best QOR.
- Scratchpad emits error messages for bank conflicts in pre-HLS simulation.
- ArbitratedScratchpad handles more complex cases that would result in scheduling failures if ac\_interleave\_array was instead used.

#### ac\_array

- Should never be used to model RAMs/ROMs.
- Supports from 1 to 3 dimensions.
- By default, does not provide index bounds checking.
- Does not provide good QOR for "non-power of 2" memories.
- ac\_bank\_array\* provides better QOR for cases where number of dimensions is 2 or more.

### Discussion of Scratchpad and ArbitratedScratchpad

The rest of this document provides an in-depth discussion of Scratchpad and ArbitratedScratchpad. We begin by introducing three simple design examples to frame the discussion.

#### Example Design #1

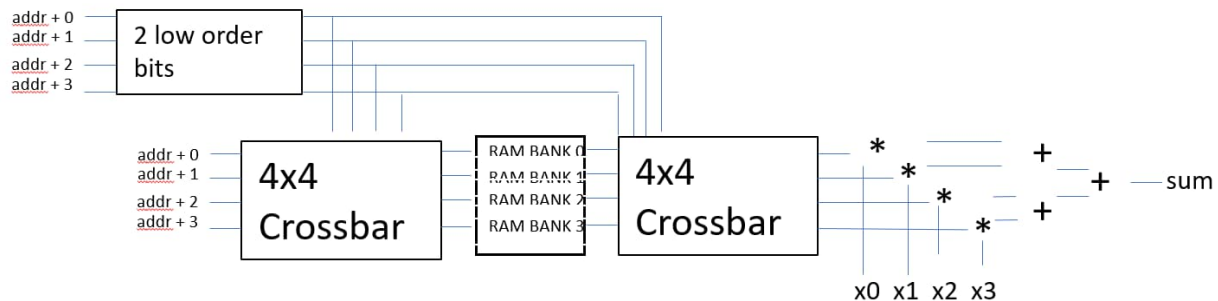
```
uint16 coeffs[1024];

void main() {
    wait();

#pragma hls_pipeline_init_interval 1
#pragma pipeline_stall_mode flush
    while (1) {
        uint16 addr = addr_in.Pop();
        uint16 x0 = in0.Pop();
        uint16 x1 = in1.Pop();
        uint16 x2 = in2.Pop();
        uint16 x3 = in3.Pop();
        uint16 sum = (x0 * coeffs[addr + 0]) +
                     (x1 * coeffs[addr + 1]) +
                     (x2 * coeffs[addr + 2]) +
                     (x3 * coeffs[addr + 3]);
        out1.Push(sum);
    }
}
```

This is a sketch of a simple multiply accumulate (or MAC) design to be synthesized thru HLS. This design is like example 41\* in the Catapult Matchlib examples. The design goal is an II=1, so a new "sum" output will be produced on every clock cycle. For this to be achieved the coeffs array must support 4 uint16 reads per clock cycle. We can meet this goal by using a banked memory architecture where the bank selection is the two low order address bits. It is easy to analyze the indexes of the array accesses and prove that there will be no bank conflicts on each cycle. This is true even if "addr" is not evenly divisible by 4, which is a case that the HW for this design would need to support.

The HW that must be generated for this design will be similar to the following:



The crossbars shown above are designed under the assumption that there are never any conflicts in routing inputs to outputs. This enables the HW to be simple, and also means that there is no need for arbitration or queueing of any of the memory read requests. All the read requests have a fixed delay to produce their outputs. None of the HW in the diagram above ever applies backpressure on upstream requests.

## Example Design #2

```
uint16 coeffs[1024];

void main() {
    wait();

#pragma hls_pipeline_init_interval 1
#pragma pipeline_stall_mode flush
    while (1) {
        uint16 addr = addr_in.Pop();
        uint16 x0 = in0.Pop();
        uint16 x1 = in1.Pop();
        uint16 x2 = in2.Pop();
        uint16 x3 = in3.Pop();
        uint16 sum = (x0 * coeffs[func0(addr)]) +
                     (x1 * coeffs[func1(addr)]) +
                     (x2 * coeffs[func2(addr)]) +
                     (x3 * coeffs[func3(addr)]);
        out1.Push(sum);
    }
}
```

This design is the same as Example #1, except that the coeffs indexes now depend on func0, func1, etc. If it can be proved that the low order address bits of these functions will never result in any bank conflicts, then the same HW as shown in example #1 can be used. However, if it cannot be proved that there will not be any bank conflicts, then more complex HW will need to be created. This HW will look similar to the HW in Example #1, but it will need to insert arbiters to handle competing requests to the same RAM bank. Because there may be bank contention, bank requests may need to be delayed, which means that the coeffs read operations now need to apply backpressure to the calling process when contention occurs. This possible backpressure needs to be handled carefully so that the overall design can continue executing and making progress, otherwise the whole design would deadlock.

When there is bank contention, clearly a new "sum" output will no longer be guaranteed to be produced on every clock edge.

### Example Design #3

```
uint16 coeffs[1024];

void main() {
    wait();

#pragma hls_pipeline_init_interval 1
#pragma pipeline_stall_mode flush
    while (1) {
        uint16 addr0 = addr_in0.Pop();
        uint16 addr1 = addr_in1.Pop();
        uint16 addr2 = addr_in2.Pop();
        uint16 addr3 = addr_in3.Pop();
        uint16 x0 = in0.Pop();
        uint16 x1 = in1.Pop();
        uint16 x2 = in2.Pop();
        uint16 x3 = in3.Pop();
        uint16 sum = (x0 * coeffs[addr0]) +
                     (x1 * coeffs[addr1]) +
                     (x2 * coeffs[addr2]) +
                     (x3 * coeffs[addr3]);
        out1.Push(sum);
    }
}
```

This design is similar to Example #1 and Example #2, but now all of the indexes into the coeffs array are addresses that are read from an input channel. Thus, there is no way for an HLS tool to prove that there will never be any bank contention on the memory read operations. It is possible that the designer may know that the externally provided addresses will never result in bank contention. If this is the case, then the simpler and smaller HW described in Example #1 perhaps could still be used, rather than using the arbitration logic described in Example #2.

### Catapult Memory Banking Directives

Catapult supports the INTERLEAVE and BLOCK\_SIZE directives to automatically build banked memories. With the directive "-INTERLEAVE 4" set on the coeffs memory shown in Example #1, Catapult will do the index analysis and prove there are no bank conflicts and build HW similar to what is described in Example #1.

For Example #2, with the same INTERLEAVE directive, Catapult may or may not succeed in proving there are no bank conflicts. If it cannot prove there are no bank conflicts, it will not build a banked memory with arbitration, instead it will effectively give up on building a banked memory altogether. This may not result in any error or warning messages, but it will result in the  $II=1$  goal not being met.

For Example #3, with the same INTERLEAVE directive, Catapult definitely will not succeed in proving there are no bank conflicts. It will not build a banked memory, and the  $II=1$  goal will not be met.

So, we see that use of the Catapult INTERLEAVE directive for Examples #2 and #3 do not meet our top two goals for the memory architecture modeling methodology:

1. Meet PPA (power, performance, area) goals without any compromises.
2. Discover and resolve all functional and performance issues related to memories in the pre-HLS model, not in the post-HLS model.

Because of these problems, we recommend that Catapult Matchlib users never directly use BLOCK\_SIZE, INTERLEAVE, and WORD\_WIDTH directives in their models. (BLOCK\_SIZE and WORD\_WIDTH have very similar issues as those described directly above).

### Recommended Matchlib Memory Architecture Modeling Methodology

The recommended Matchlib memory architecture modeling methodology is to use memory architecture models in the pre-HLS design that explicitly specify the desired memory architecture. This guarantees that HLS will build exactly what the user wants, and it also enables all performance and functional problems related to the memory architecture to be easily observed in the pre-HLS simulation.

### Scratchpad

A commonly-used Matchlib memory architecture model is called Scratchpad, and the model is currently located in the Catapult Matchlib examples toolkit include directory, called ScratchpadClass.h.

This model implements the simple crossbar structure described in the Example #1 scenario above. It has no arbitration or queueing of requests, thus all bank requests must be free of conflicts. In the pre-HLS simulation, if there are any bank conflicts, it will emit an error message at the time of the conflict. Because the model emits errors when there are bank conflicts, it is safe to use for the scenarios described in Example #2 and Example #3. In the latter case, if the user knows that the externally-provided addresses are free of bank conflicts, then the optimized HW can still be used.

The Scratchpad has these characteristics:

- The template parameters are the word\_type (ie the element type that the overall array stores), the number of banks, and the total capacity of the banked memory in words.



- Each request to the banked memory includes multiple inputs. The number of inputs is equal to the number of banks. Each input is on its own "lane".
- Similarly, each response includes multiple outputs.
- One new request is consumed on each invocation.
- One new response is produced on each invocation if the requesting operation is a load (i.e. a read).
- The model can be pipelined with  $II=1$  or any other desired  $II$ .
- There are no variations in delay since there is no arbitration or backpressure.
- The scratchpad can be instantiated in your model as a basic class (ScratchpadClass<>) which is inlined into the calling process, or it can be instantiated as a separate SC module (Scratchpad<>) which has its own process.


The following code for the DUT in example 41\* shows its usage, this usage is very similar to Design #1 above:

```

54 void run() {
55     in1.Reset();
56     out1.Reset();
57     wait();
58
59 #pragma hls_pipeline_init_interval 1
60 #pragma pipeline_stall_mode flush
61     while (1) {
62         // get the input request from the testbench
63         dut_in_t req1 = in1.Pop();
64
65         local_mem::scratchpad_req_t sp_req; // local scratchpad request type
66
67         // copy incoming request to scratchpad request
68 #pragma hls_unroll yes
69         for (int i=0 ; i < local_mem::num_inputs; i++)
70             sp_req.set(i, req1.addr + i, req1.data[i]);
71
72         if (req1.is_load)
73         {
74             // if it is a load (i.e. read) operation, get the read data from the RAM
75             local_mem::base_rsp_t rsp = scratchpad1.load(sp_req);
76
77             // compute MAC
78             local_mem::word_type sum=0;
79 #pragma hls_unroll yes
80             for (int i=0; i < local_mem::num_inputs; i++) {
81                 sum += rsp.data[i] * req1.data[i];
82             }
83
84             // Push out the sum
85             out1.Push(sum);
86         }
87         else
88         {
89             // if it is a store (i.e. write) operation, write the data to the RAM
90             scratchpad1.store(sp_req);
91         }
92     }
93 }
94 };

```

The following are the pre-HLS waveforms for example 41\*:



During the first 4000 ns 4000 write requests are done to the banked RAM, where each request has 16 separate lanes. During the last 1000 ns, 1000 read requests are done to the banked RAM, where each request has 16 separate lanes.

You can observe that the RAM is fully (and optimally) utilized since all the rdy/vld signals are high when they can be. Note that this means that Catapult has successfully pipelined this design with an II=1. There is a non-trivial amount of logic within the pipeline due to the two crossbar instantiations. To see the actual logic generated by Catapult for scratchpad for example 41\*, we can reduce the number of banks from the original 16 down to 4 to make it more easily visible. This design then closely corresponds to Example #1. The schedule view is shown below, and you should note its close correspondence to the schematic diagram shown for Example #1.



## ArbitratedScratchpad

If your design needs a banked memory but there may be conflicts among bank request inputs, then a more complex model is needed. This model is ArbitratedScratchpad. It is very similar in concept to Scratchpad, but it has arbitration and queueing of request inputs since there may be conflicts among the banks.

ArbitratedScratchpad usage is shown in Catapult Matchlib example 42\*. This example is very similar to example 41\* described above.


The following code for the DUT in example 42\* shows its usage.

```

64 void main() {
65     out1.Reset();
66     in1.Reset();
67
68     bool was_consumed[local_mem::num_inputs];
69     bool all_consumed = 1;
70
71     wait();
72
73     local_mem::req_t req;
74     local_mem::rsp_t rsp;
75
76 #pragma hls_pipeline_init_interval 1
77 #pragma pipeline_stall_mode flush
78     while (1) {
79         // if there are remaining requests in some lanes that were not yet consumed
80         // then dont Pop a new input request
81         if (all_consumed)
82             req = in1.Pop();
83
84         mem.load_store(req, rsp, was_consumed);
85
86         bool any_valid = 0;
87         all_consumed = 1;
88 #pragma hls_unroll yes
89         for (unsigned i=0; i < local_mem::num_inputs; i++) {
90             // check if some of the requests were not consumed..
91             if (was_consumed[i] == 0)
92                 all_consumed = 0;
93
94             // if a particular request was consumed, then it is no longer valid
95             if (was_consumed[i] == 1)
96                 req.valids[i] = 0;
97
98             // check if we have any read (aka "load") responses
99             if (rsp.valids[i] == 1)
100                 any_valid = 1;
101         }
102
103         if (any_valid)
104             out1.Push(rsp);
105
106         wait();
107     }
108 }

```


The following are the pre-HLS waveforms for example 42\*:



During the first 1000 ns 1000 write requests are done to the banked RAM, where each request has 16 separate input lanes. During the last 1000 ns, 1000 read requests are done to the banked RAM, where each request and response has 16 separate lanes.

You can observe that the RAM is almost fully utilized since all the rdy/vld signals are almost always high when they can be. The rdy/vld signals are occasionally not high when they can be because of the arbitration and backpressure, and because the testbench occasionally sends empty requests. Note that Catapult has successfully pipelined this design with an  $II=1$ . There is a non-trivial amount of logic within the pipeline due to the two crossbar instantiations as well as the arbitration and queuing logic.

The following are the post-HLS waveforms for example 42\*. They are basically the same as the pre-HLS waveforms.



The ArbitratedScratchpad has these characteristics:

- Generally, any desired throughput can be achieved simply by increasing the number of banks in the memory and the number of inputs in each request
- Generally, you want the number of banks to be greater (e.g. 4-16x) than number of inputs in each request. This reduces possible bank contention at minimal additional HW cost.
- Bank contention is OK, so no error or warning message is emitted when it occurs.

- Note the "was\_completed" flags in dut.h. These indicate if requests were accepted in previous call, if they were not completed they need to be presented again to the memory.

You should be aware of these ordering concerns when using ArbitratedScratchpad:

- Write requests to the same address on different lanes may not occur in the order that they are presented to the model due to the arbitration and queuing inside the model.
- Read and write requests to the same address on different lanes may not occur in the order that they are presented to the model due to the arbitration and queuing inside the model.
- In the two scenarios above, in all cases, the pre-HLS and post-HLS ArbitratedScratchpad designs will always behave identically (i.e. there will not be any differences).

These ordering issues are inherent to banked memories like this which support arbitration and queueing. They can be properly handled by appropriate synchronization and coordination at higher levels in the design. In example 42\*, the testbench performs all writes to independent addresses, then flushes the writes (by sending several empty write requests), then proceeds to only issue read requests. This is one possible strategy for insuring expected behavior with this model.

In summary, ArbitratedScratchpad is a very powerful and capable model, however it does add additional design considerations and additional HW beyond Scratchpad. When considering using ArbitratedScratchpad, you should always also consider if your design requirements can be changed such that Scratchpad can be used instead.

What about the Catapult BLOCK\_SIZE directive?

Scratchpad and ArbitratedScratchpad effectively implement the Catapult INTERLEAVE directive. If you instead want the higher order address bits to determine bank selection, simply use `ac_int::get_slc<>` and `ac_int::set_slc<>` to form a new address with the bank selection bits located as the low order bits in a transformed address before presenting this transformed to Scratchpad and ArbitratedScratchpad.

What about the Catapult WORD\_WIDTH directive?

We do not recommend using the WORD\_WIDTH directive within Catapult Matchlib models. Instead, use the ScratchpadClass model described above. In this case, the Scratchpad `word_type` should be the same as the data type used for the original array accesses in the pre-HLS model. The Scratchpad `num_banks` parameter should be: `WORD_WIDTH / number_of_bits(word_type)`, where WORD\_WIDTH is the number of bits that would have been used in the Catapult WORD\_WIDTH directive.